

Replica Sets



Nuri Halperin

@nurih | www.plusnconsulting.com


“

Two is One,
and One is None

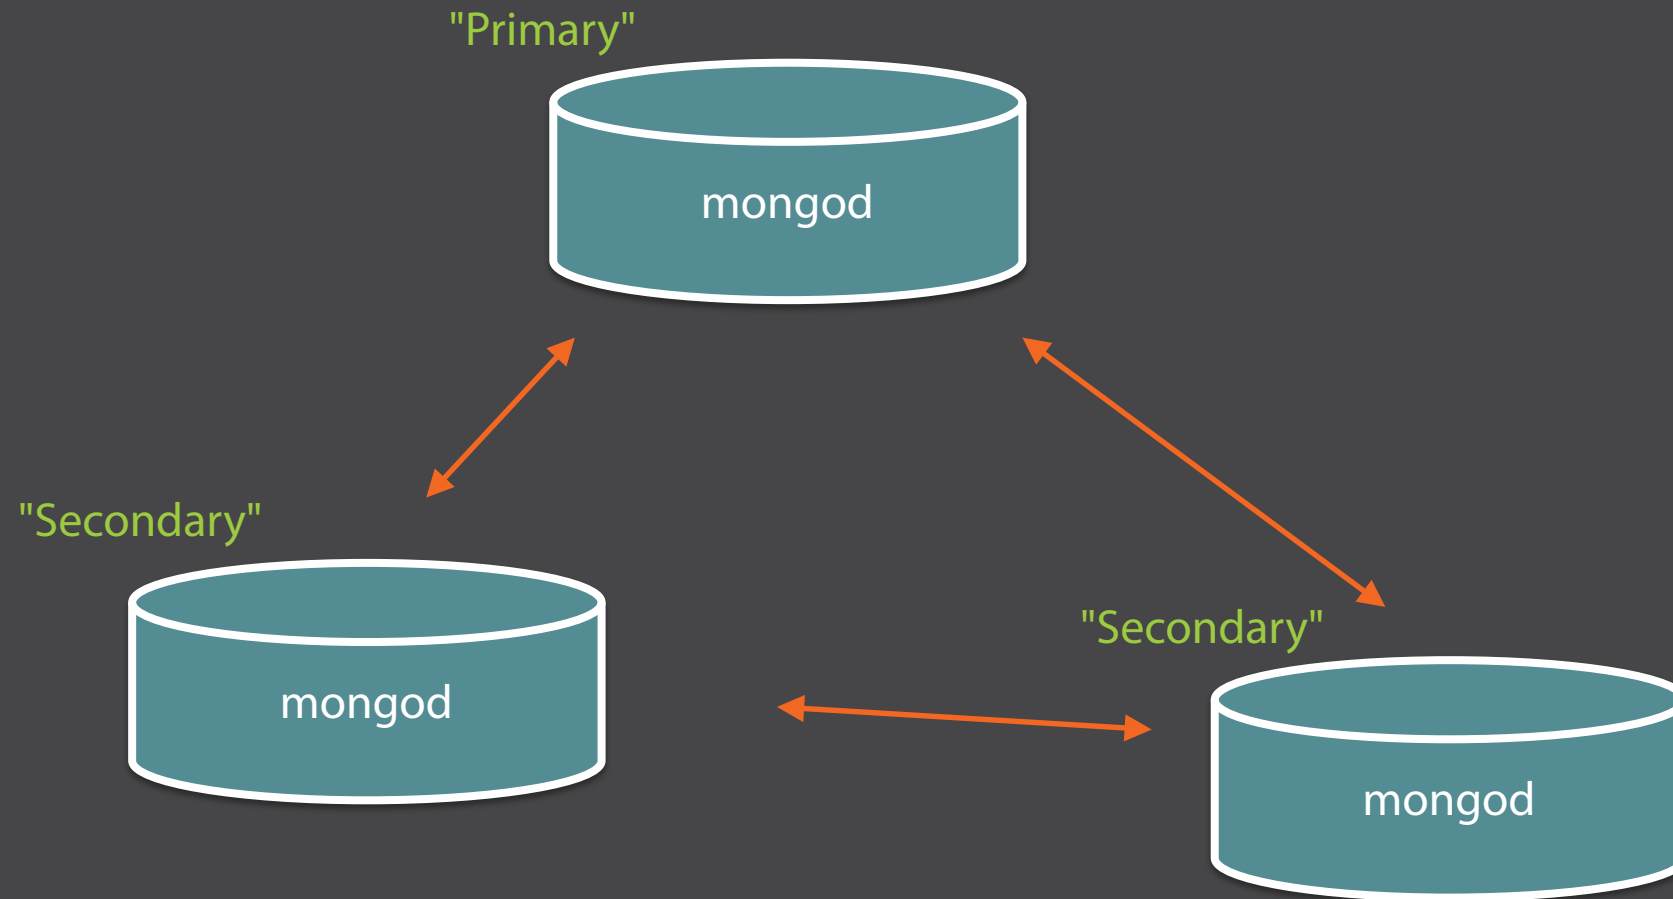
- Unknown

”

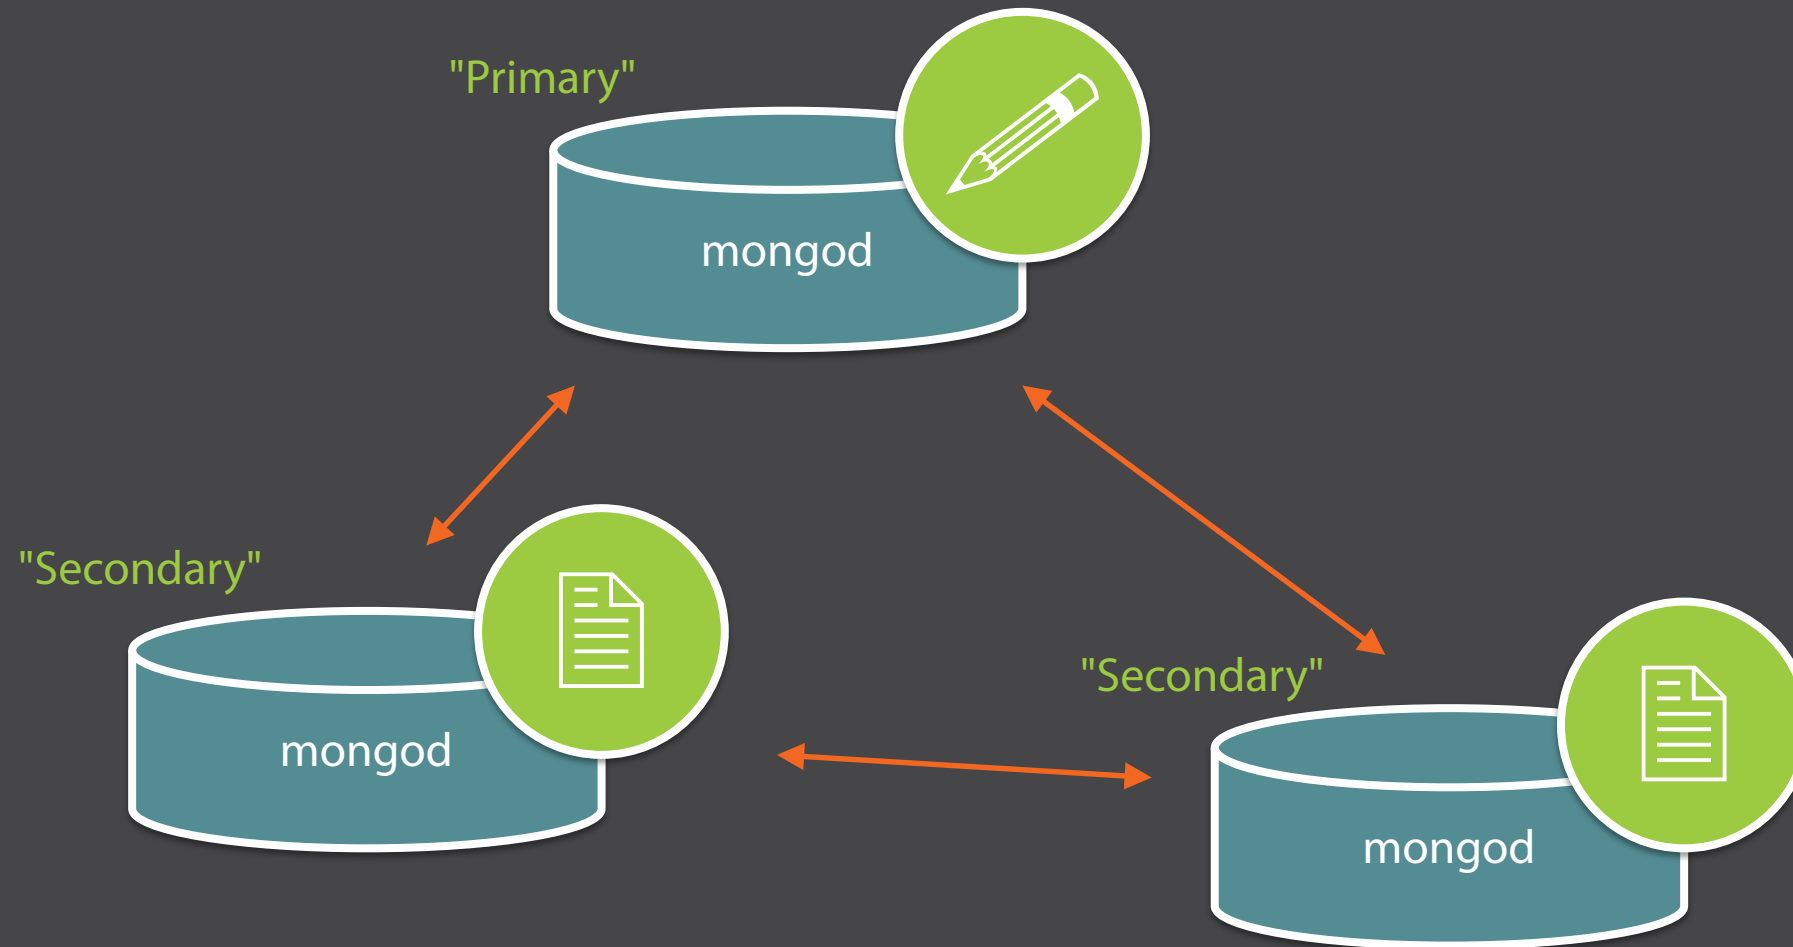
Replica Set



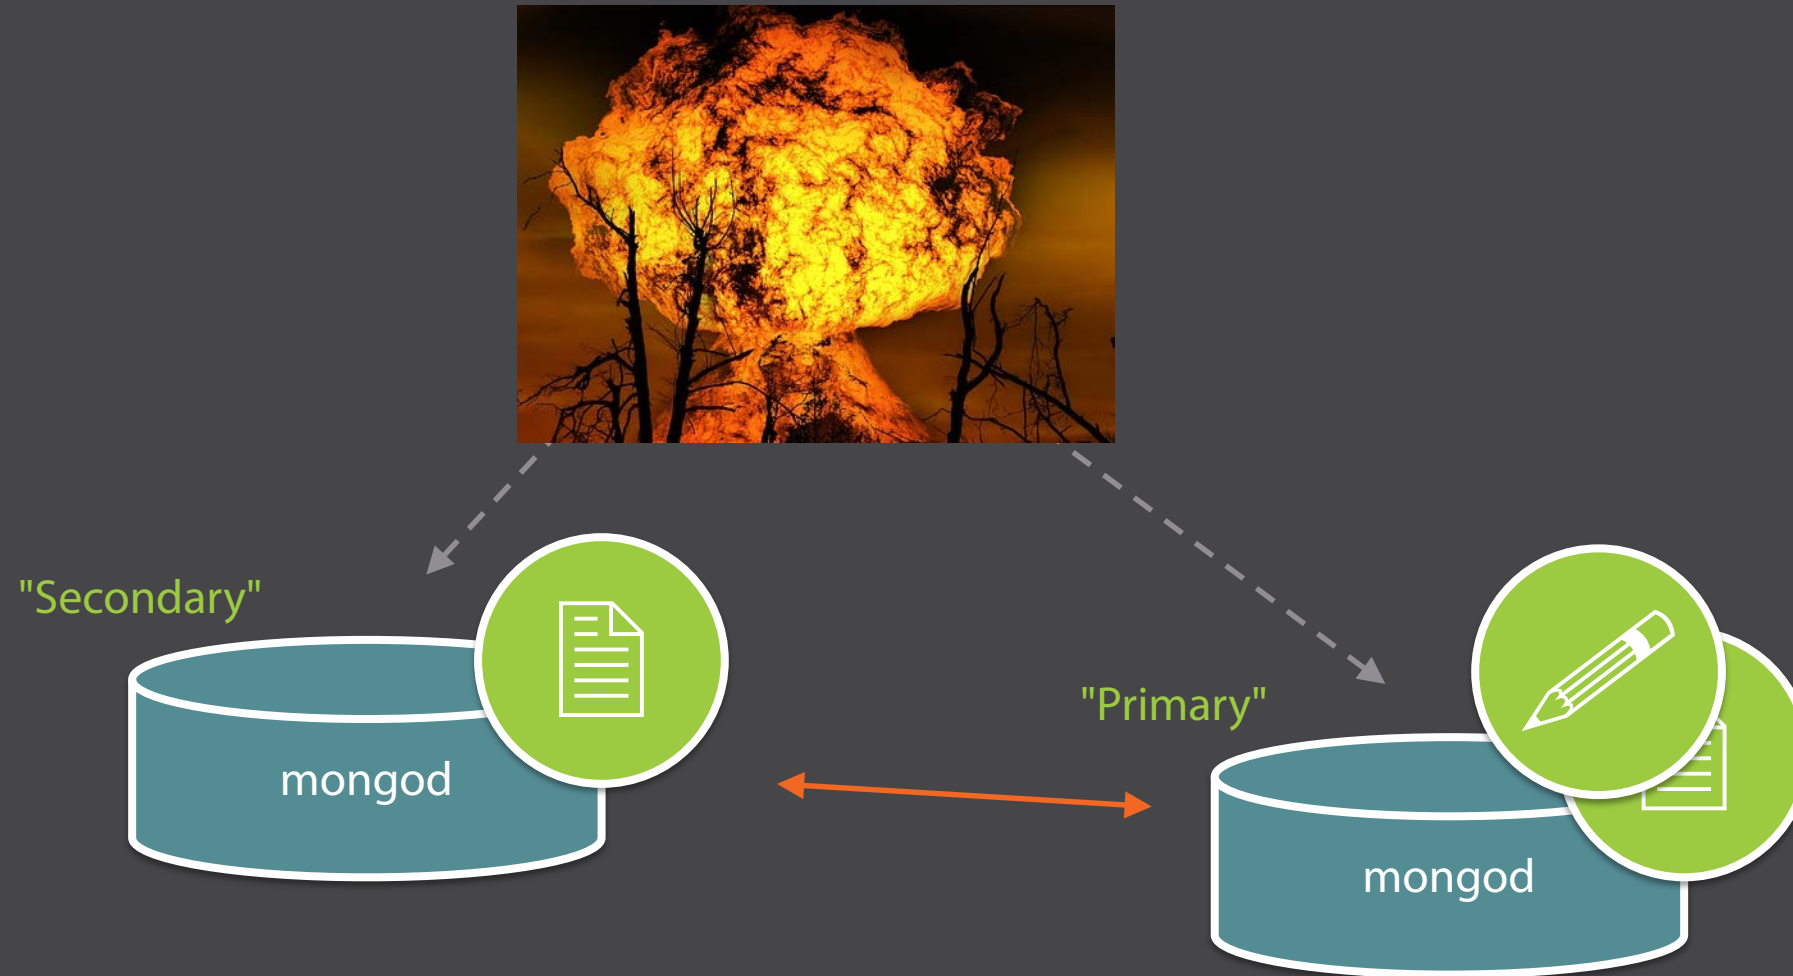
Replica Set




Replica Set



Replica Set



Capped Collection



Capped Collection




Capped Collection



Oplog

`db.demo.insert(...)`



`db.demo.insert(...)`

`db.demo.update(...)`

Replica Set Thus far...

```
// Initialize
```

```
rs.initiate();
```


```
// Get configuration
```

```
var cfg = rs.config();
```


```
// Change configuration
```

```
rs.reconfig(cfg);
```


Elections




Elections




Elections




Elections



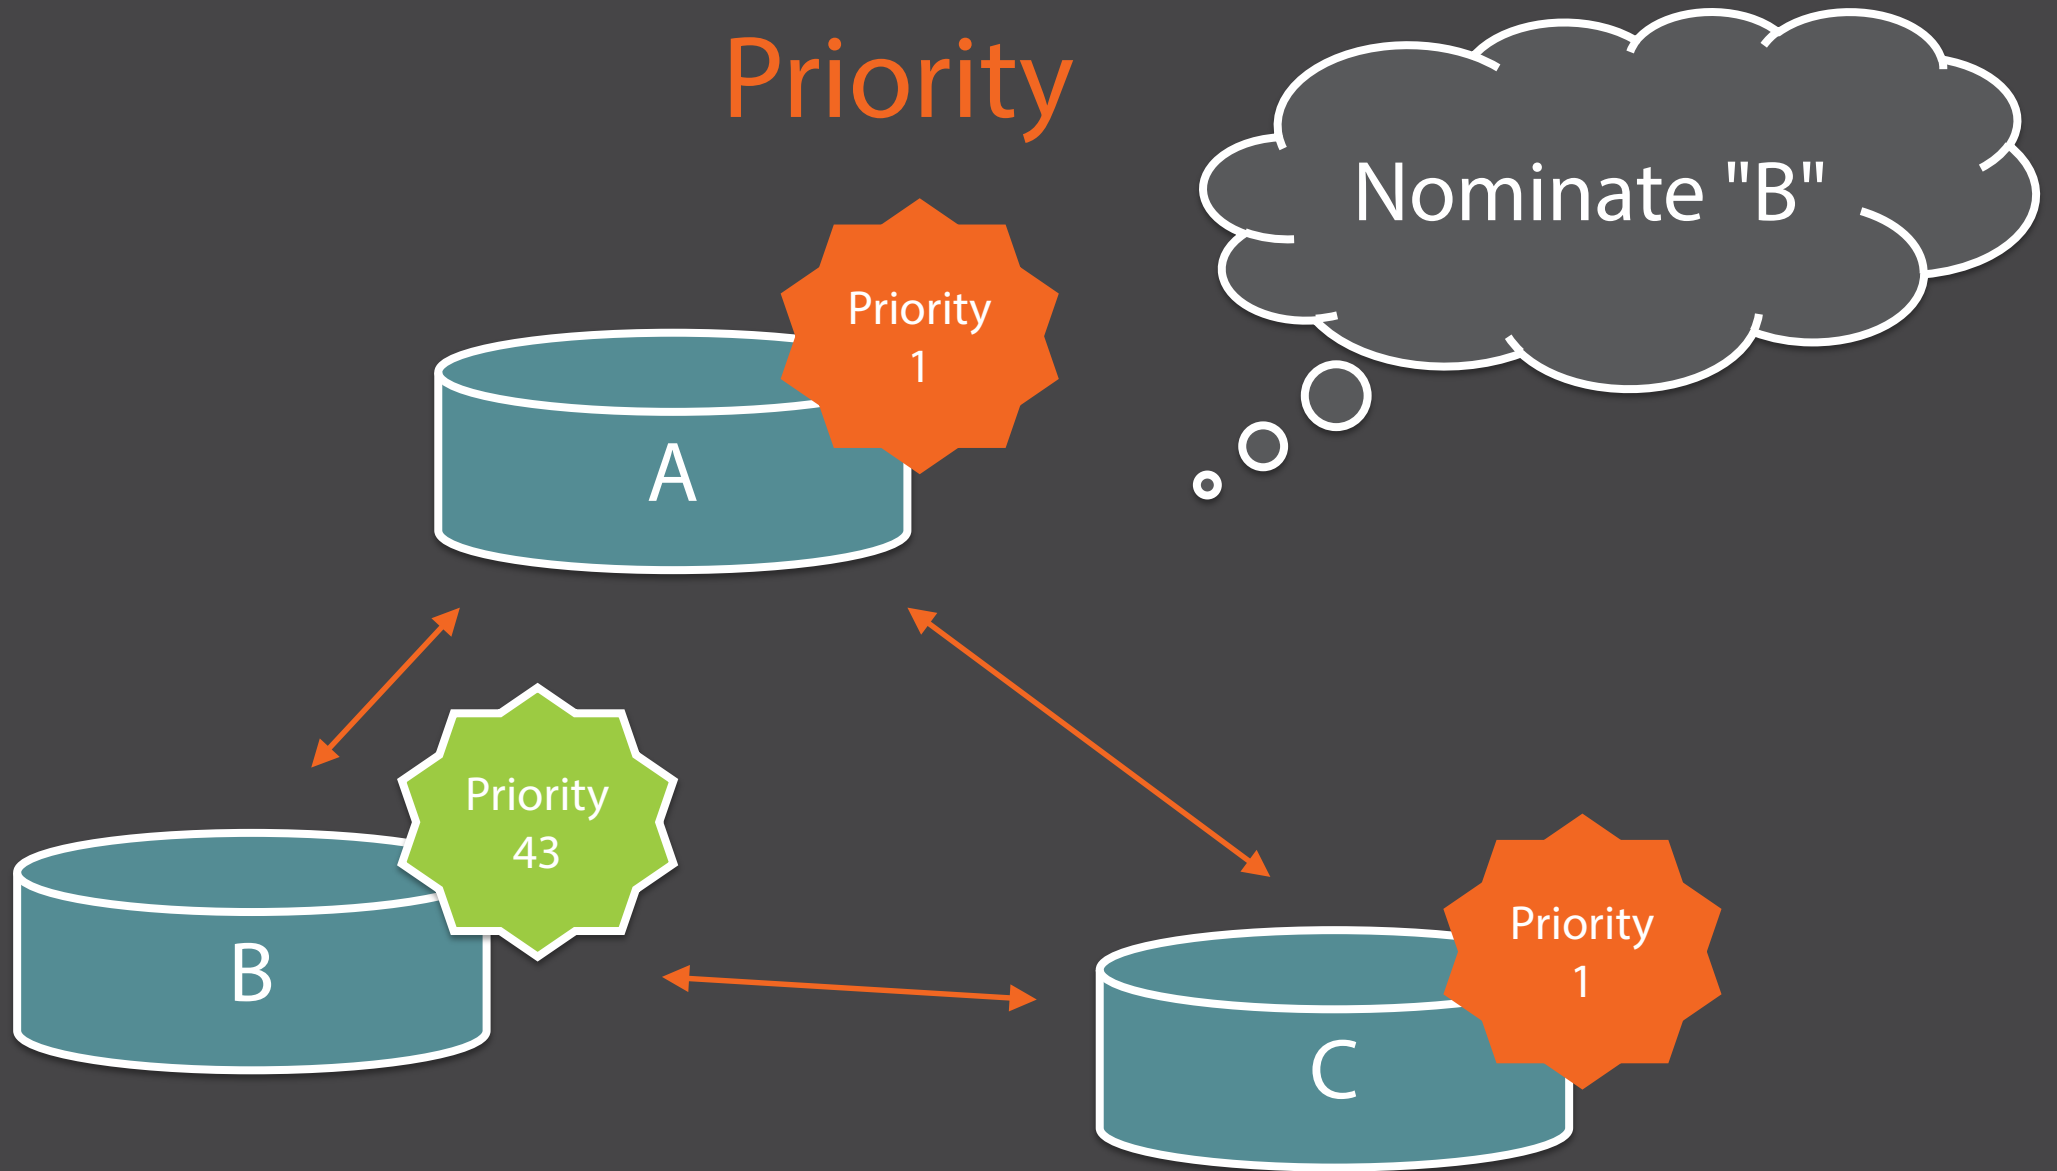
Elections



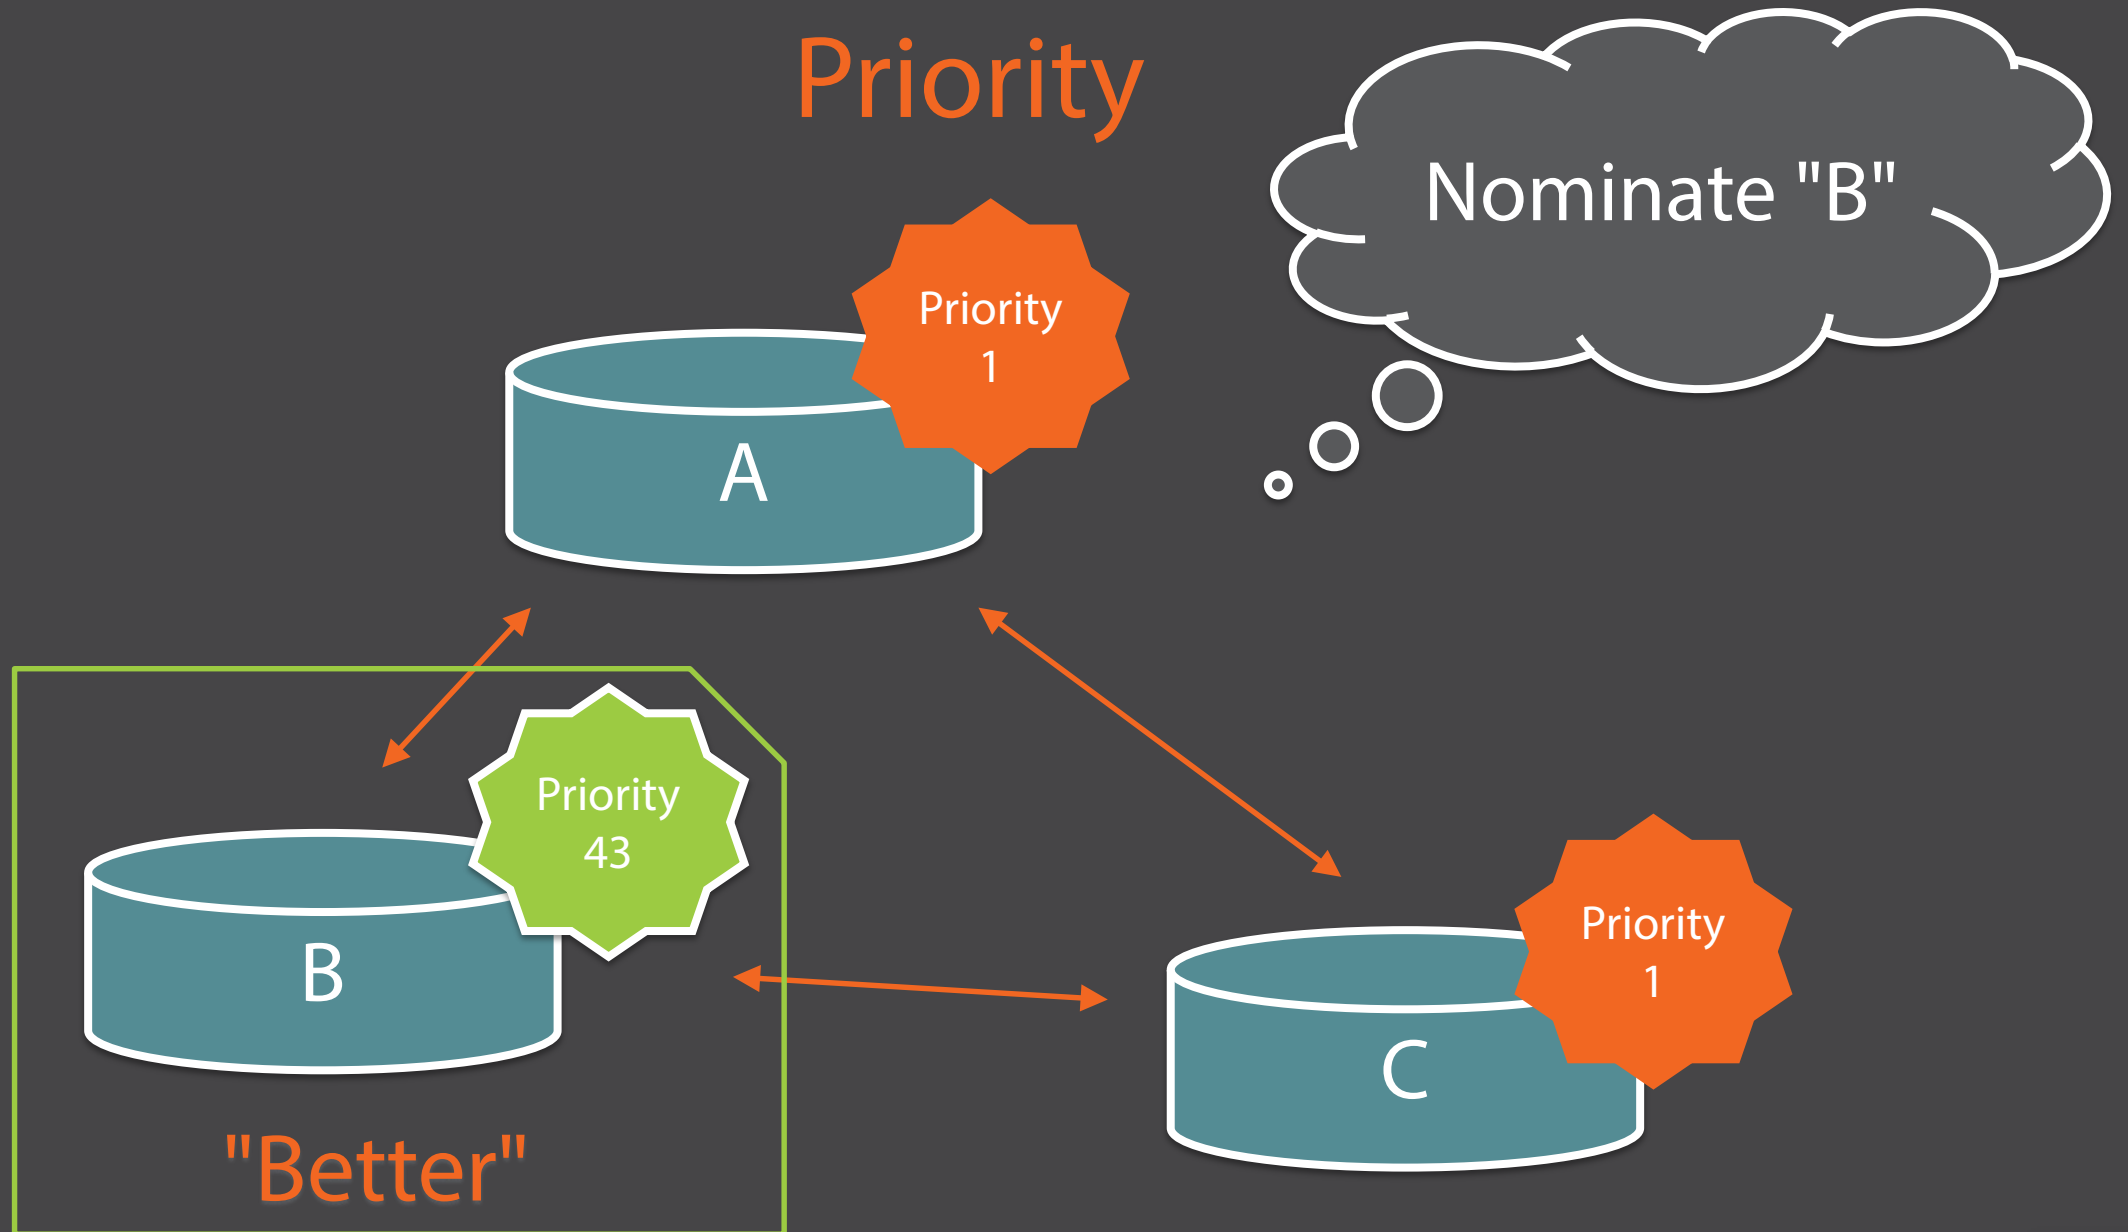
Quorum & Fault Tolerance



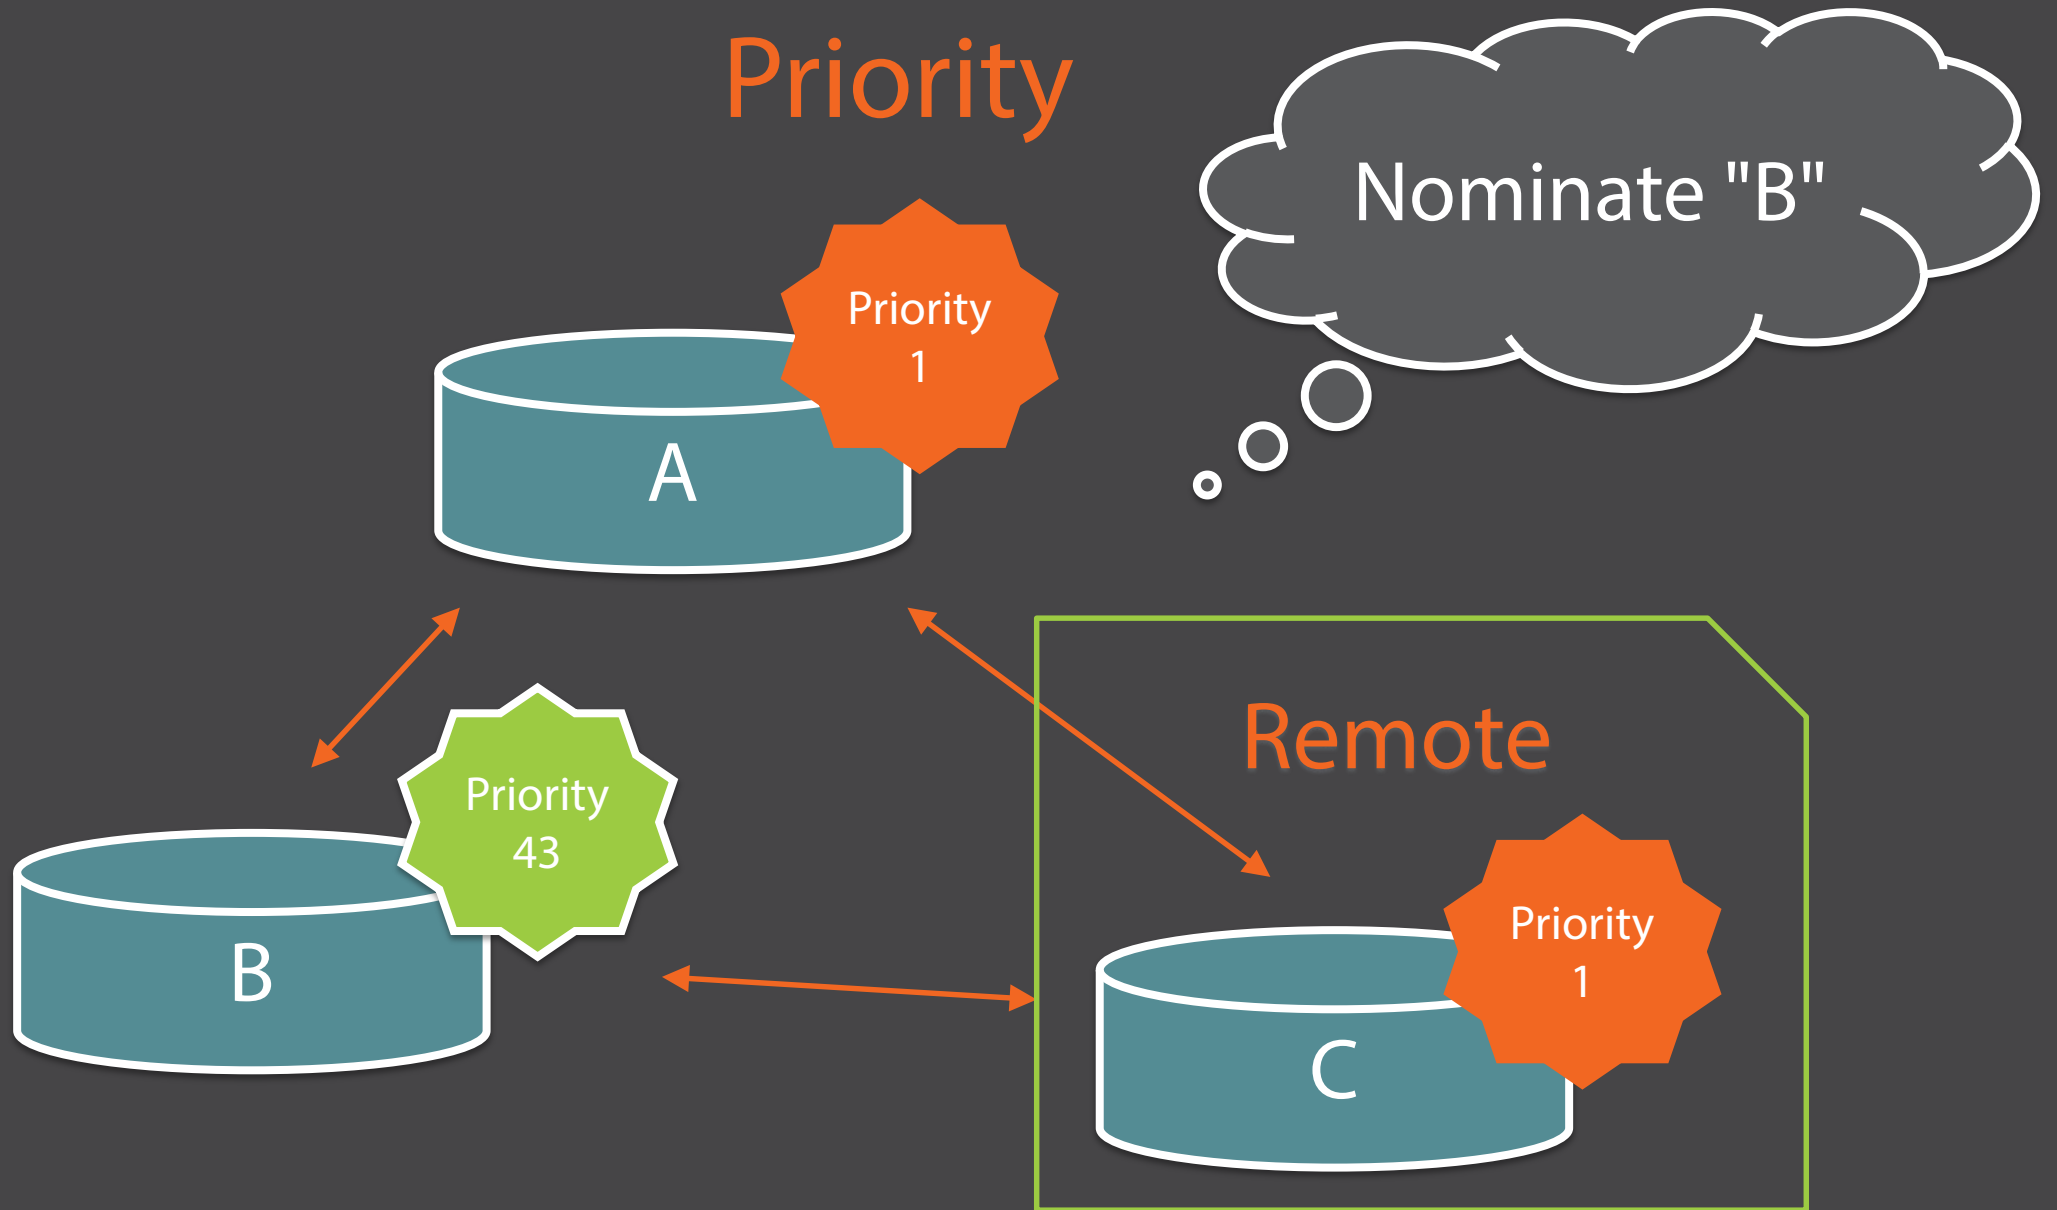
Priority



Priority



Priority



Step Down




"Temporary"
primary




"Usual"
primary



Freeze




Hidden



Hidden

Read Only Application

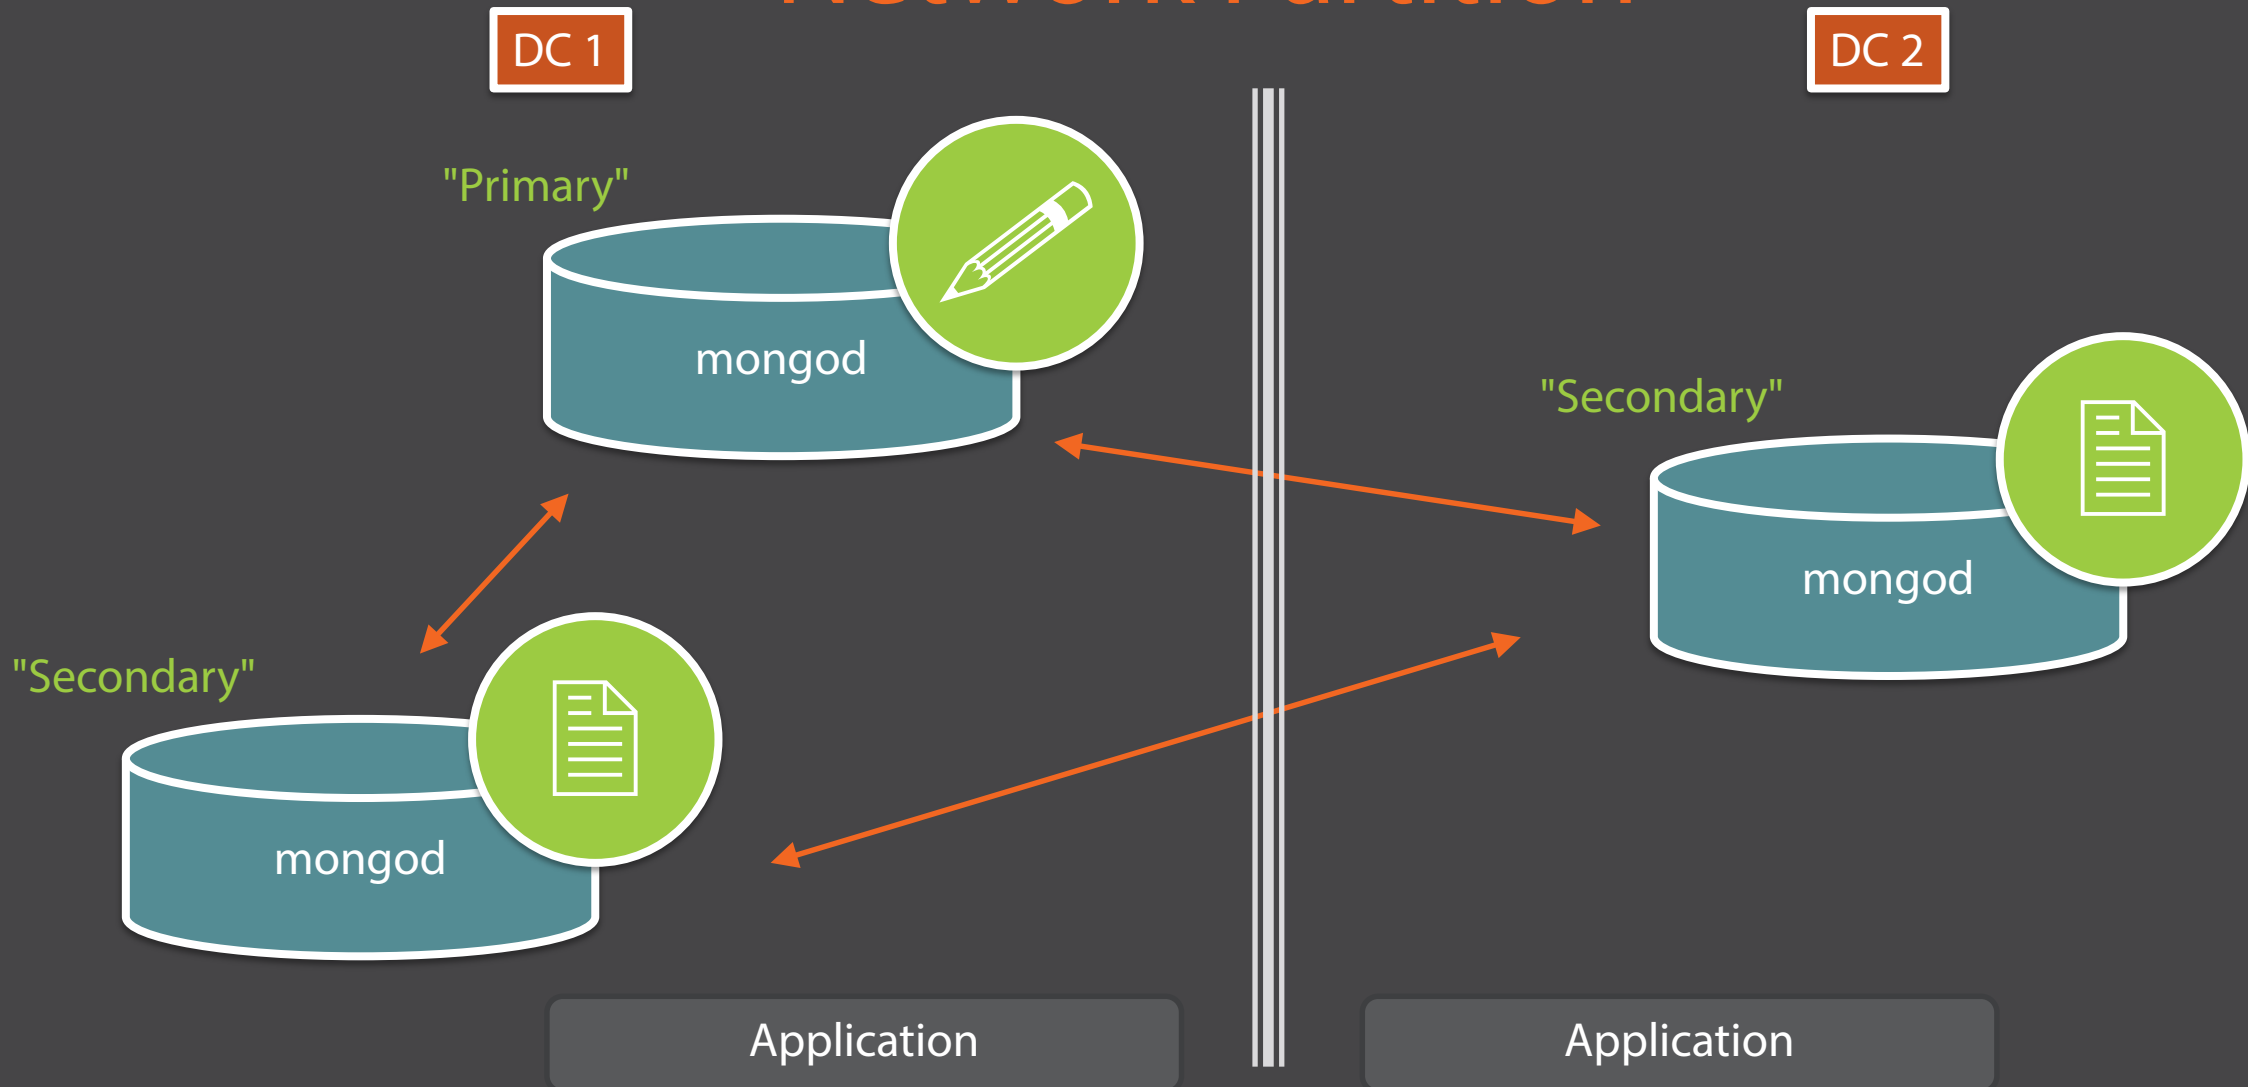
Direct Single
Connection



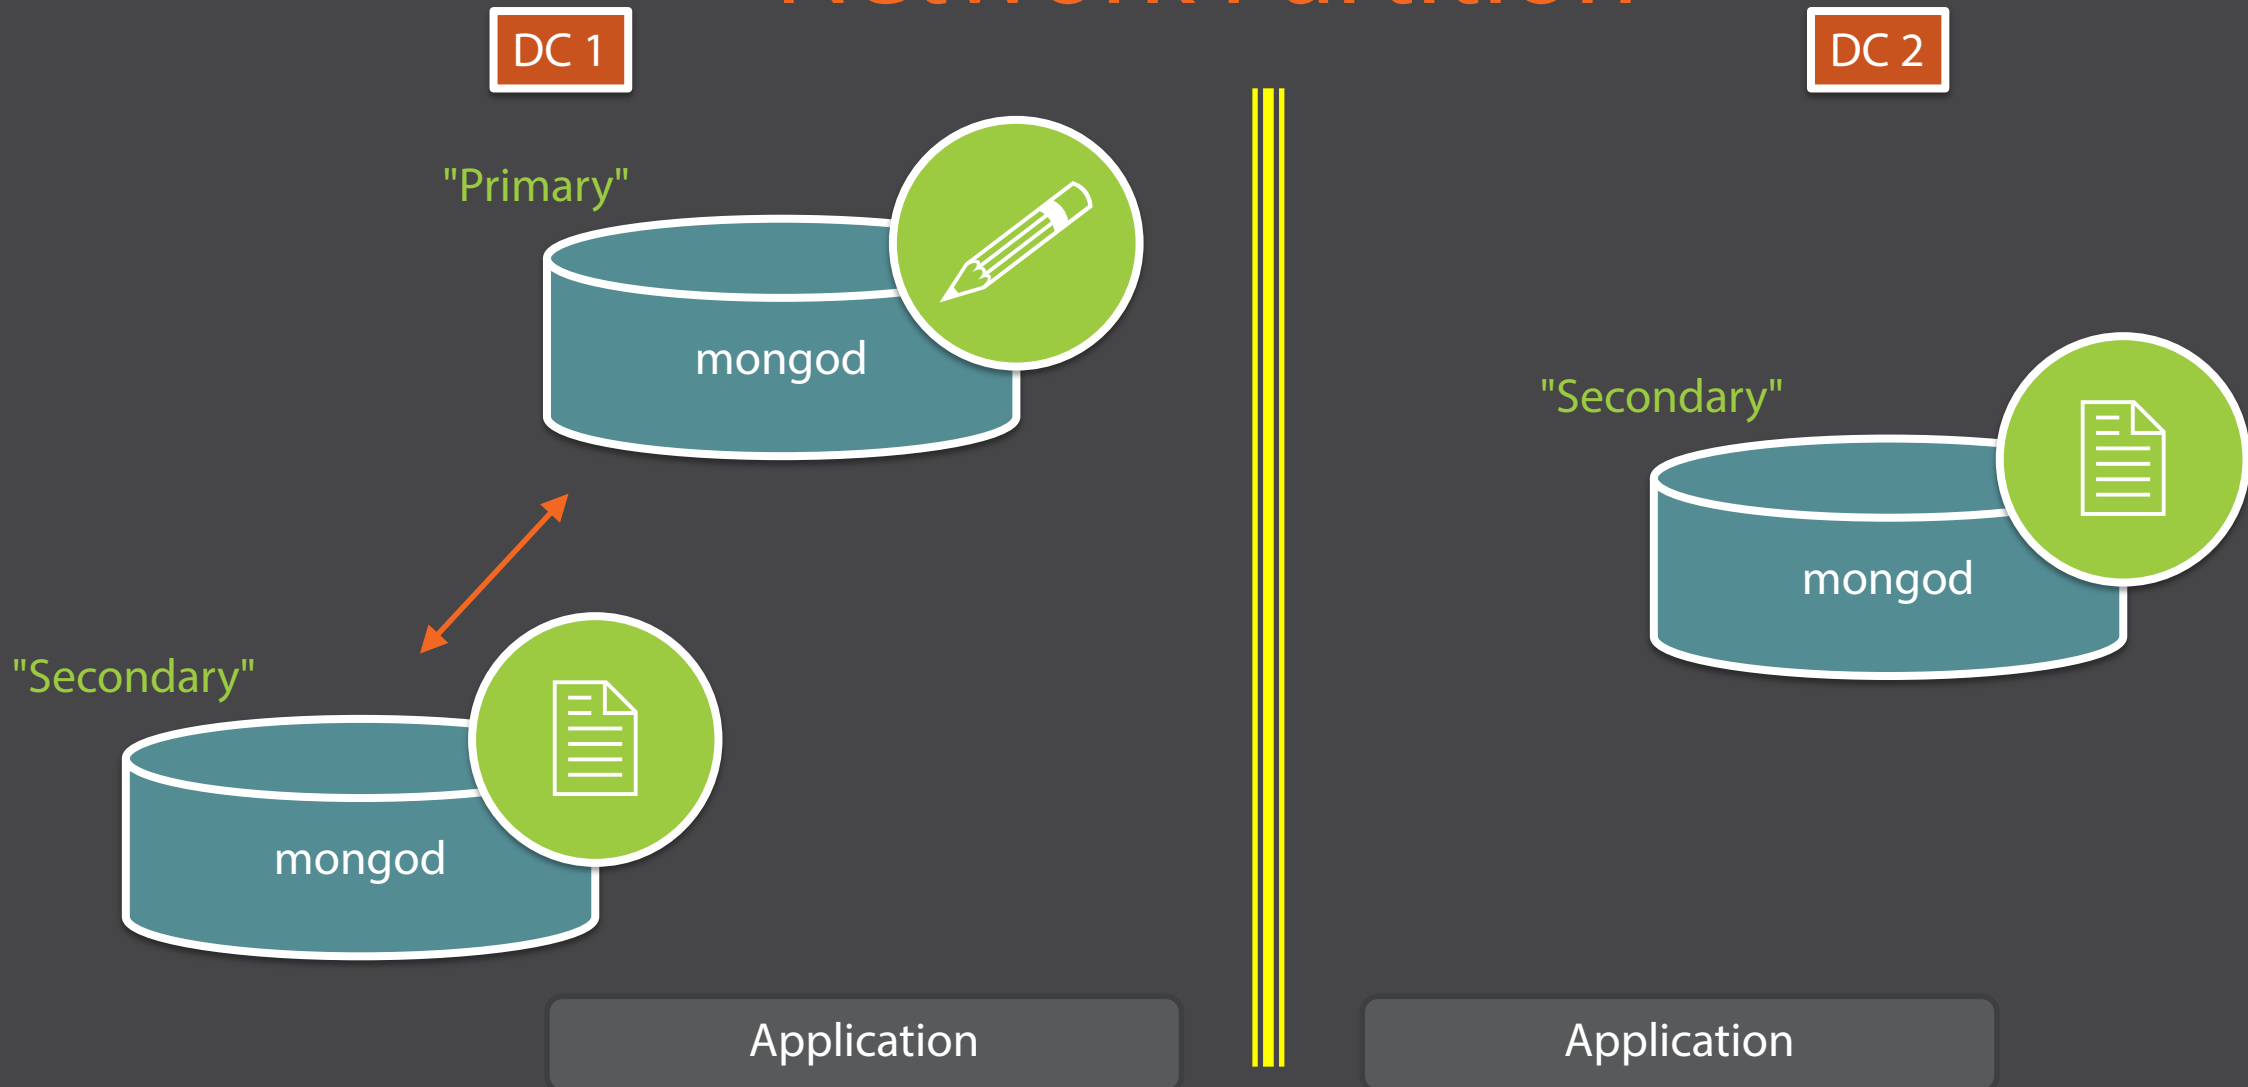
Delayed



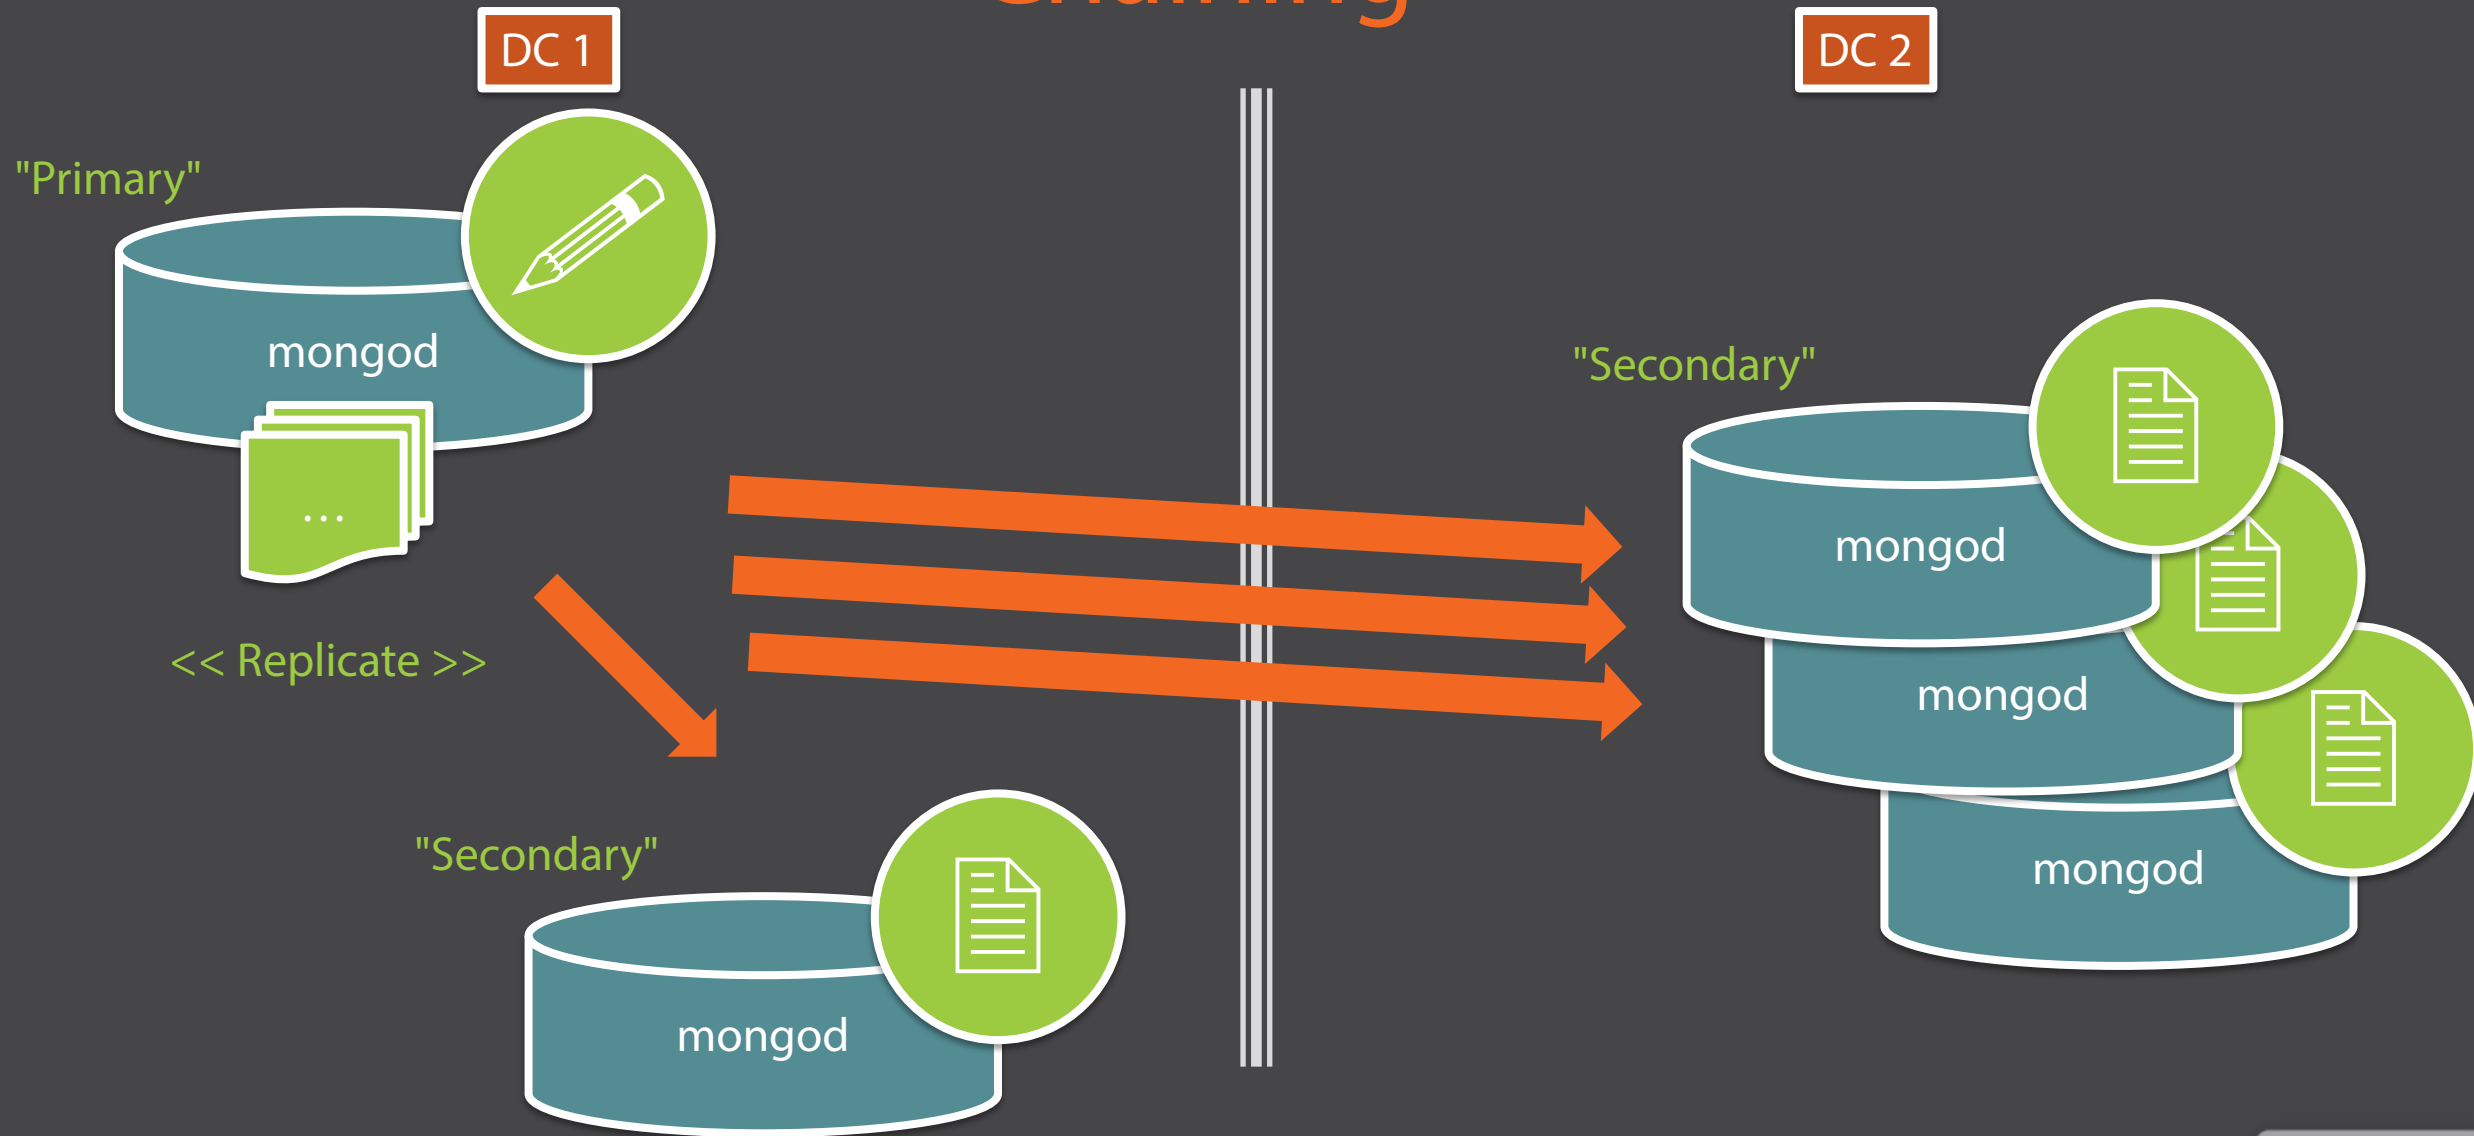
Network Partition



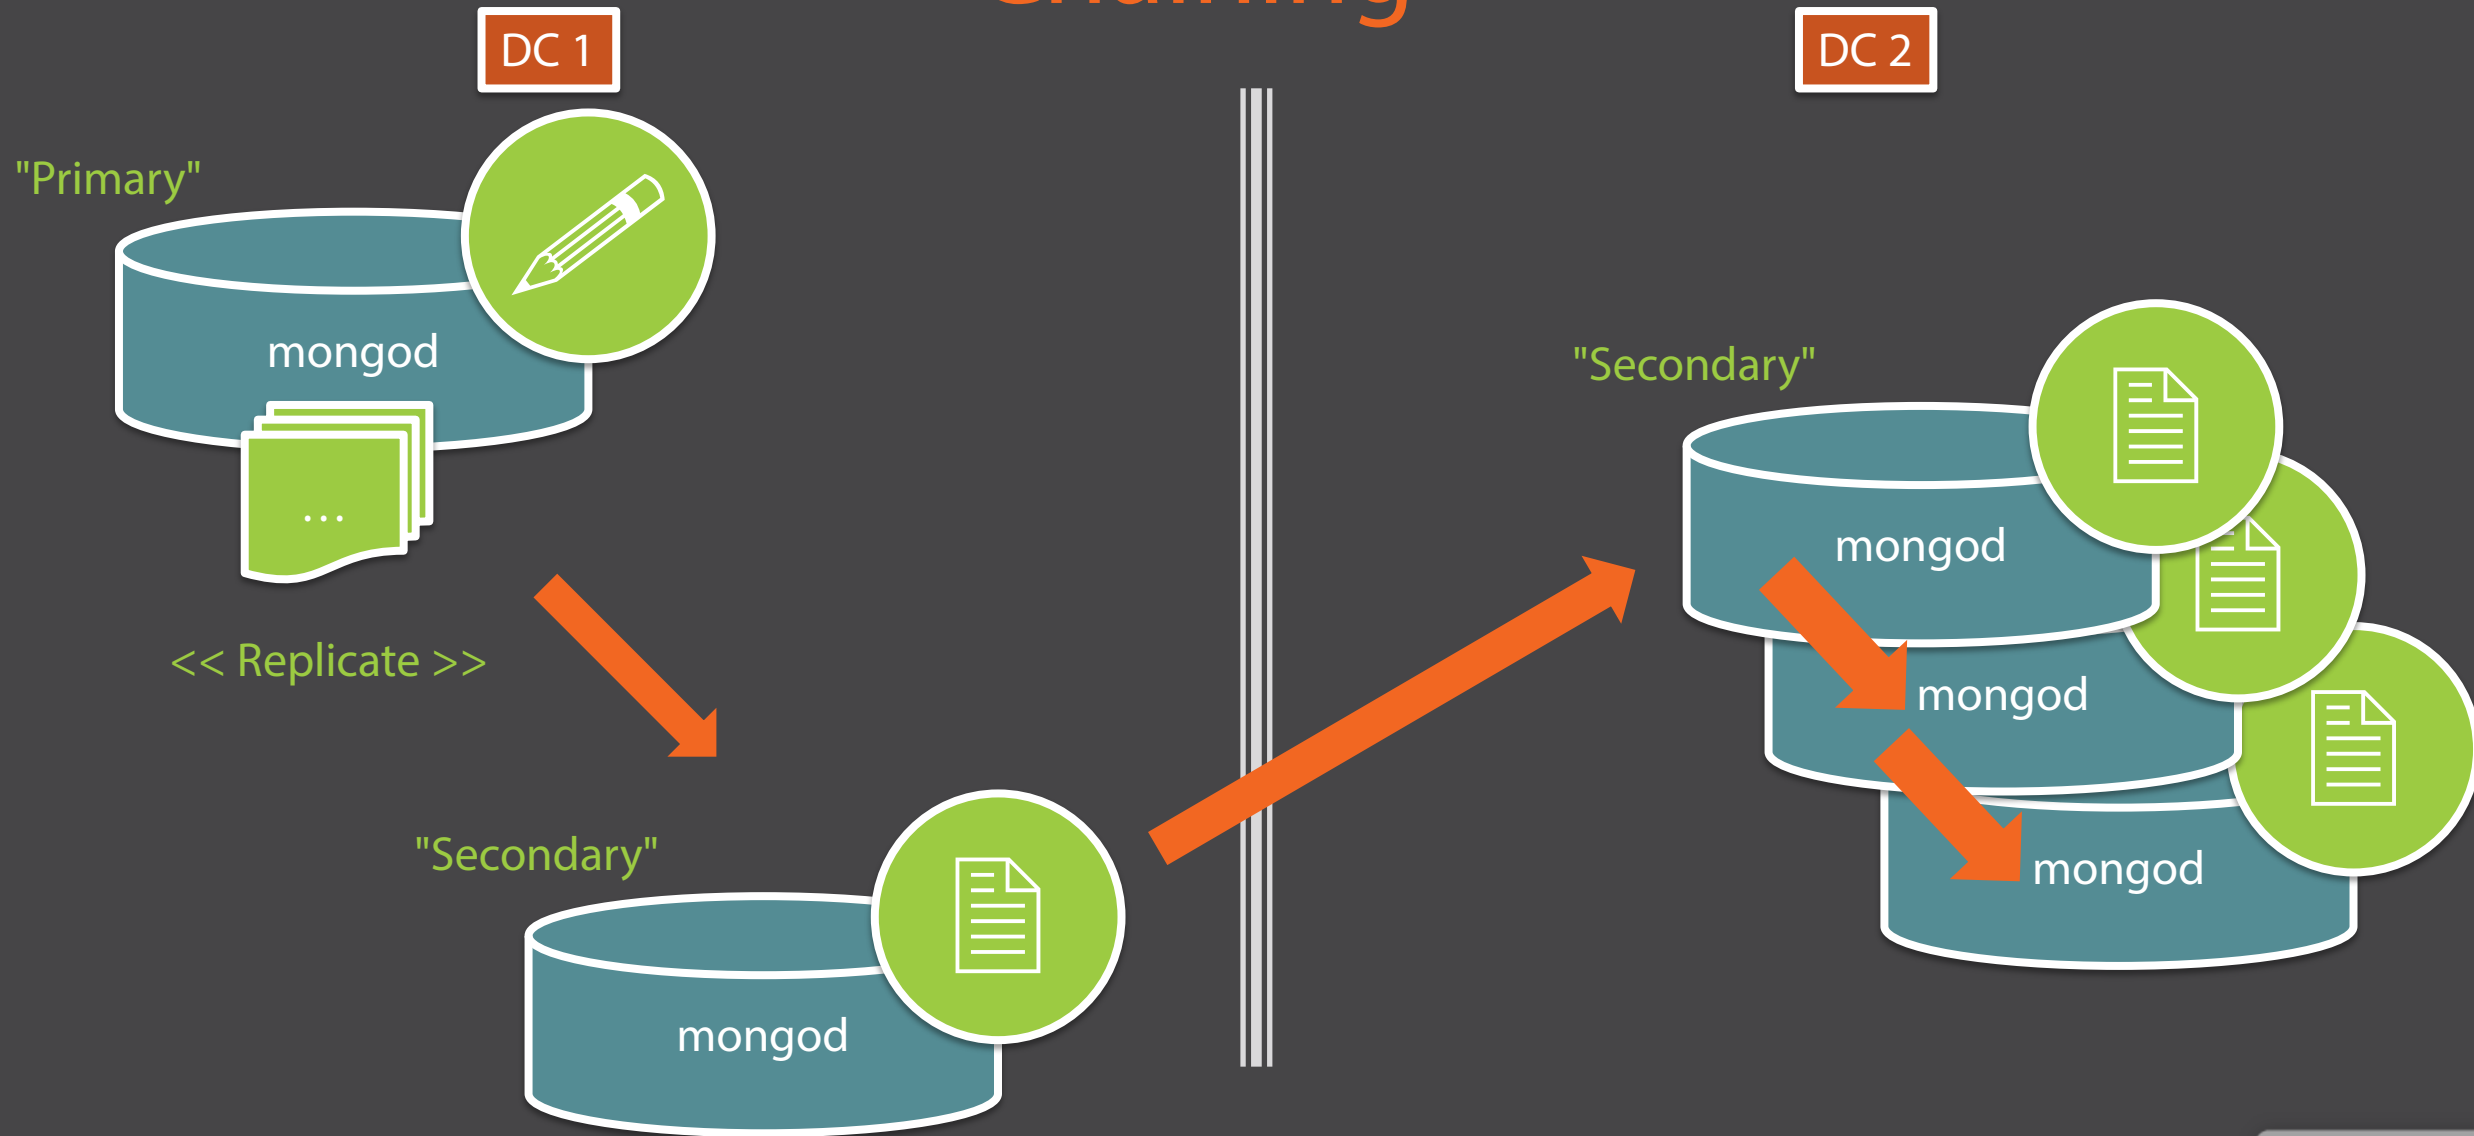
Network Partition



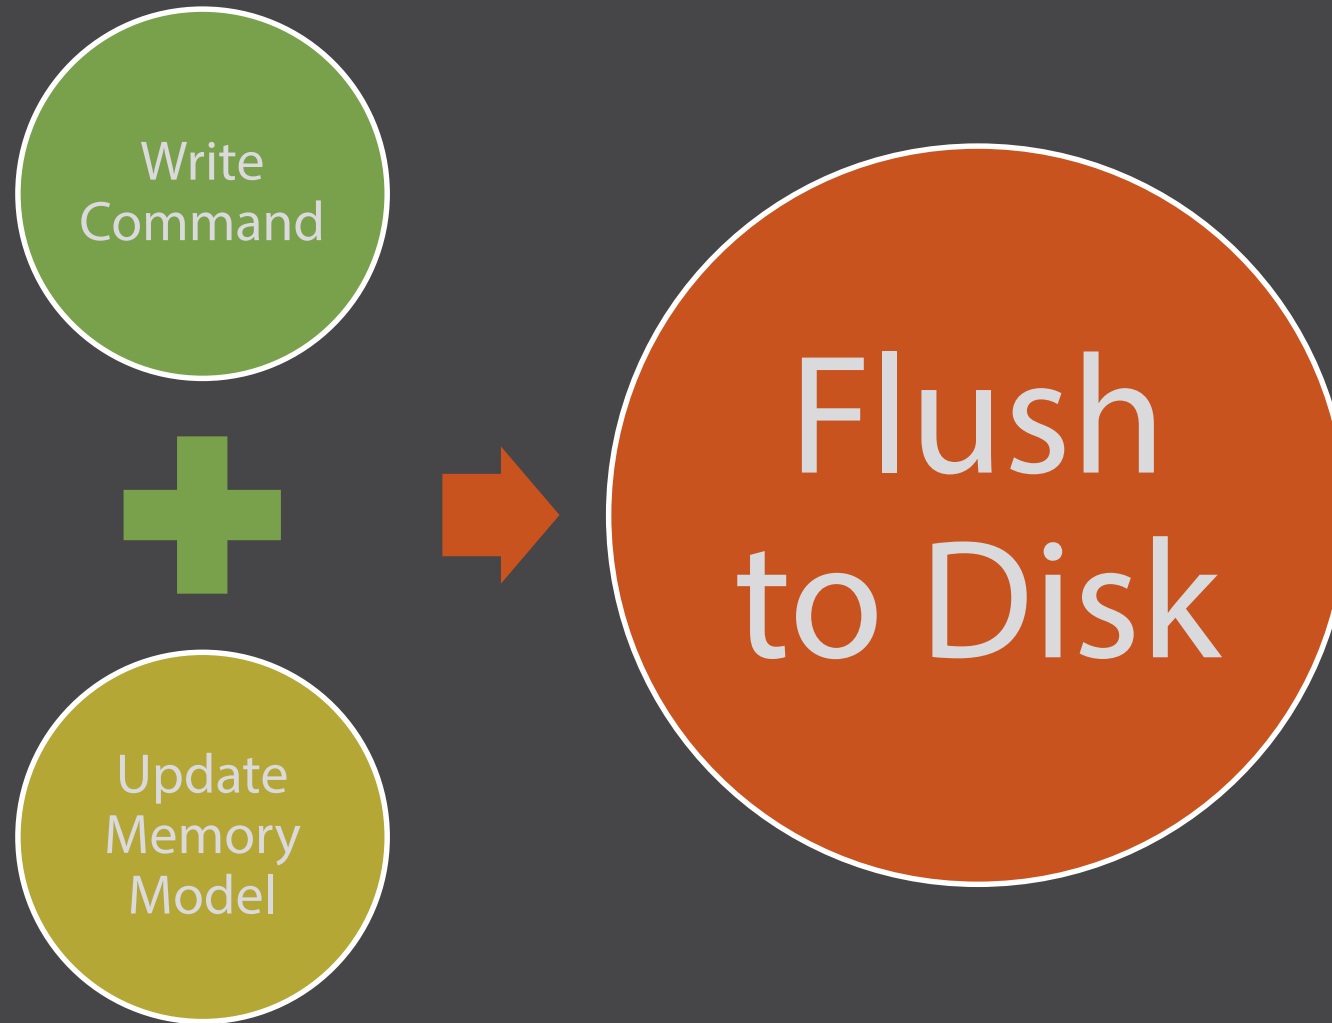
Chaining



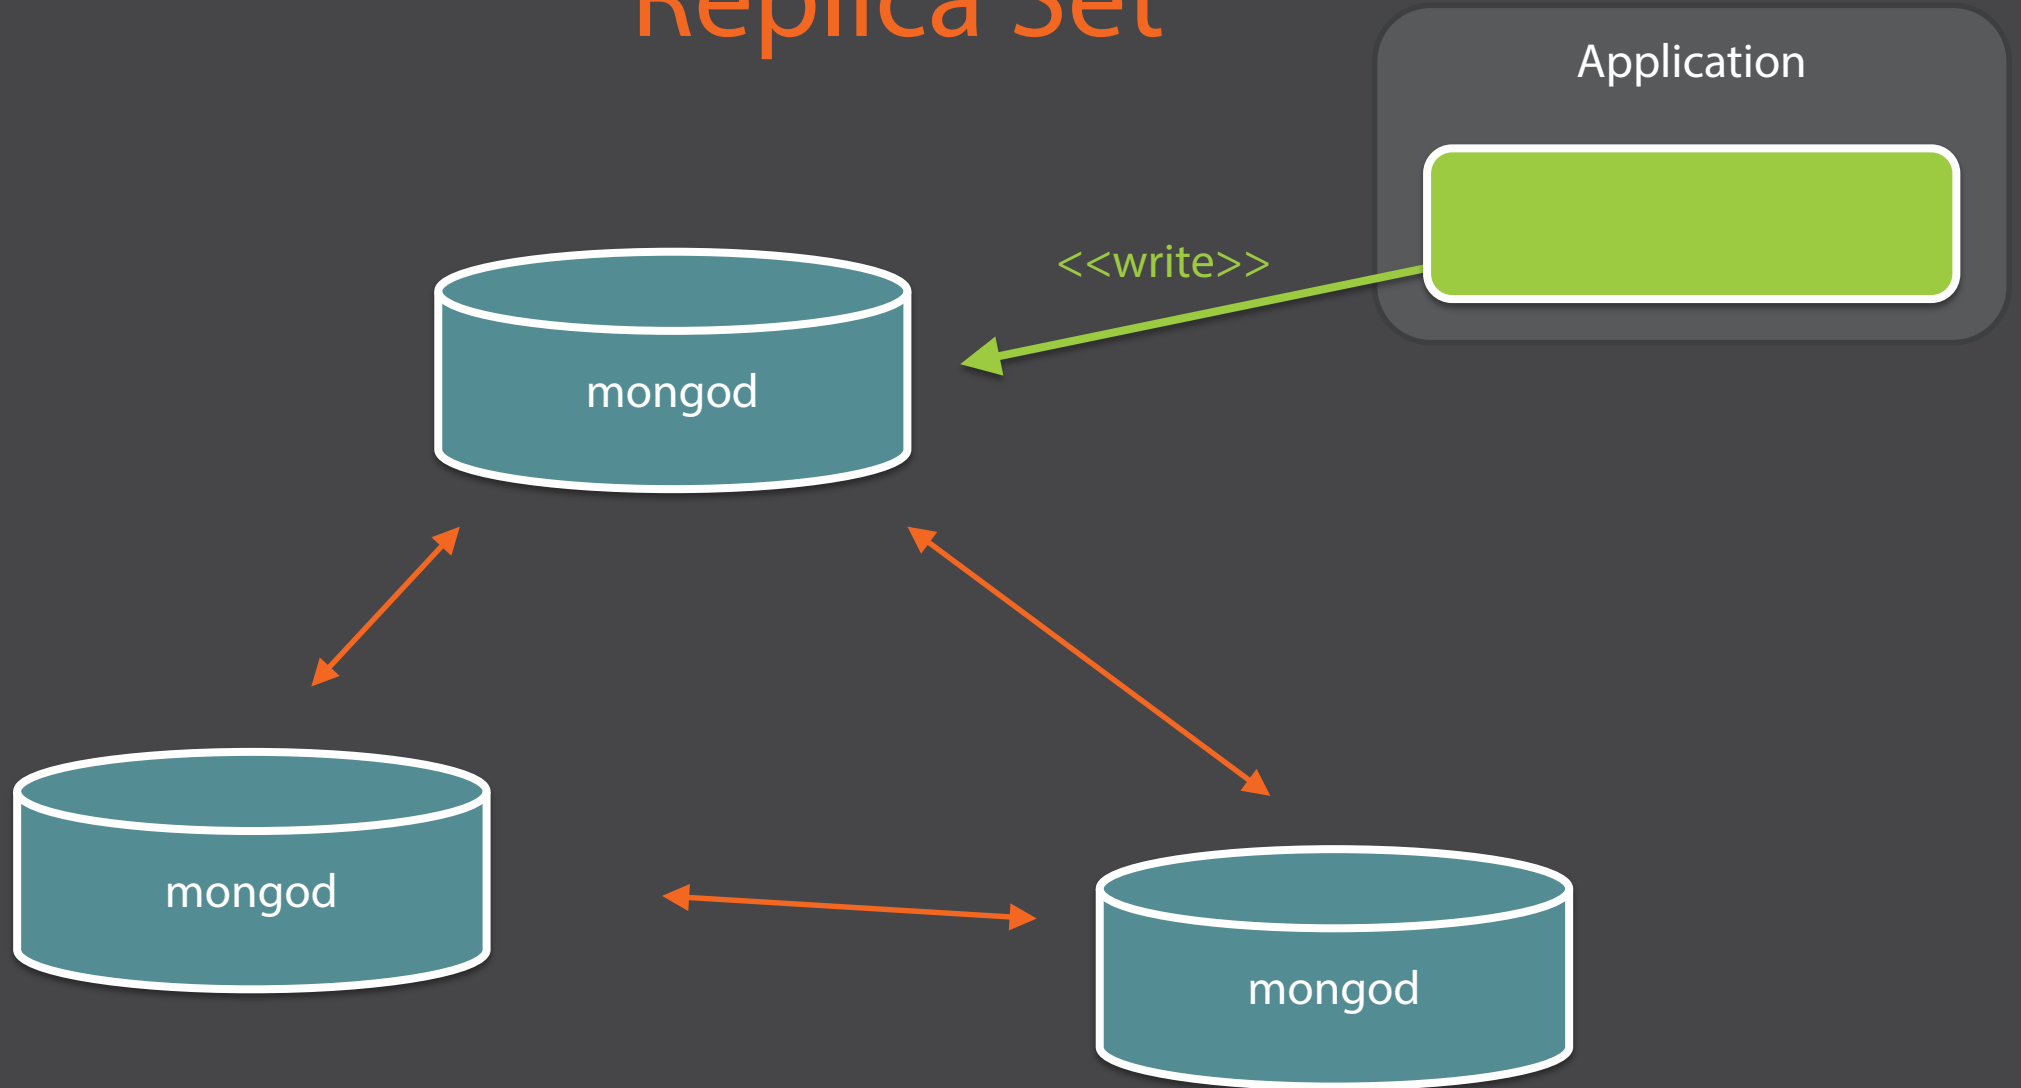
Chaining



Durability



Replica Set





"Acknowledged" write concern



"Unacknowledged" write concern




"Journaled" write concern



[w2|w3..|majority] write concerns

W1+ Write Concerns






"Majority of voting members"

majority write concern

W1+ Write Concerns



```
db.demo.insert(  
    {data...} ,  
    { w: [0..7|'majority'],  
      j: [true|false],  
      wtimeout:[milliseconds] } )
```

Write Concerns

Summary

