Sharding

Nuri Halperin

@nurih | www.plusnconsulting.com

Not Sharded

Application

Sharded

Application

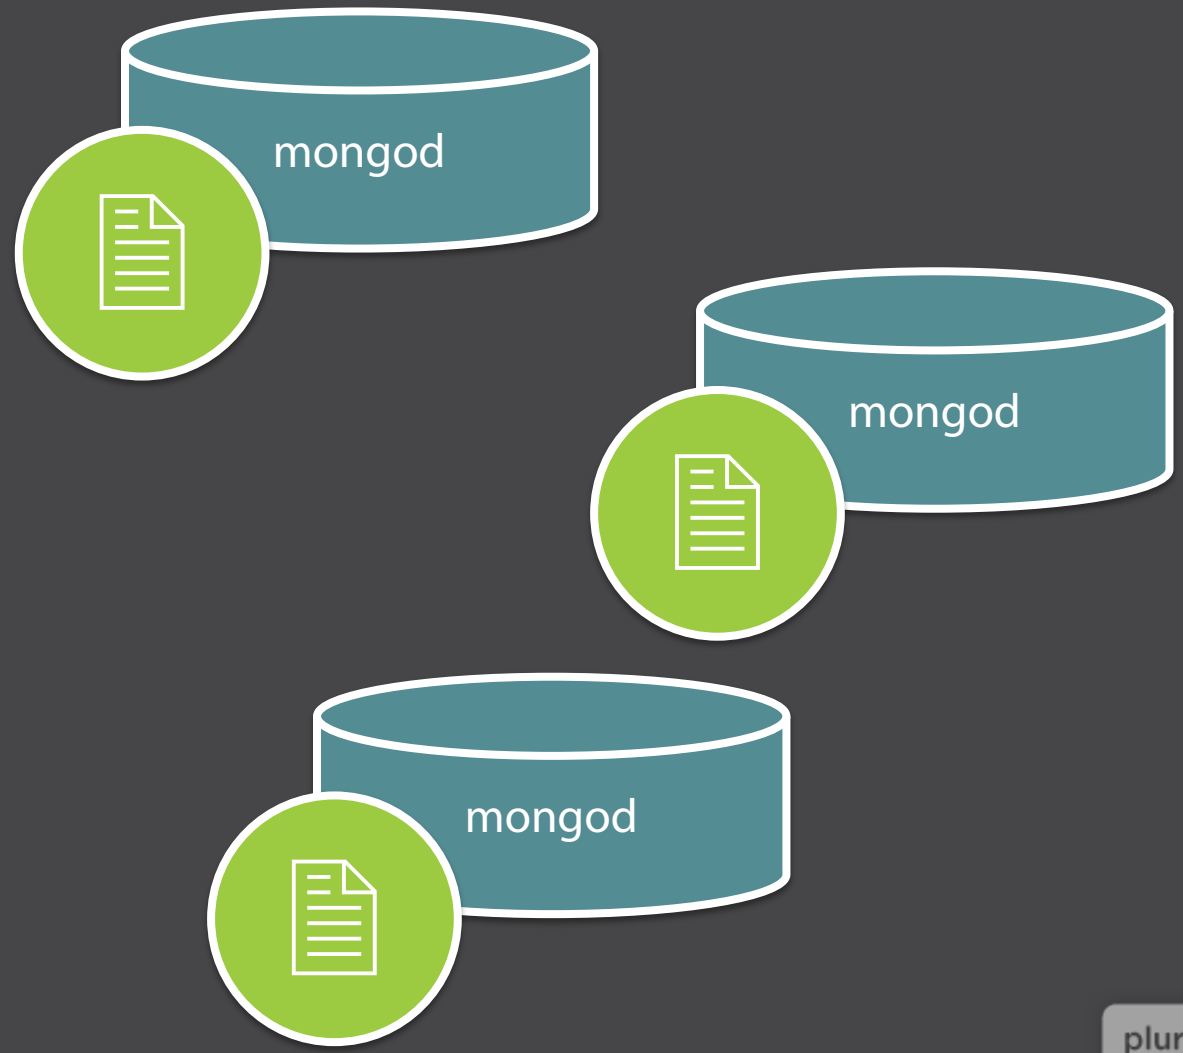

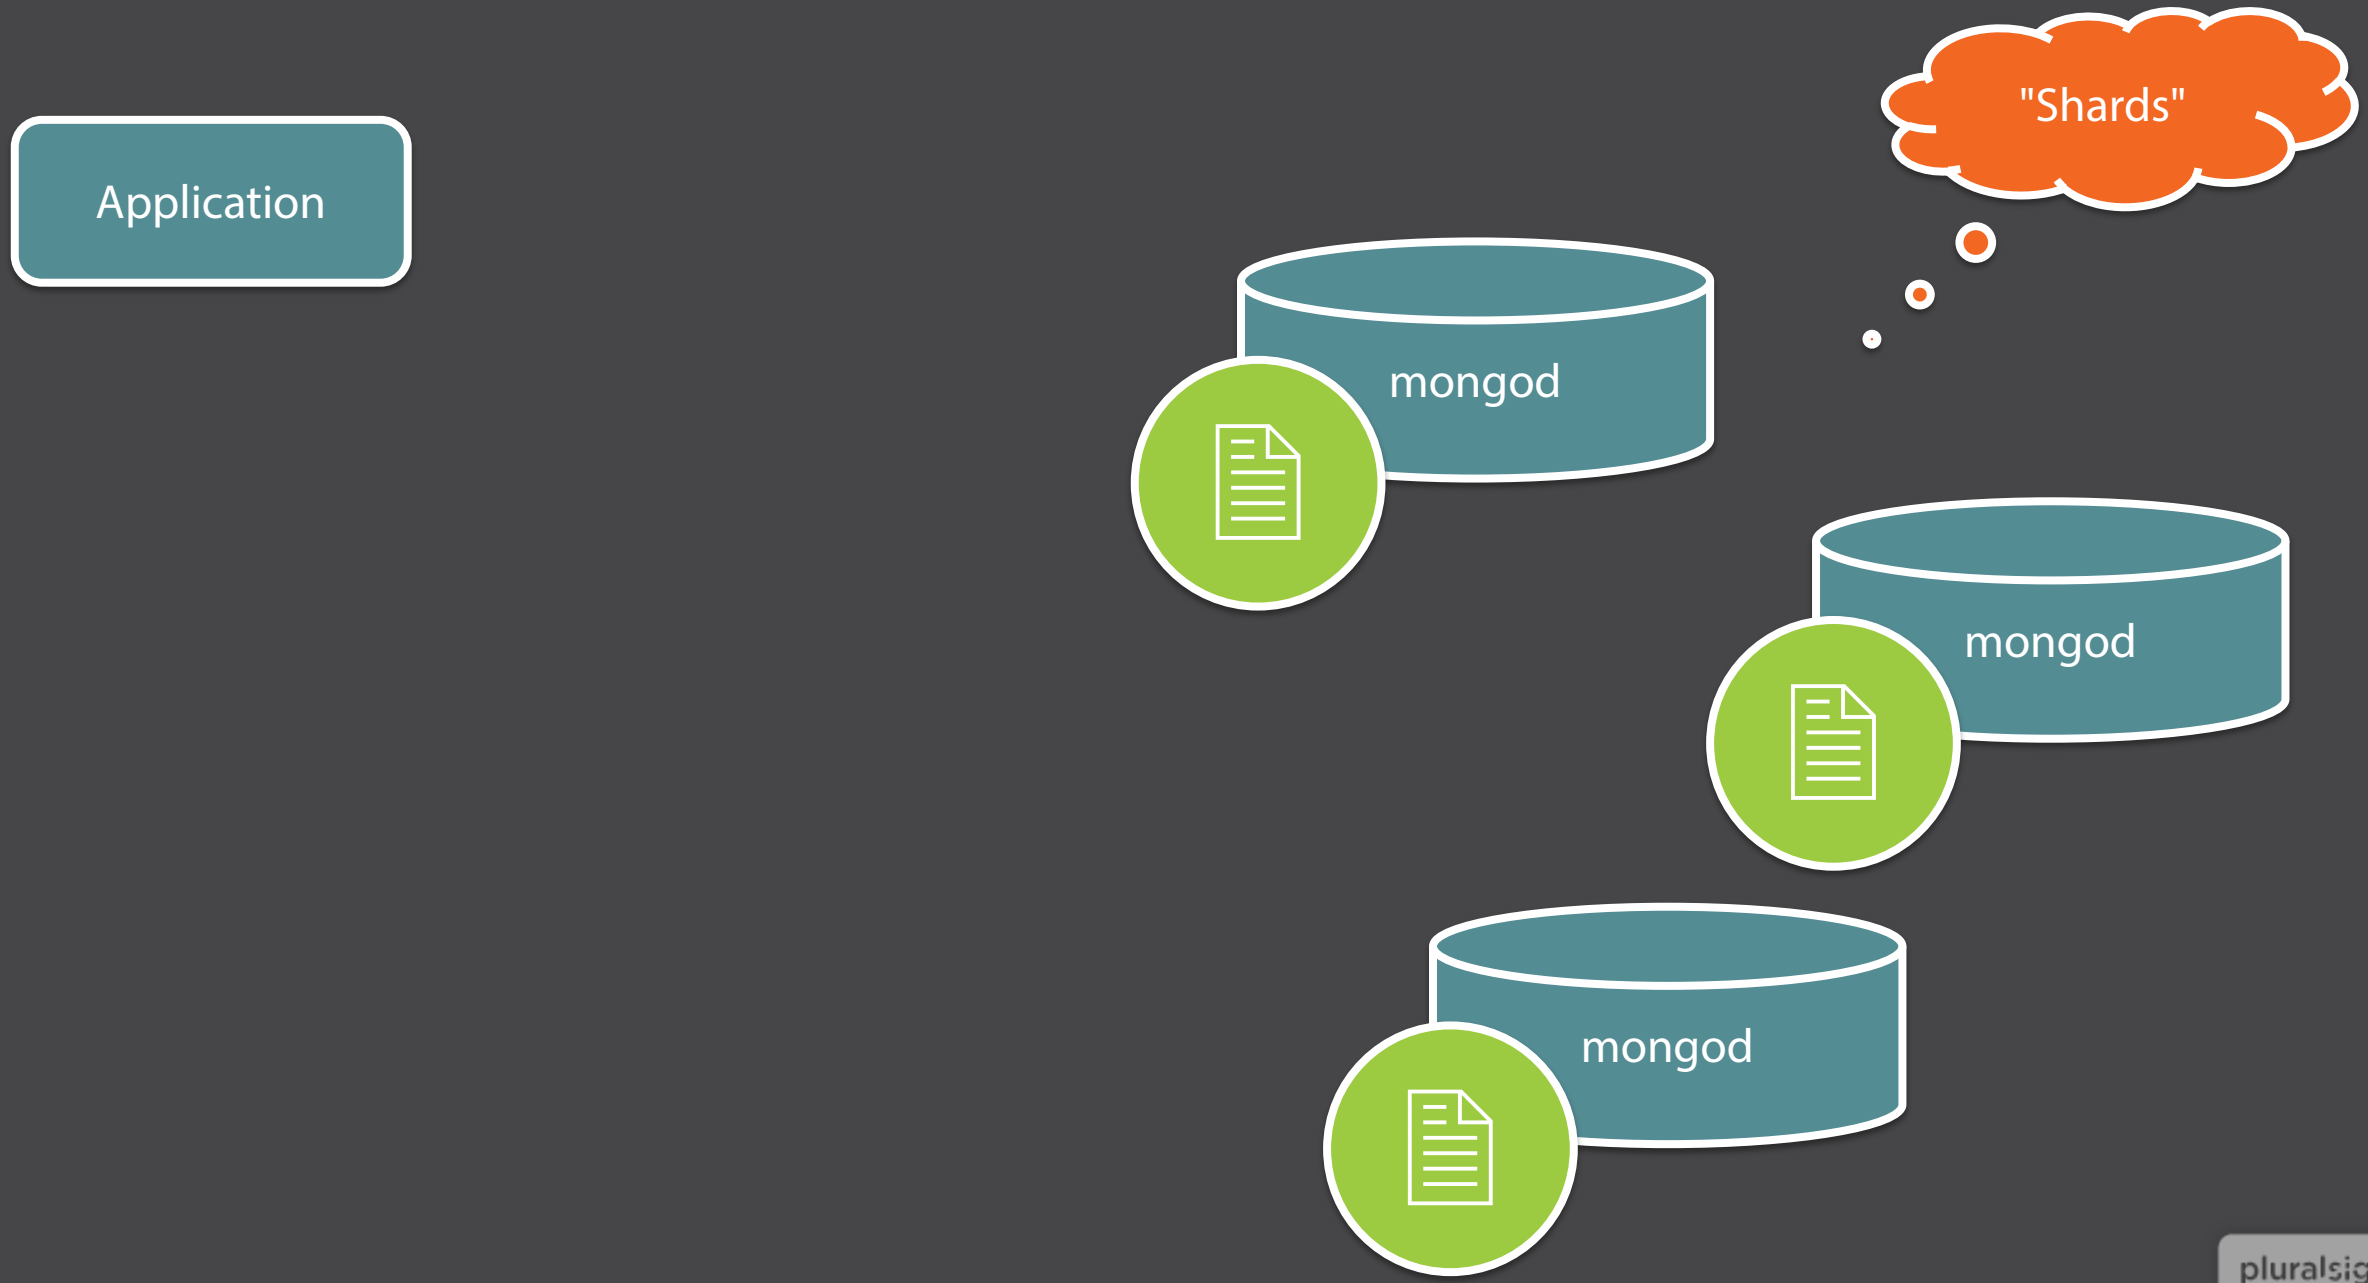
Application
"Shards"
mongod
mongod
mongod

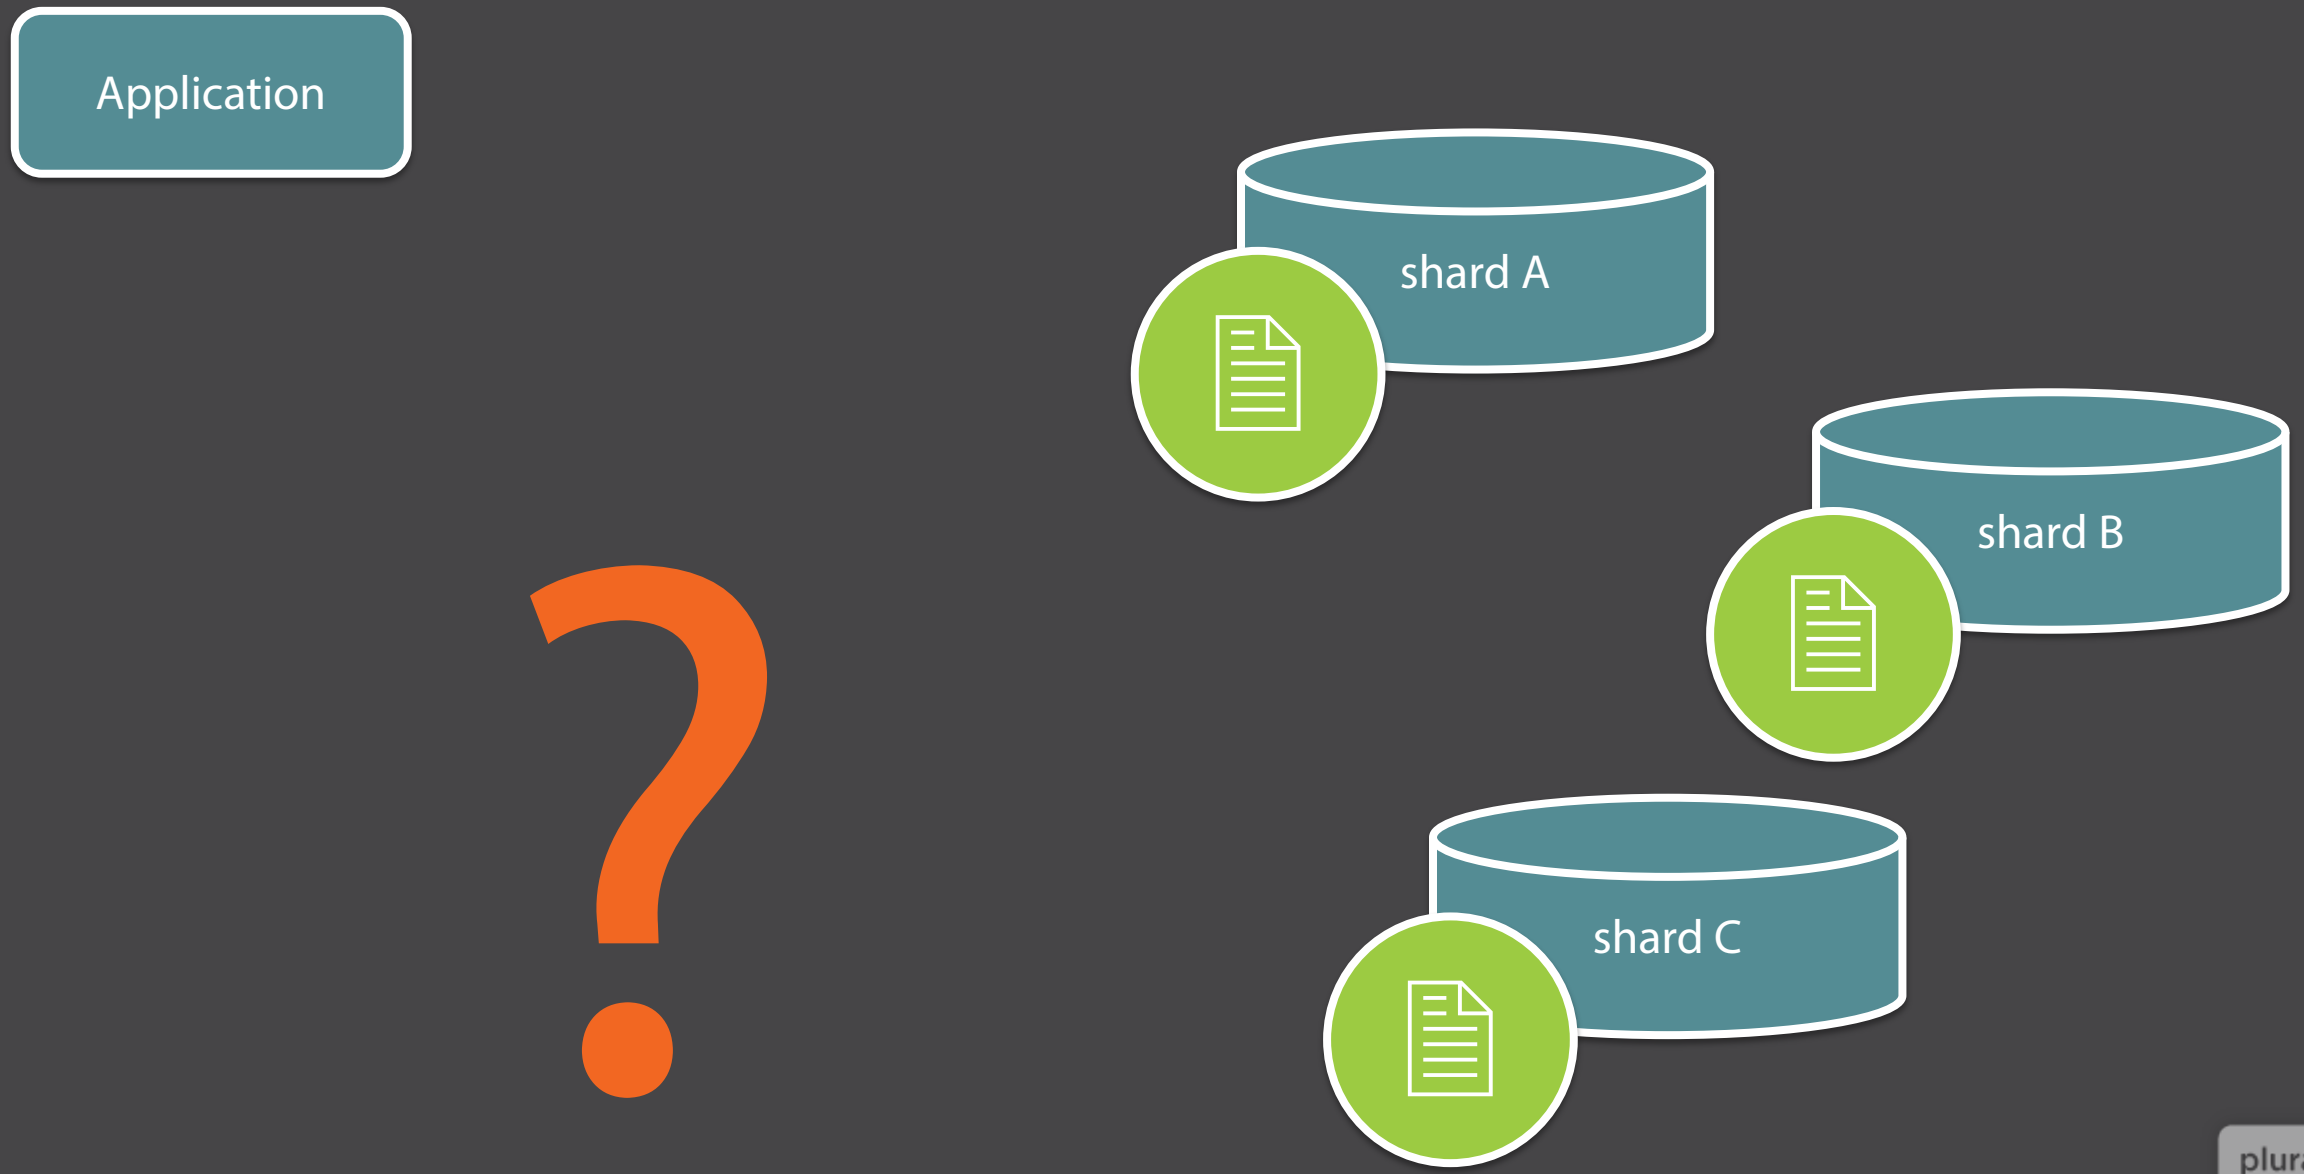
Application
shard A
shard B
shard C

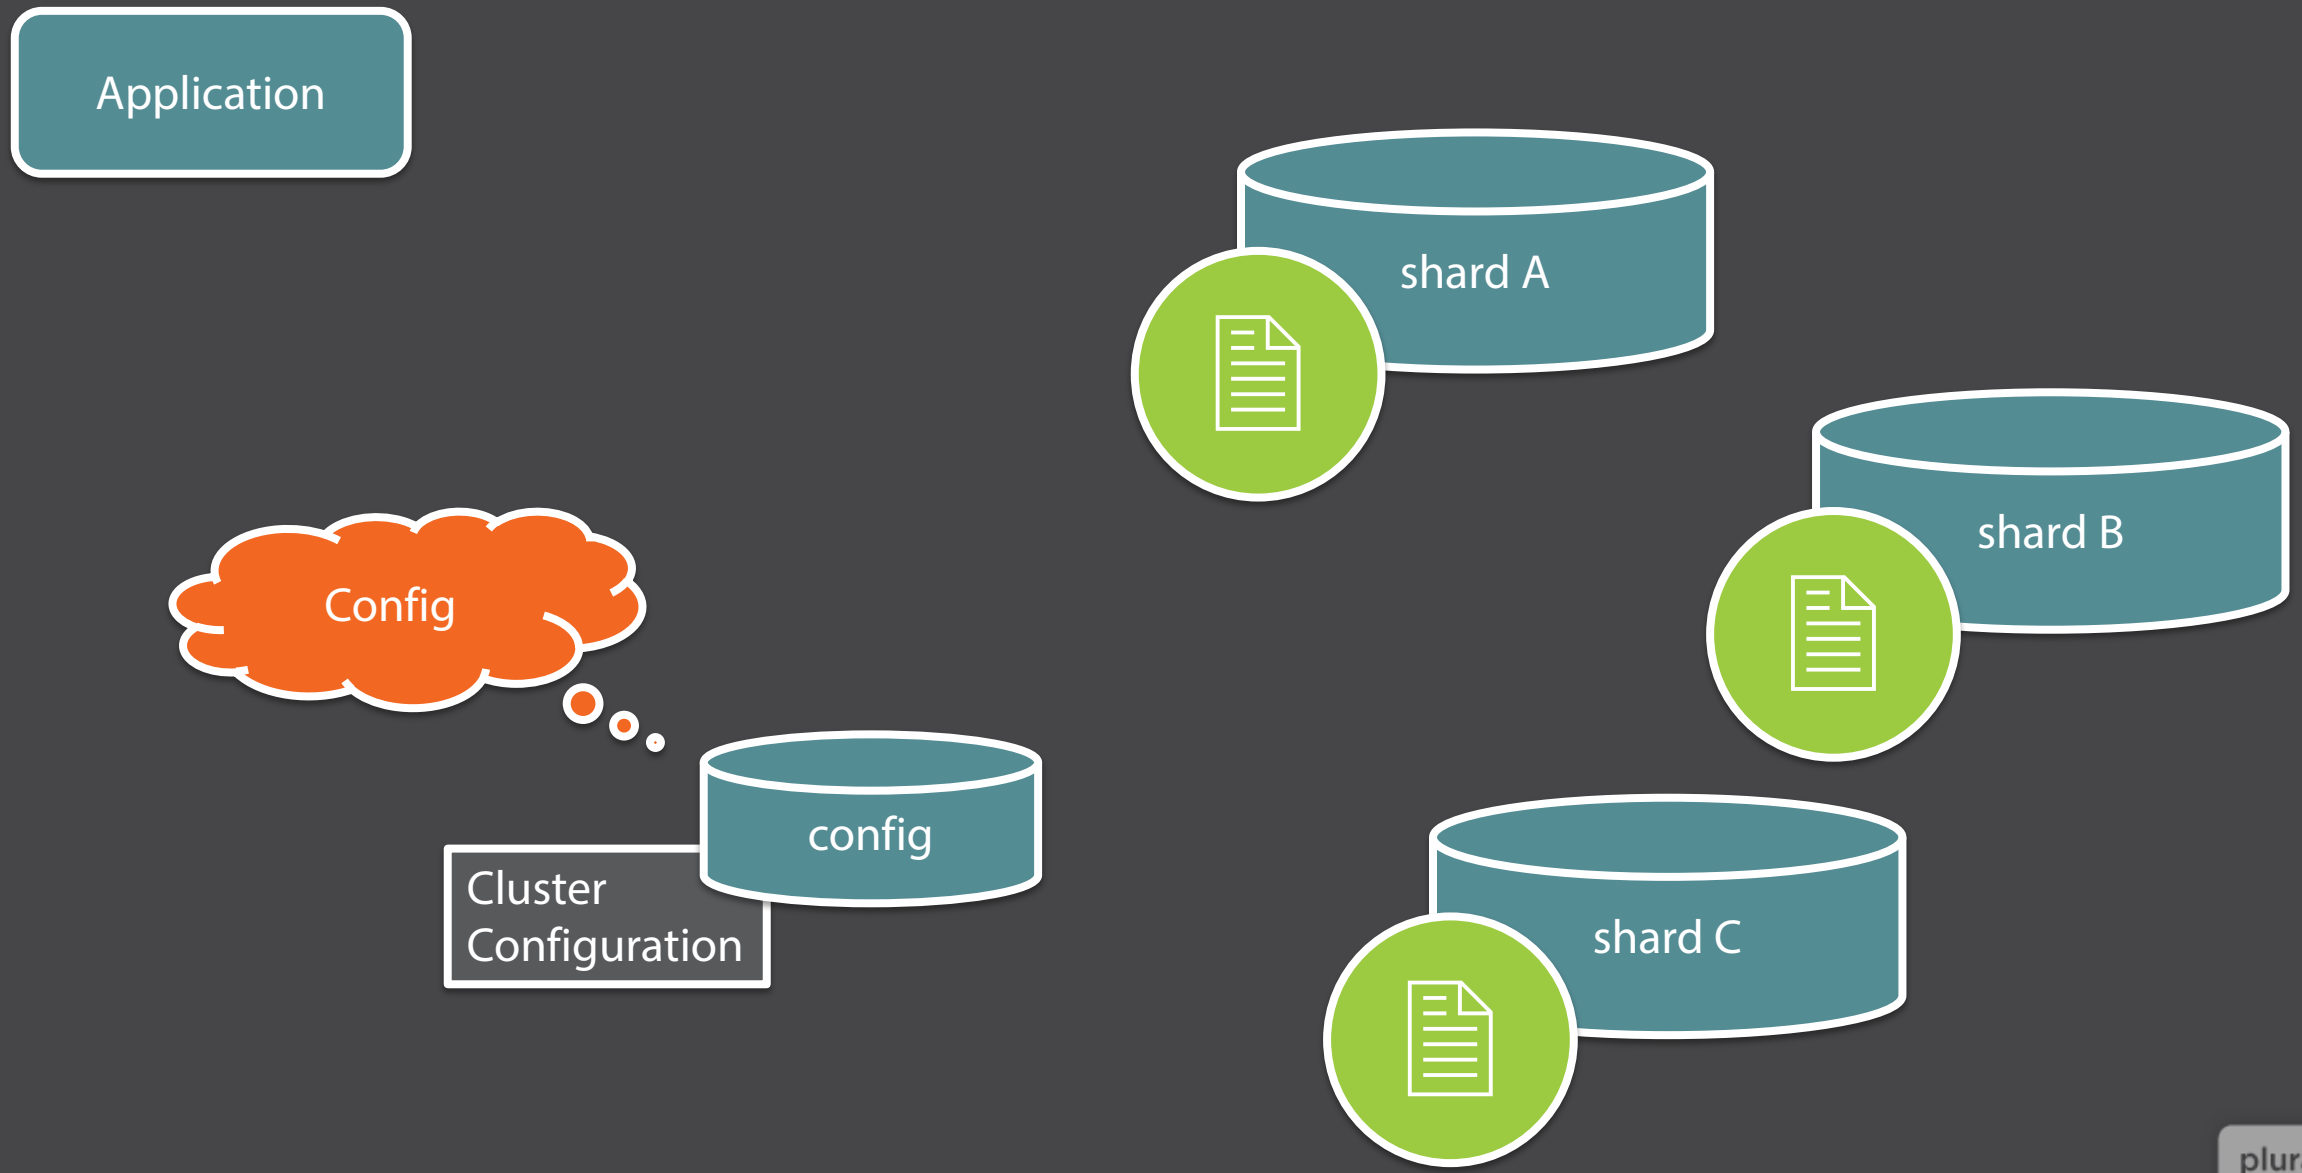
Application
shard A
shard B
shard C
Config
config
Cluster
Configuration

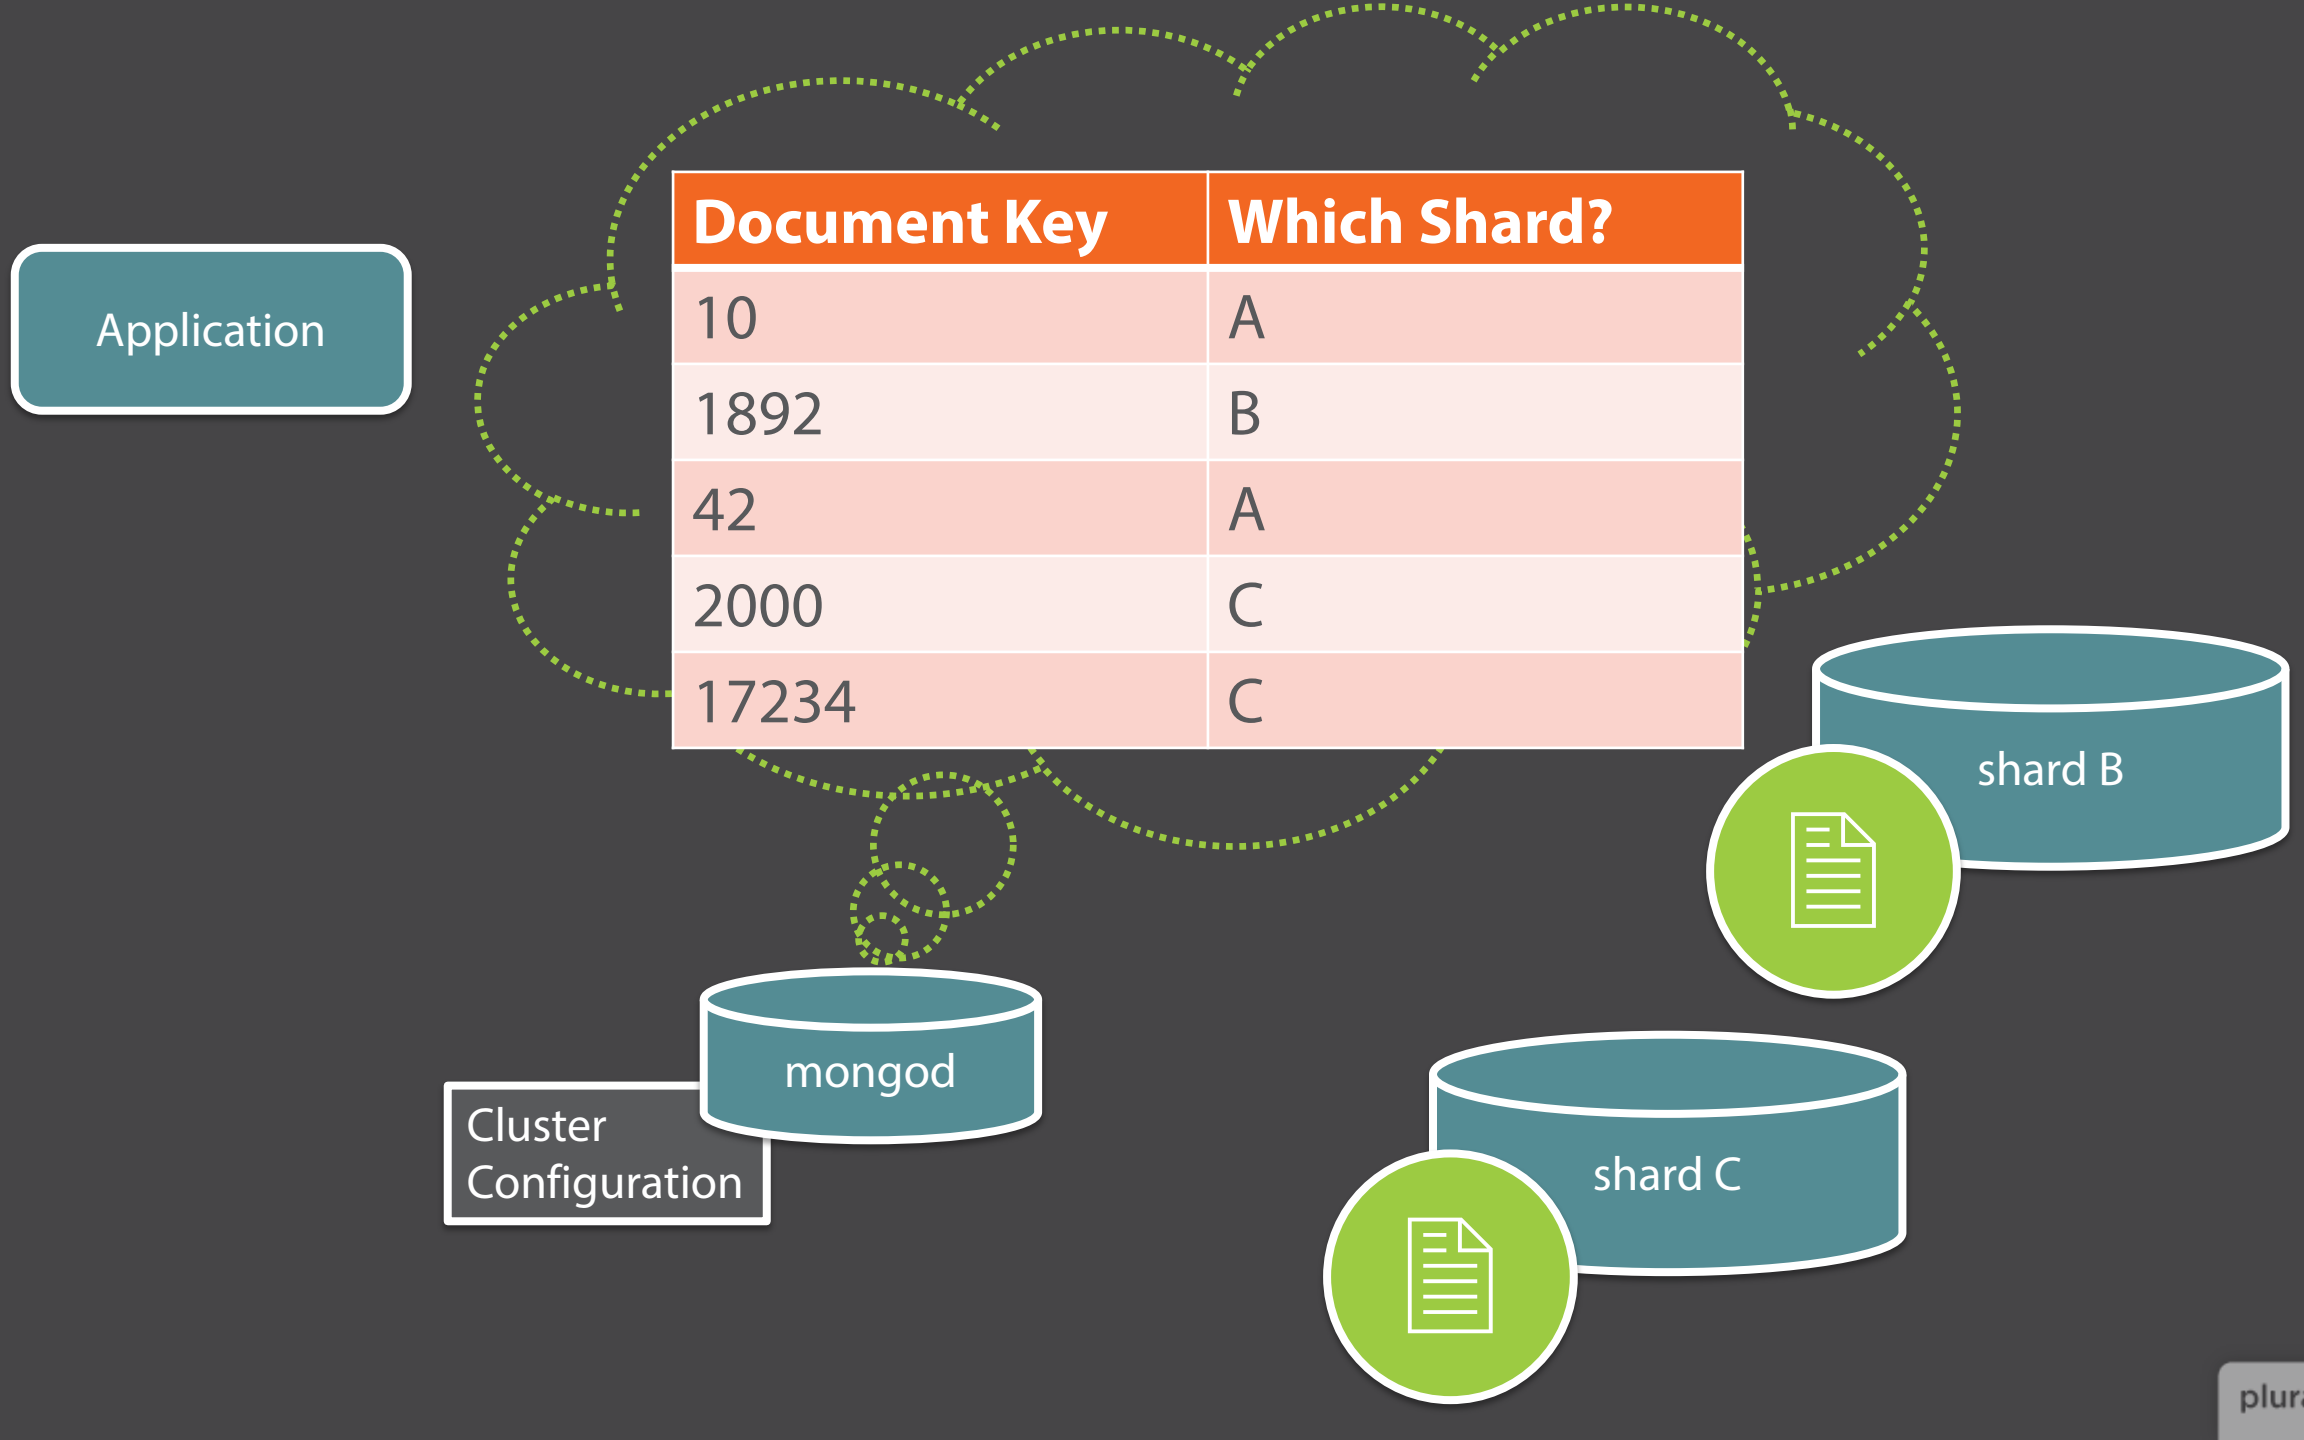
Application
Document Key
Which Shard?
10
A
1892
B
42
A
2000
C
17234
C
shard B
mongod
Cluster Configuration
shard C

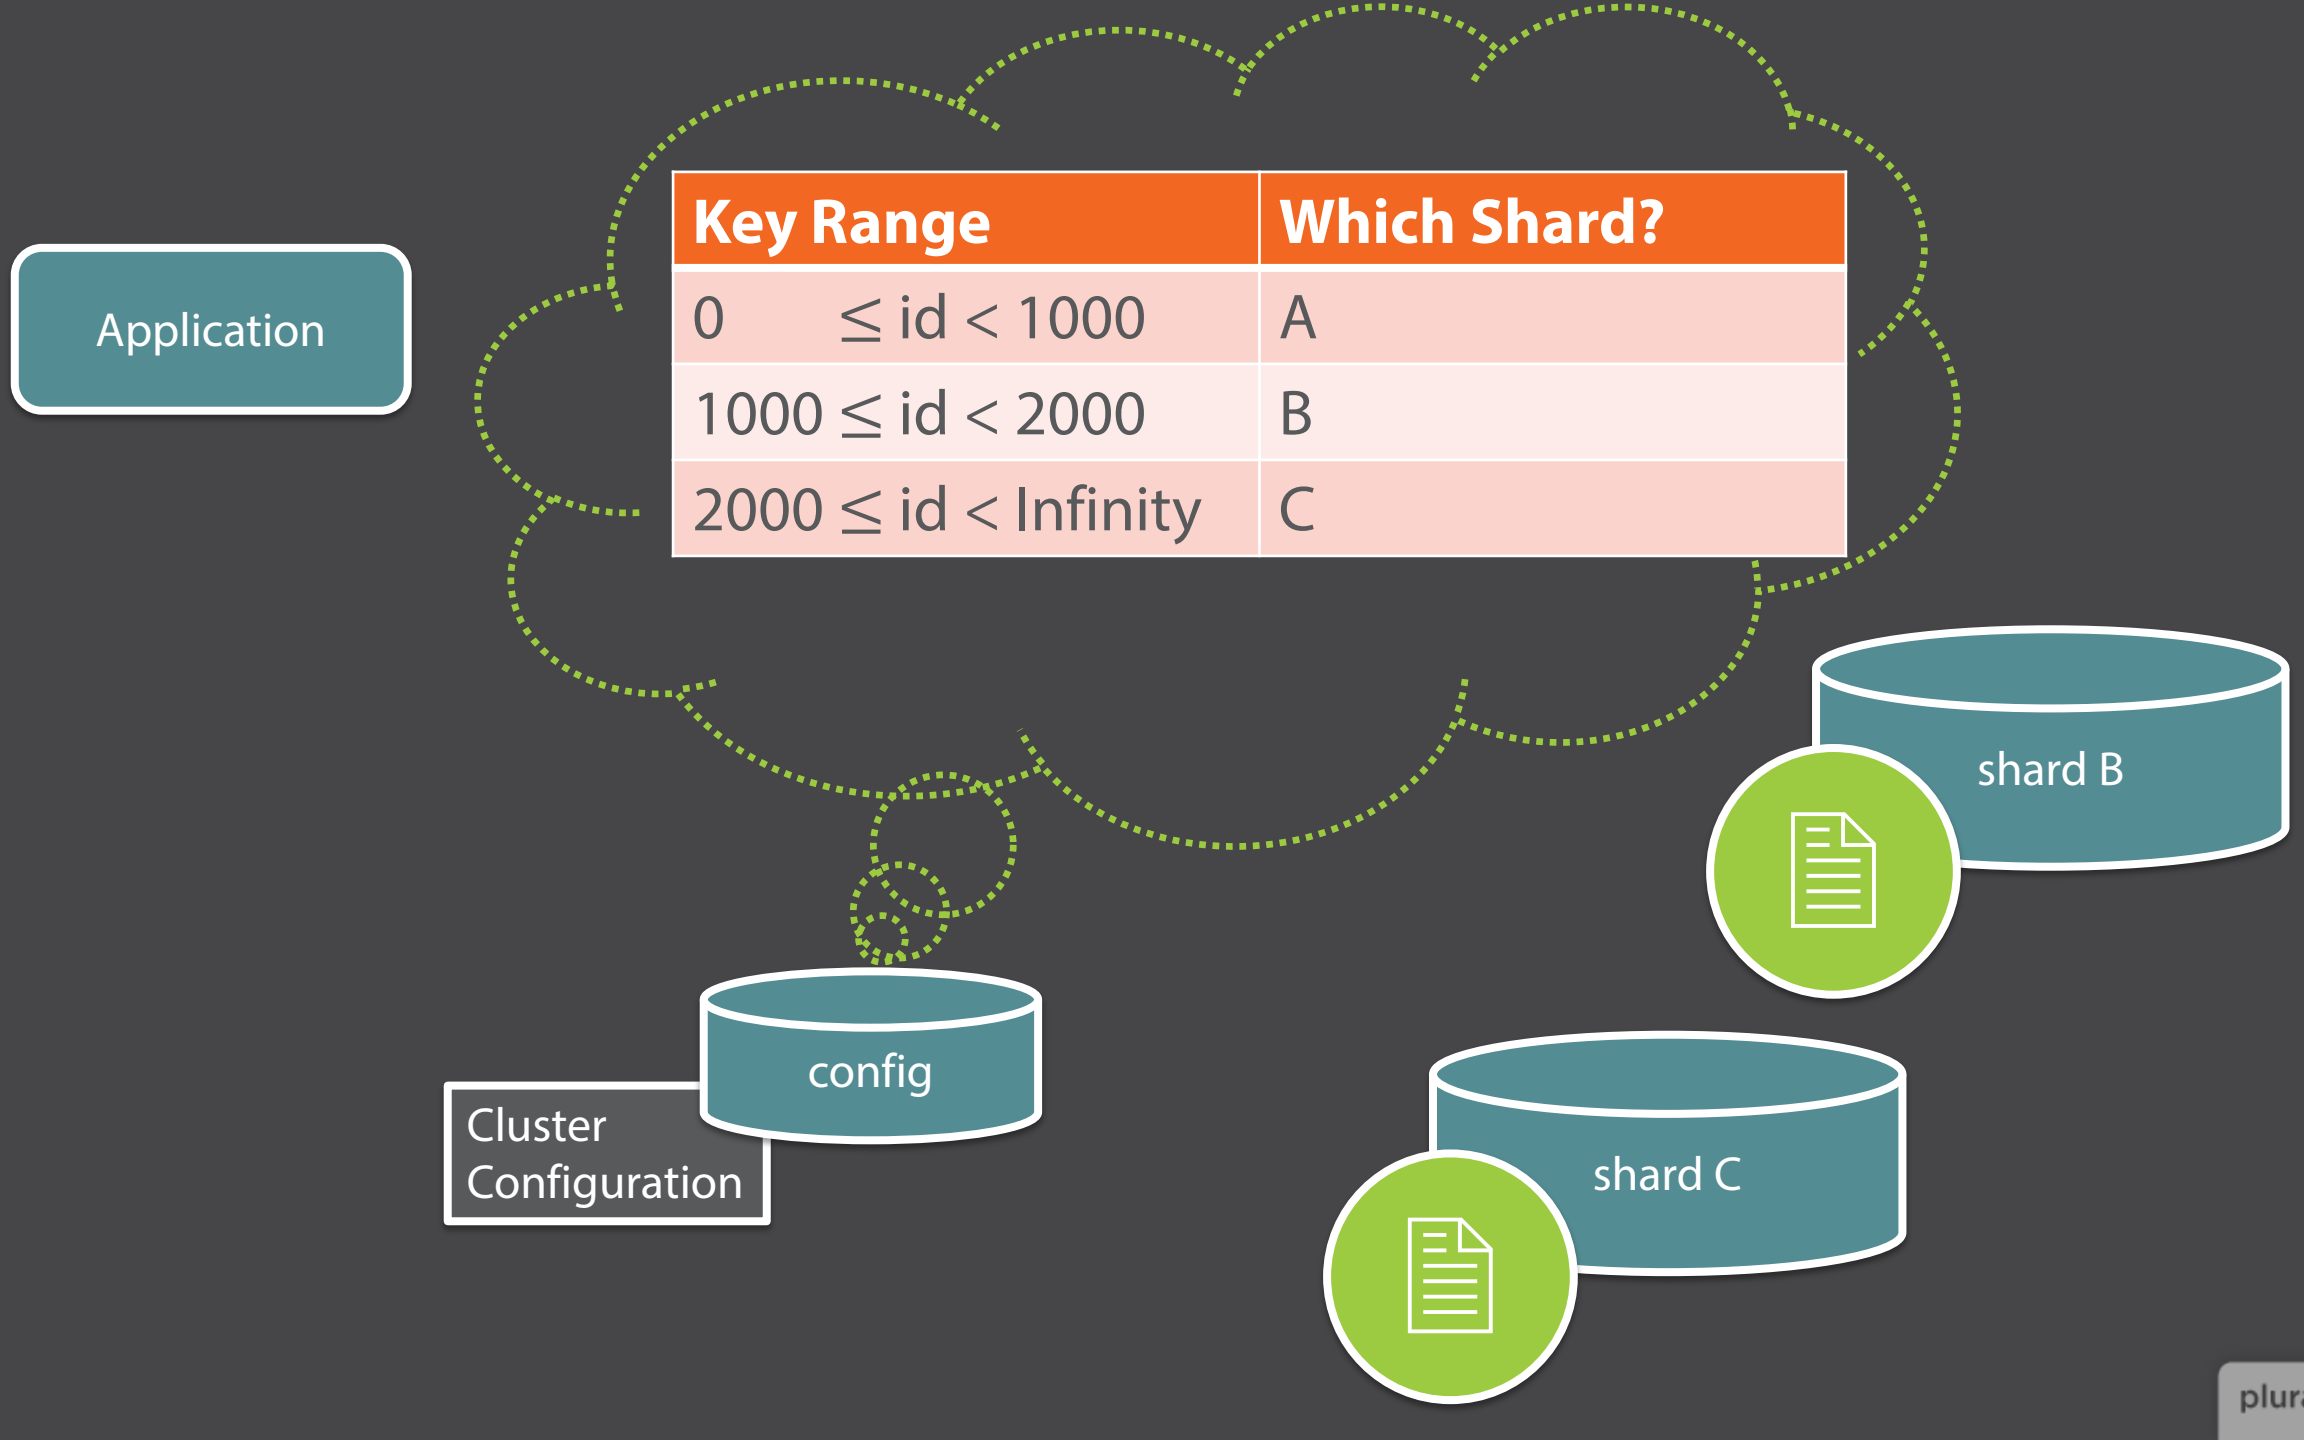
Application
Key Range
Which Shard?
0 ≤ id < 1000
A
1000 ≤ id < 2000
B
2000 ≤ id < Infinity
C
shard B
config
Cluster Configuration
shard C

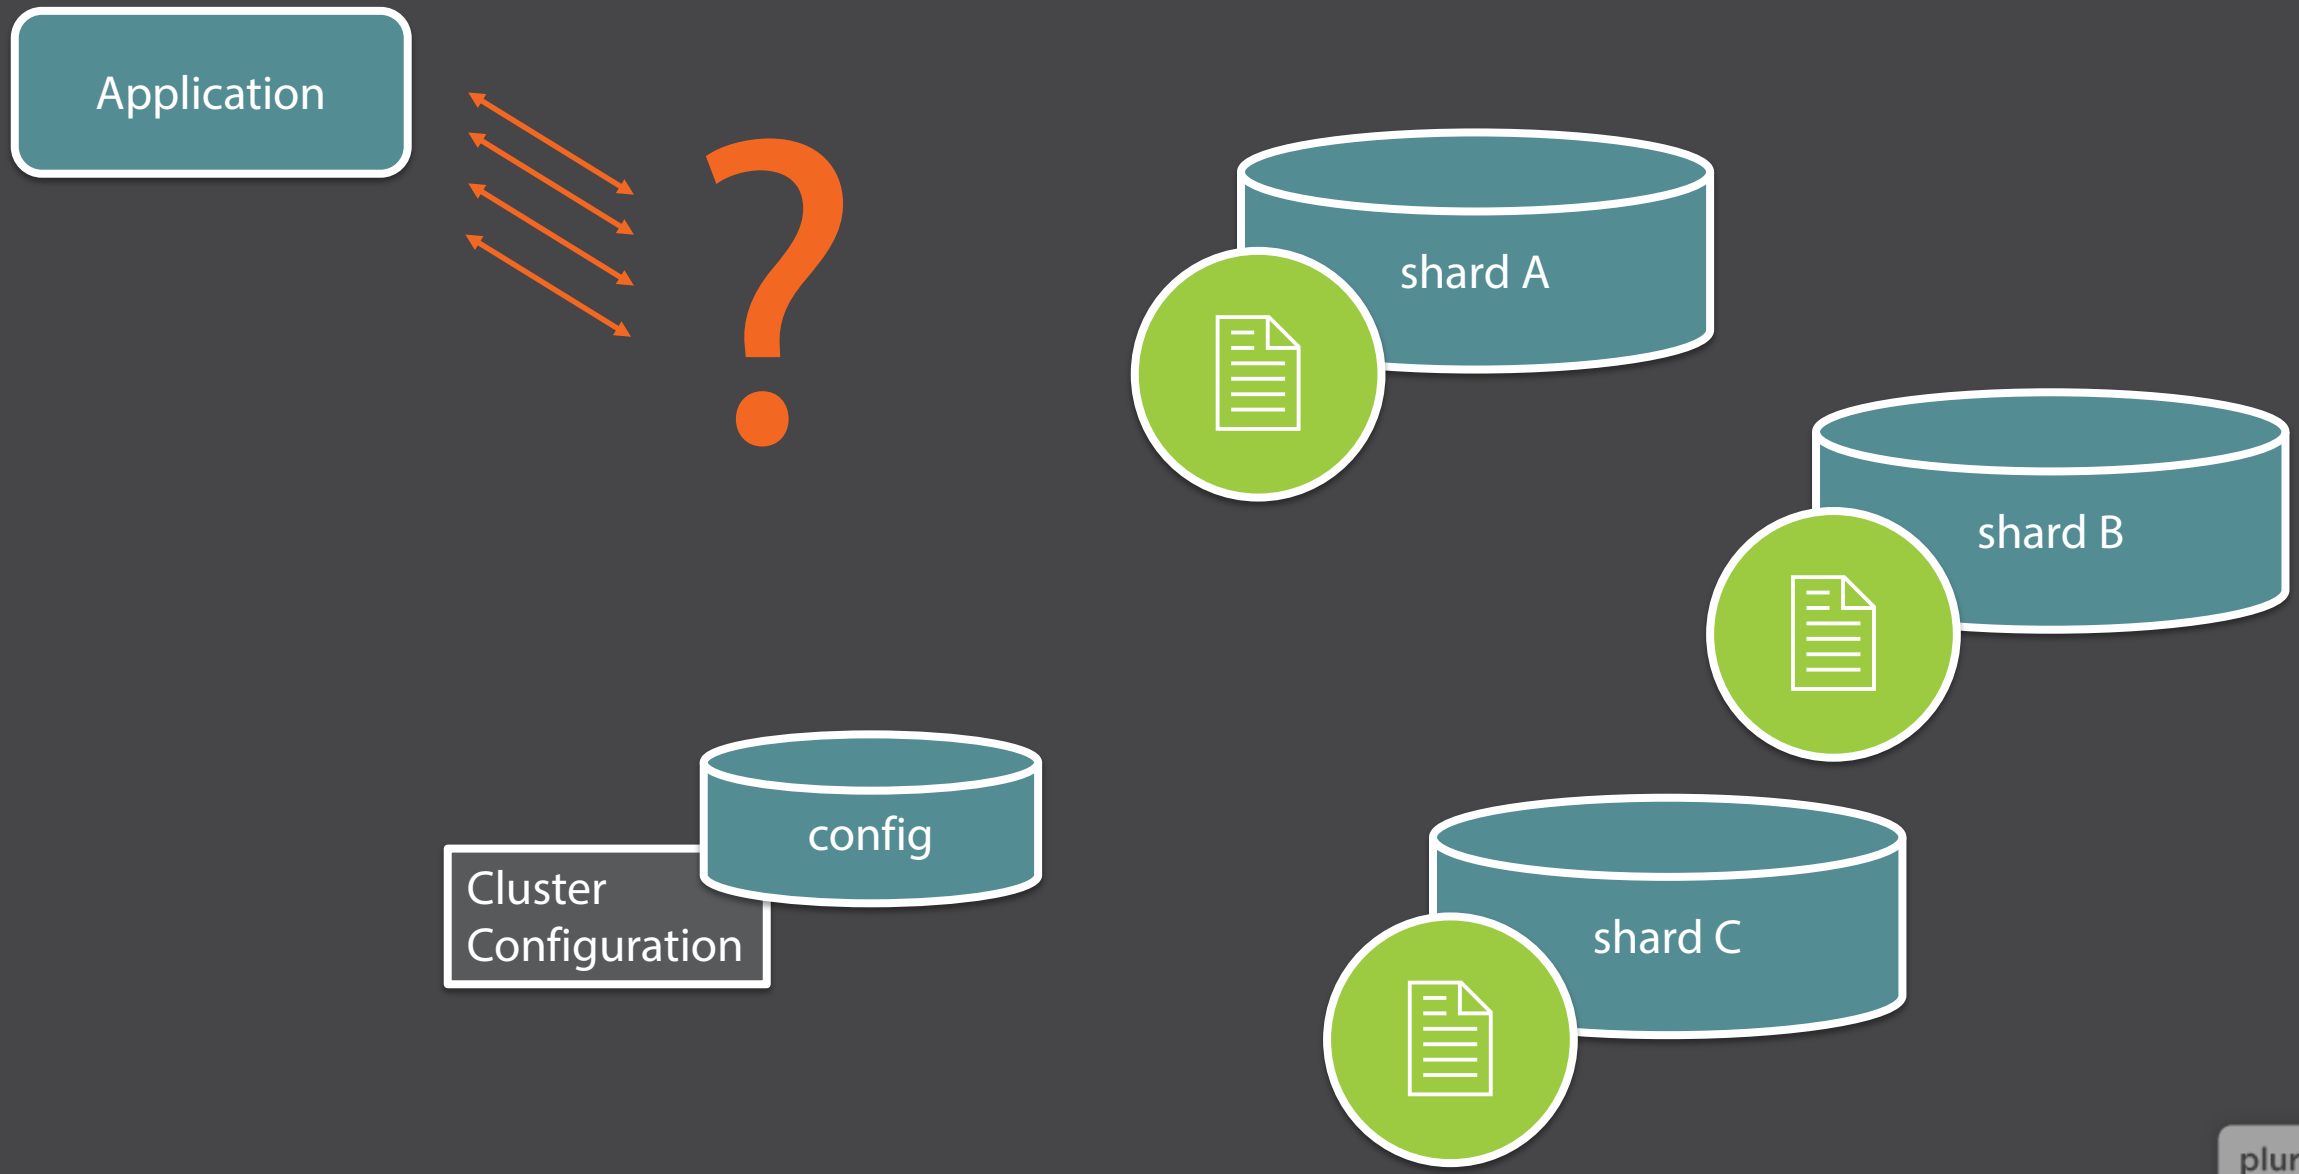
Application
?
shard A
shard B
shard C
config
Cluster
Configuration

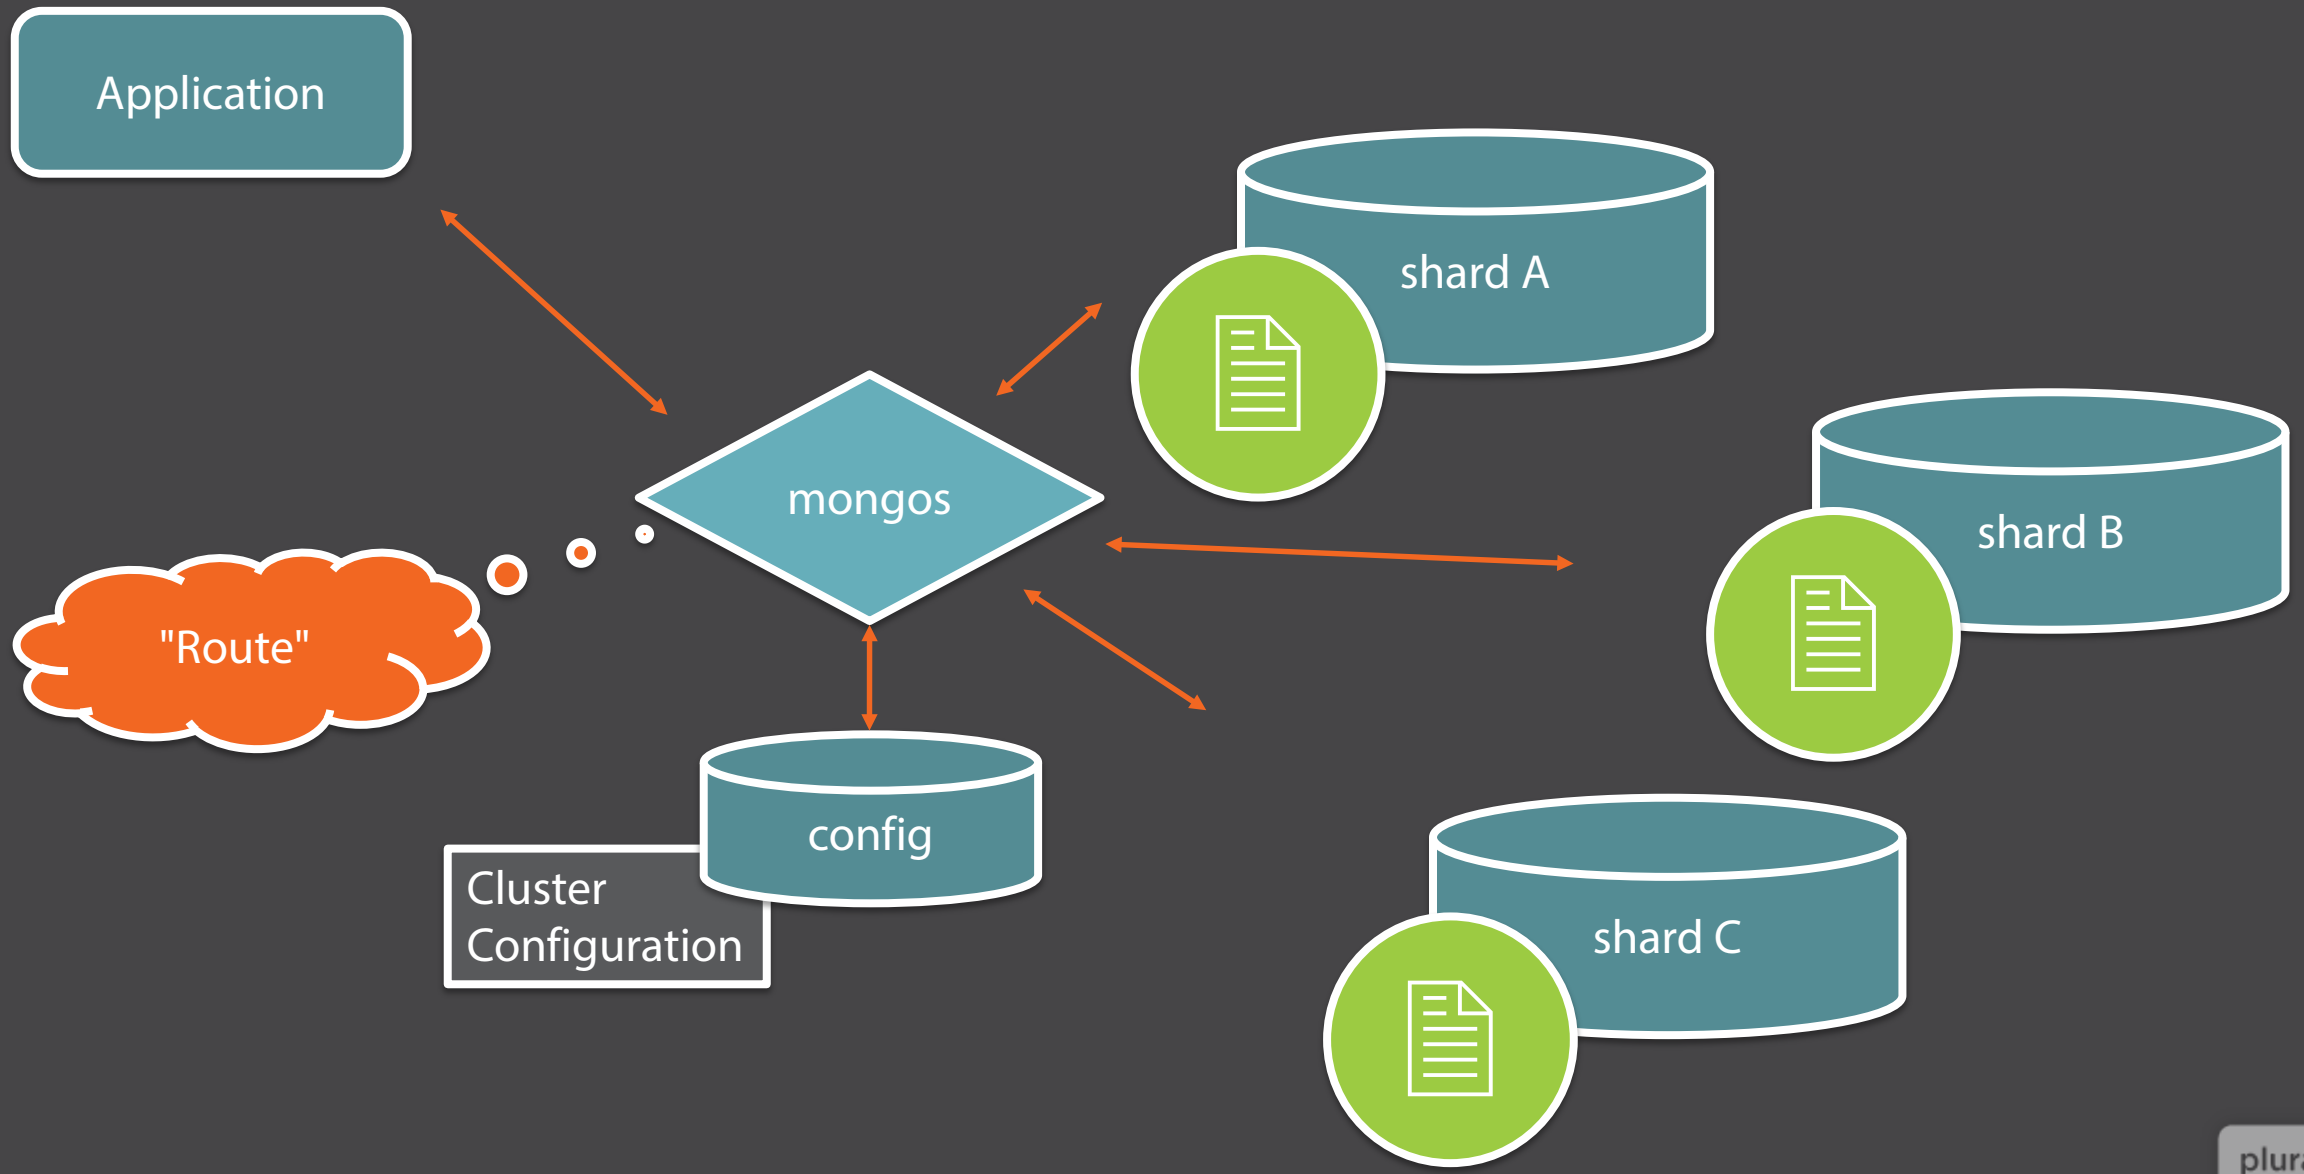

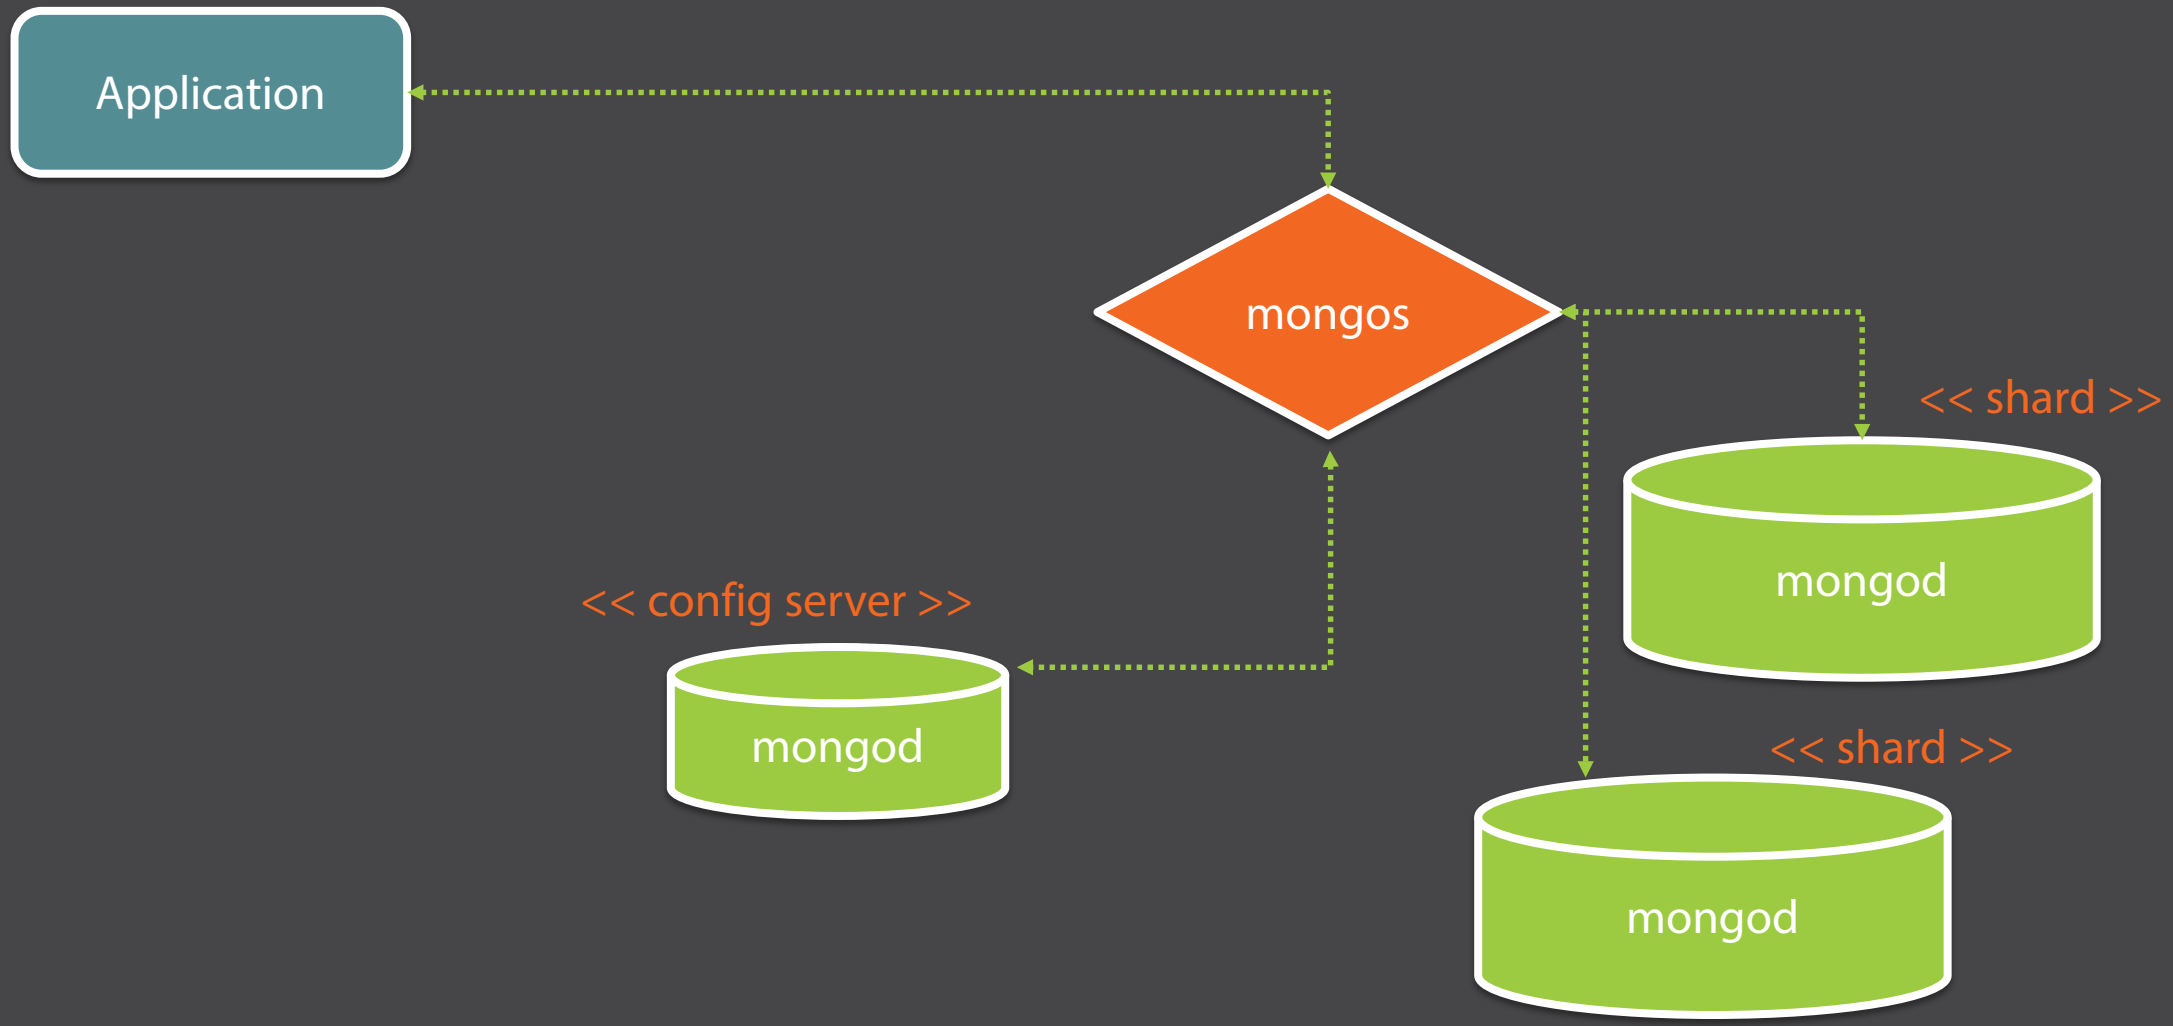

At a Glance

```
mongos --configdb configServer

// inside the shell:
mongos> sh.addShard('mongodServer:port')

mongos> sh.enableSharding('dbName')

mongos> sh.shardCollection('dbName.coll', {'field':1})
```

“

Production Gig?

Three Config!

- Me ”

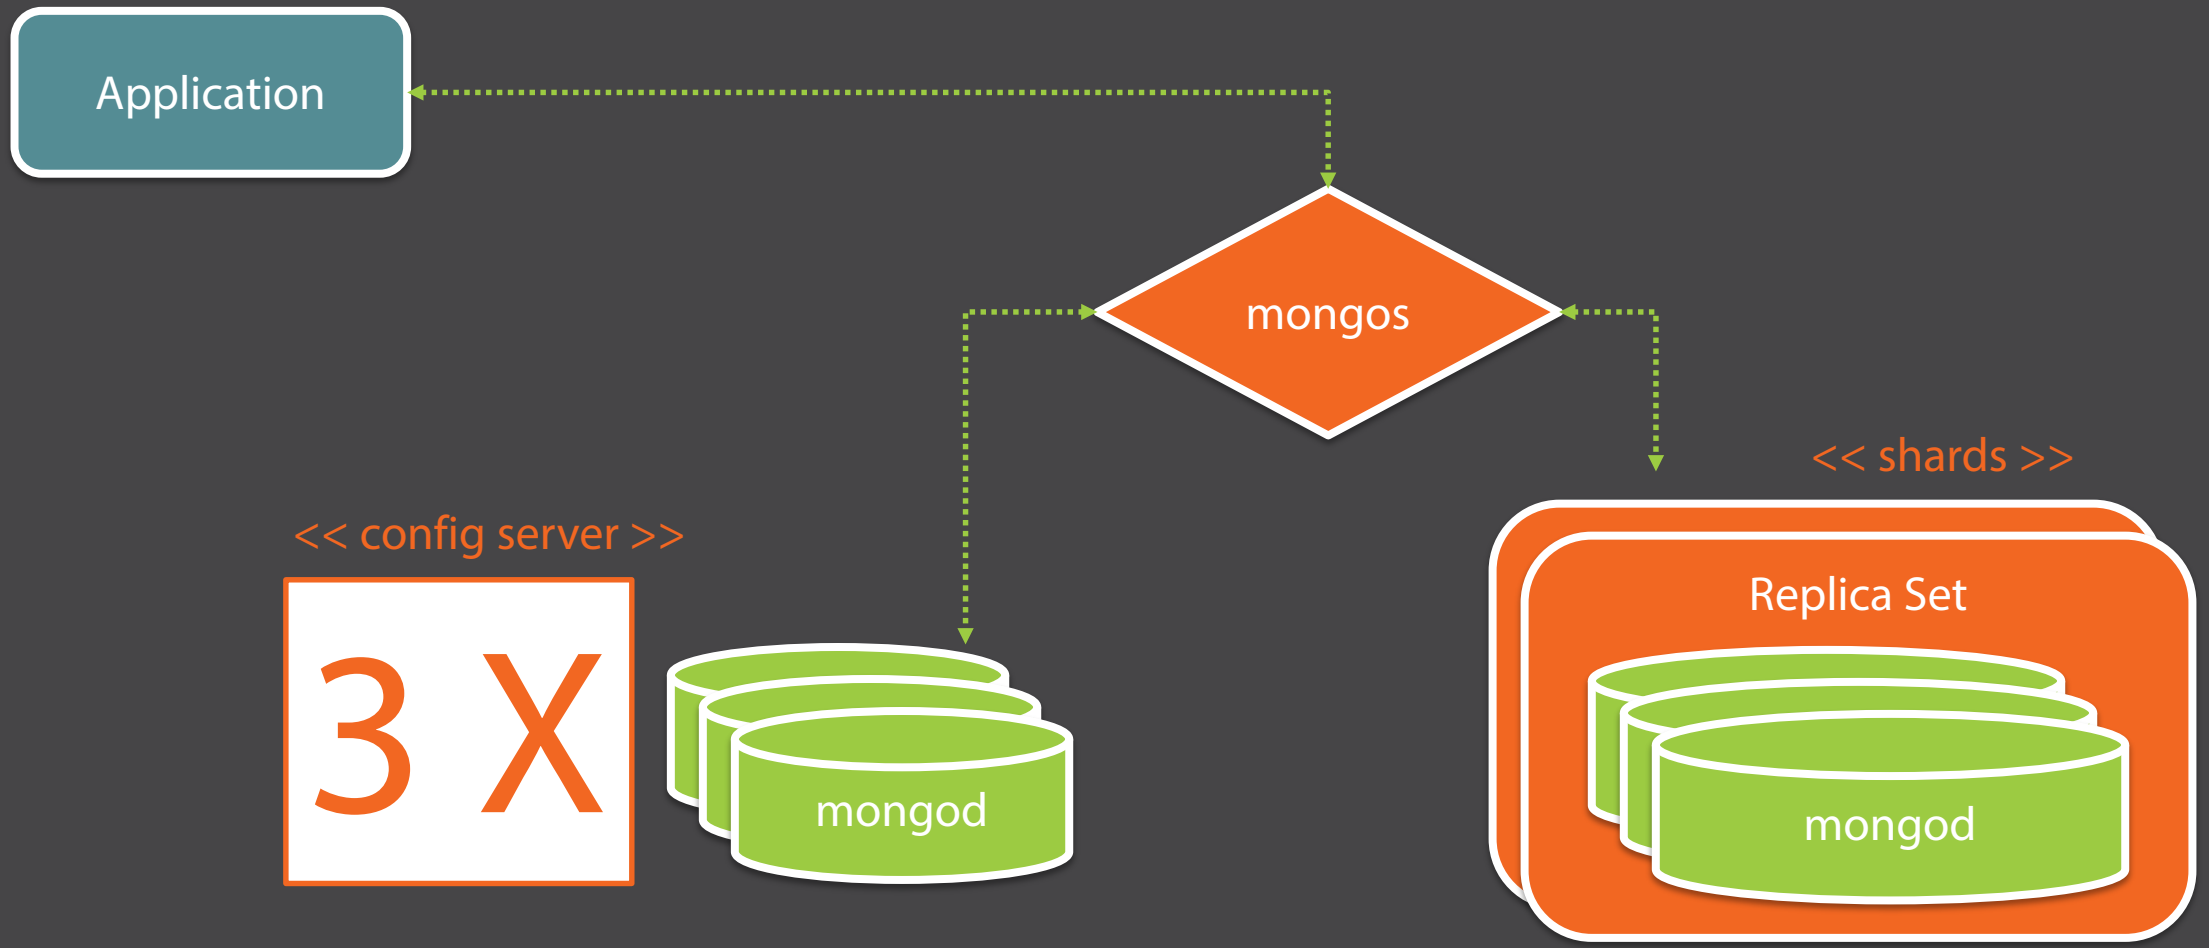

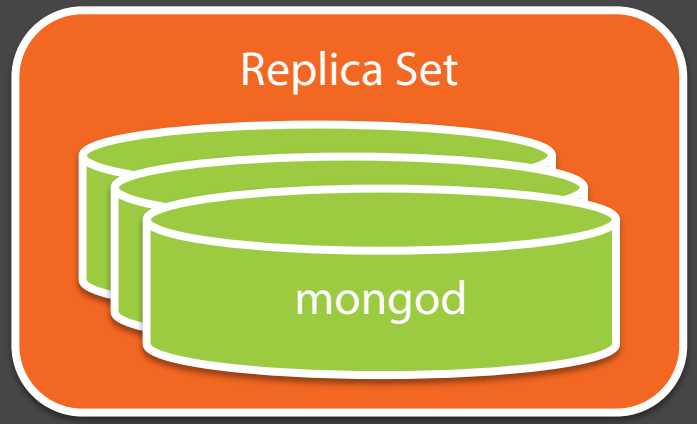

- Durability
- Availability

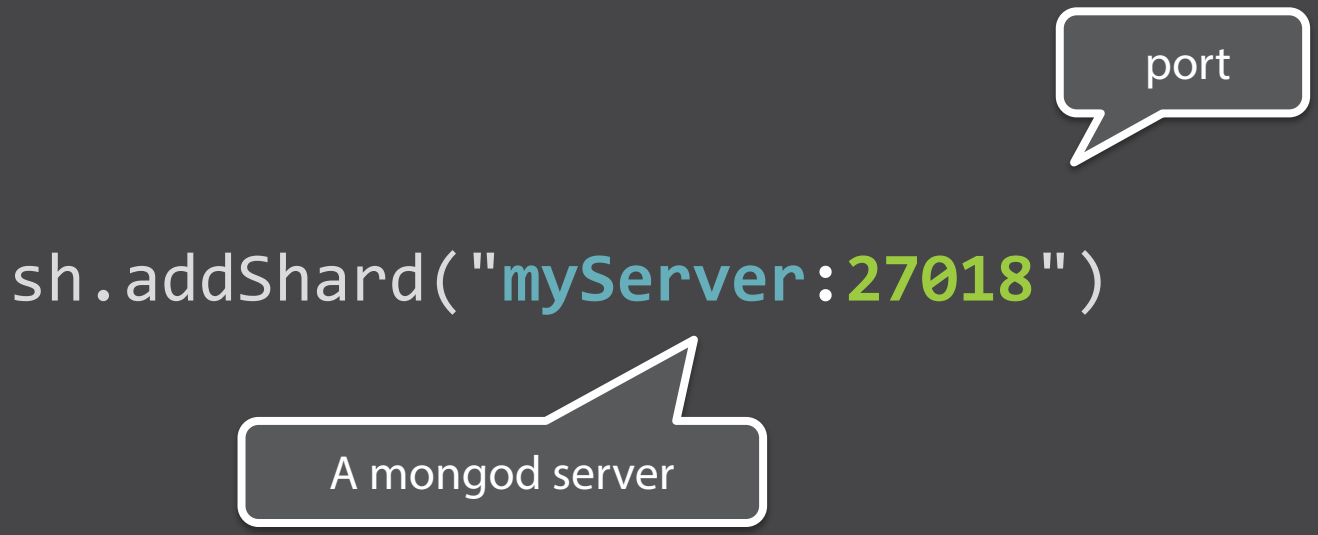

Adding a Shard

- Specify server + port

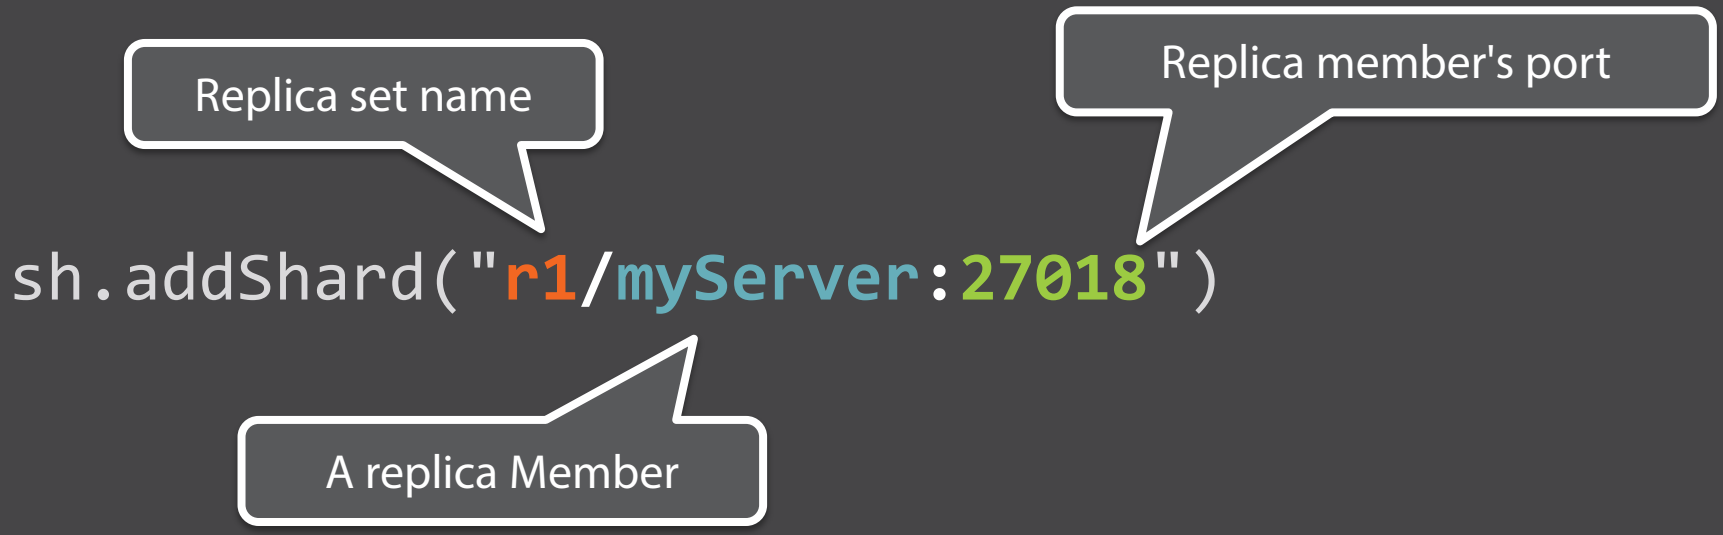

Adding a Replica Set as a Shard

- Prefix with replica set name
- Specify one of the replica set member + port
- The rest of the members in the replica set are discovered & registered as well

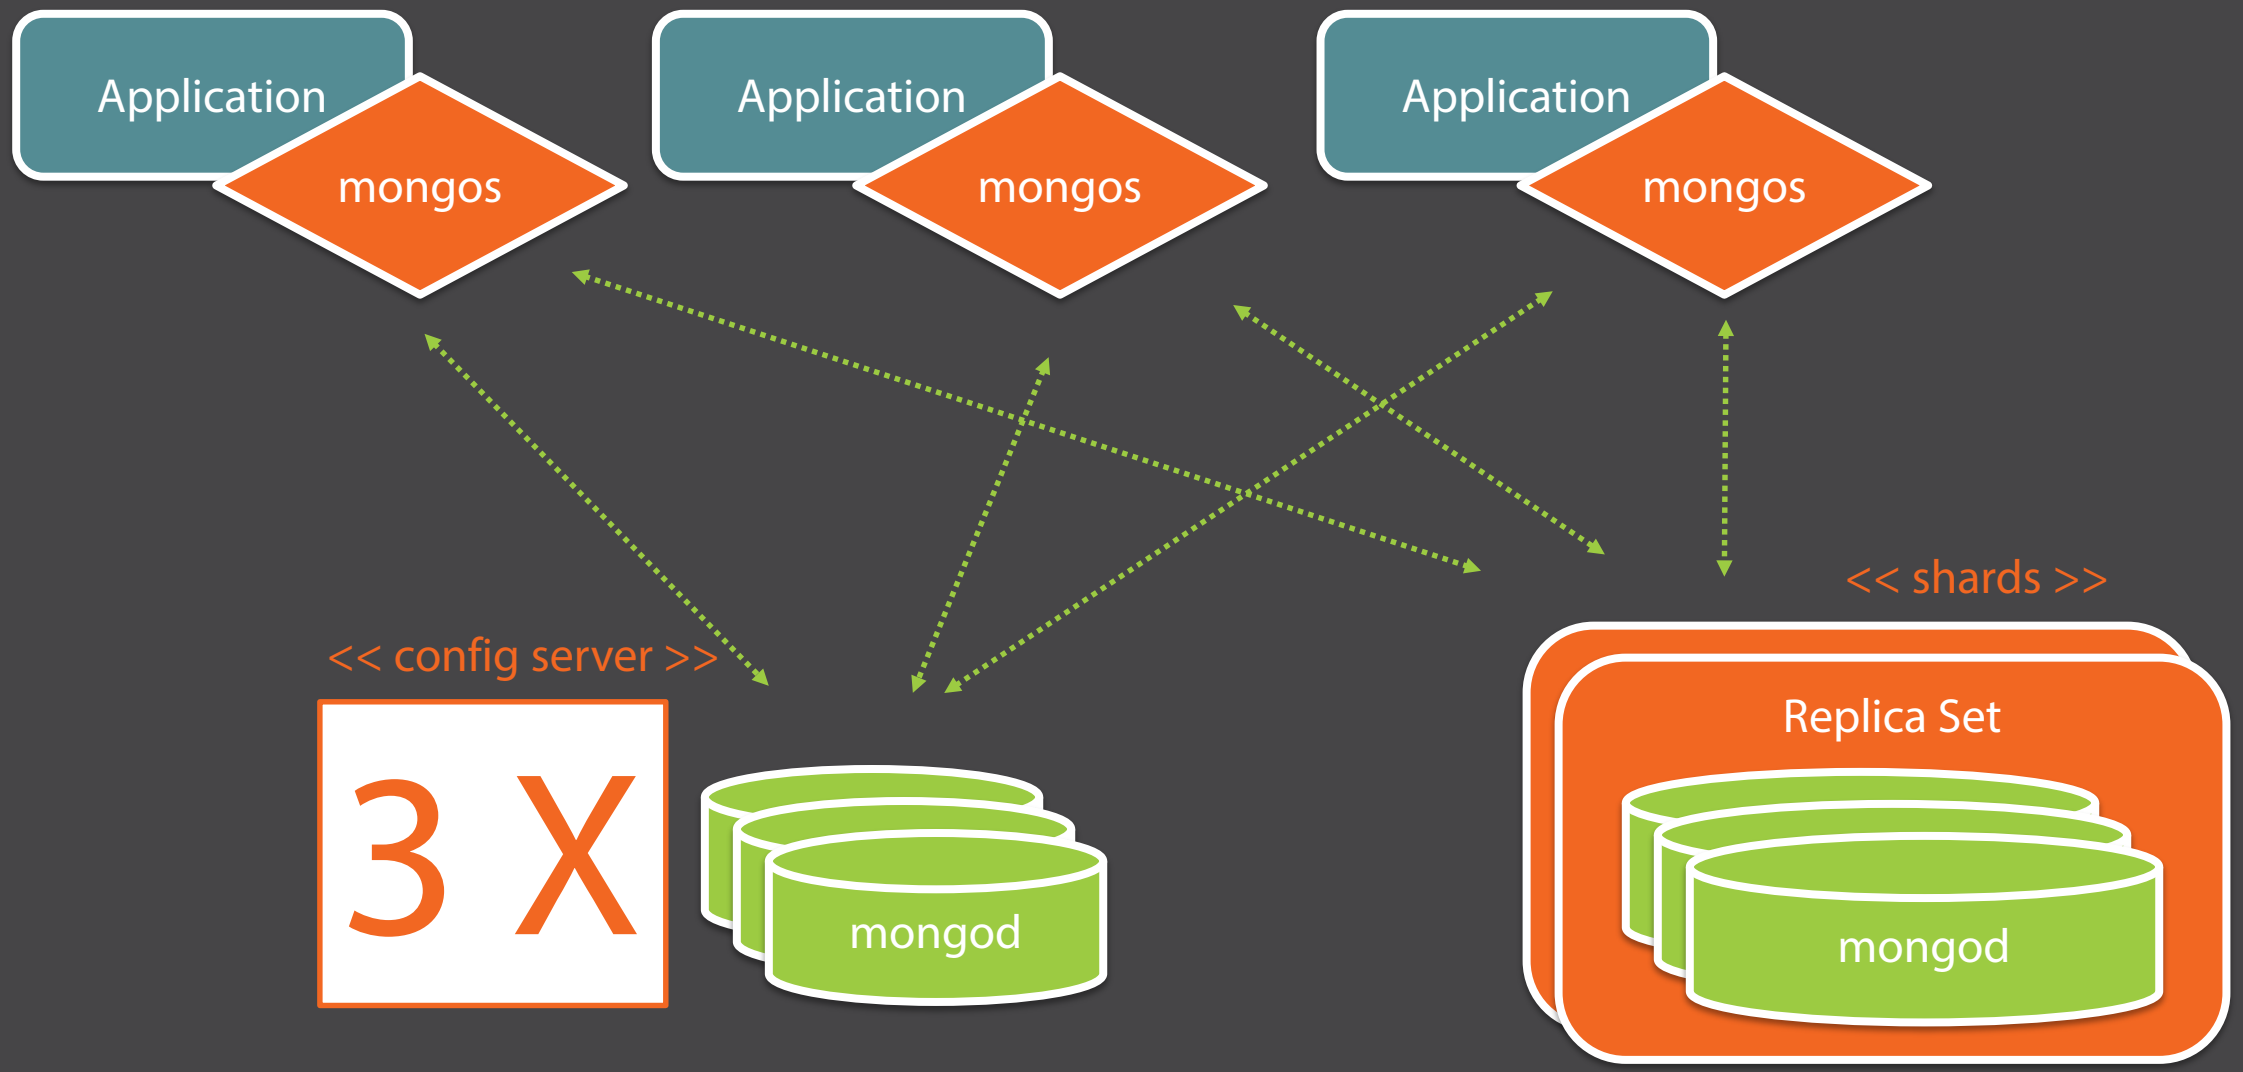

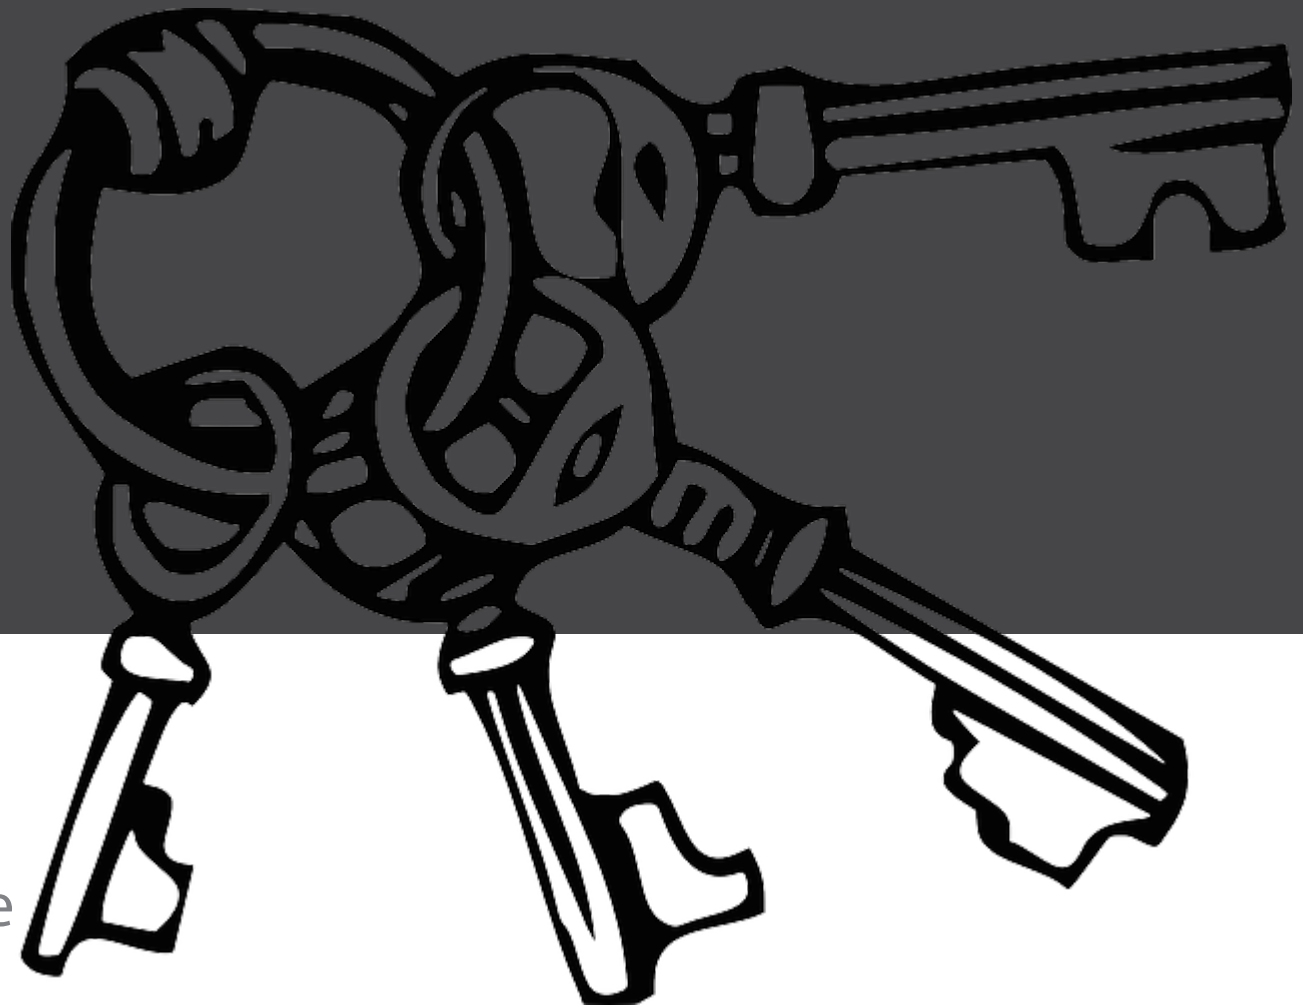

Shard Key Choices

- Important to get "right" the first time
- A balance act to get "right"

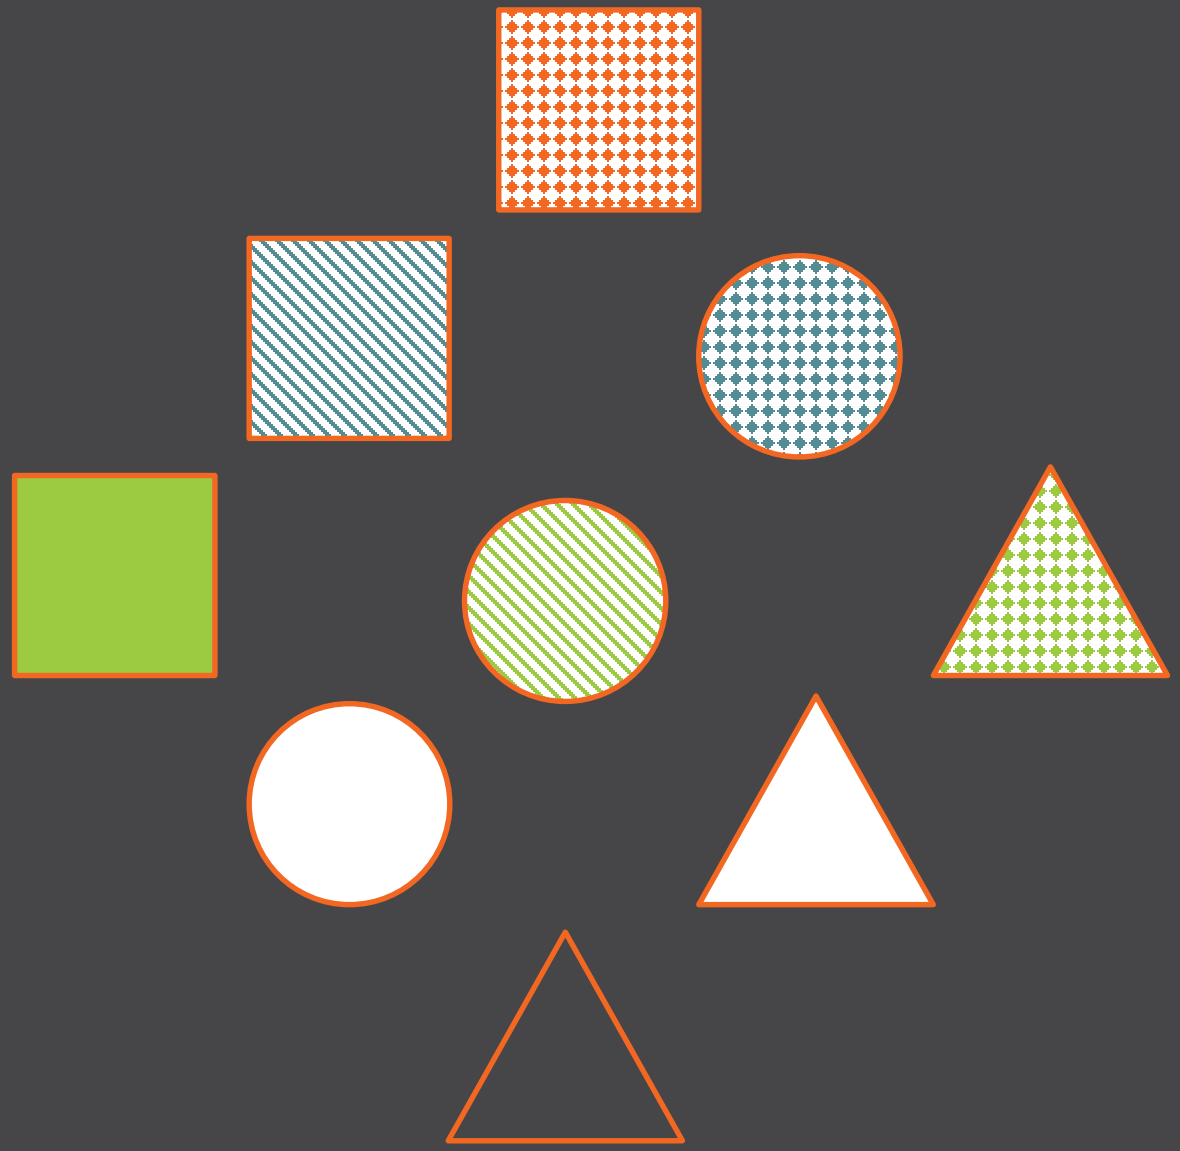

Picking a Shard Key

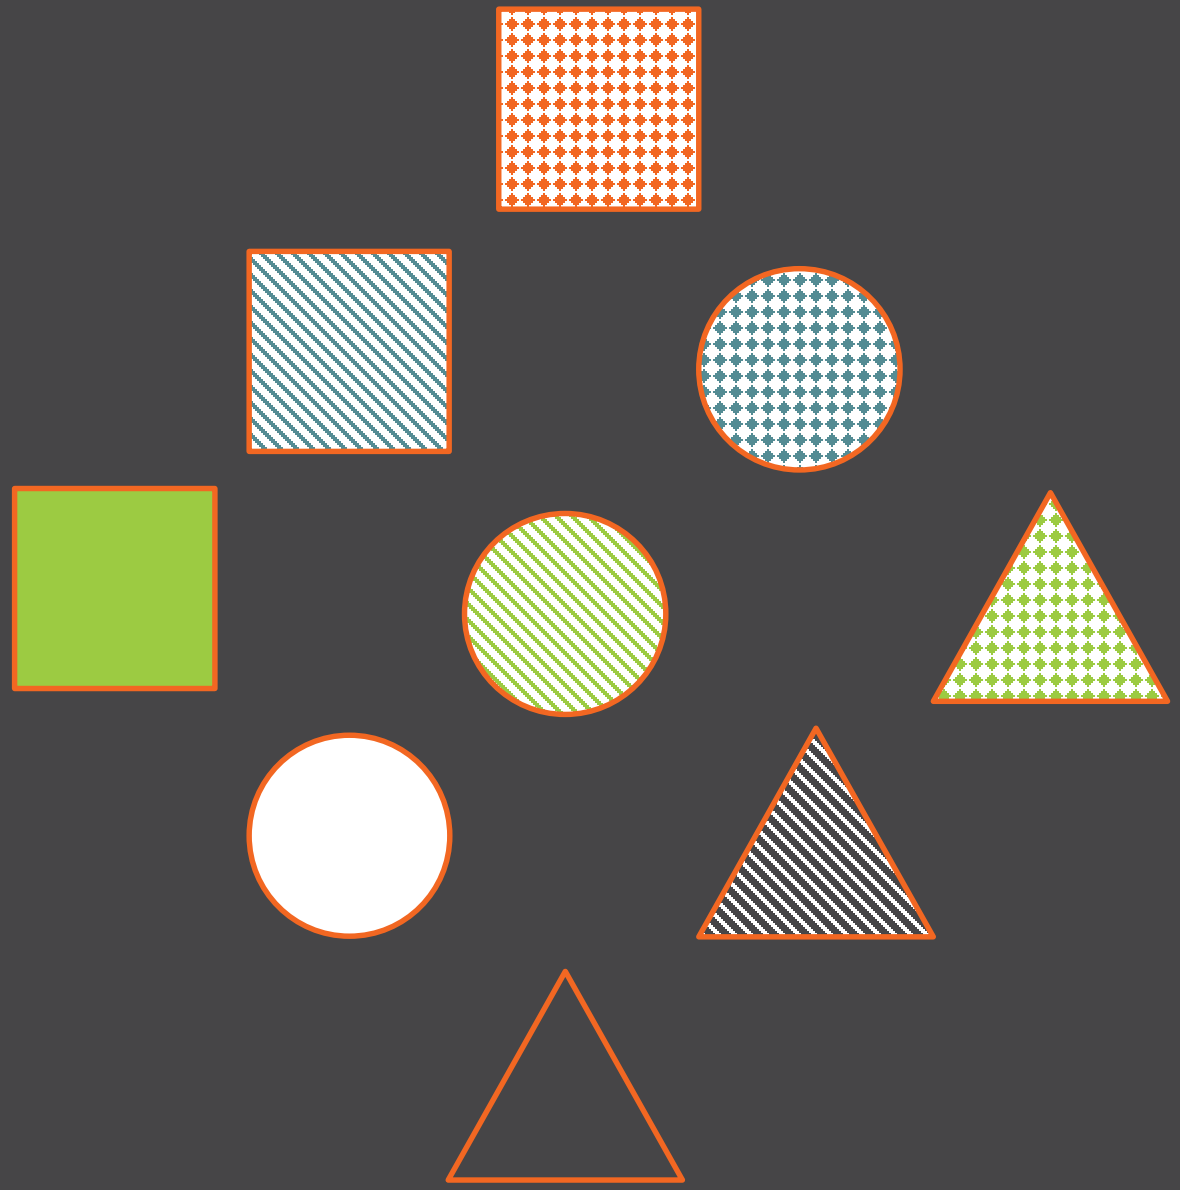

Picking a Shard Key

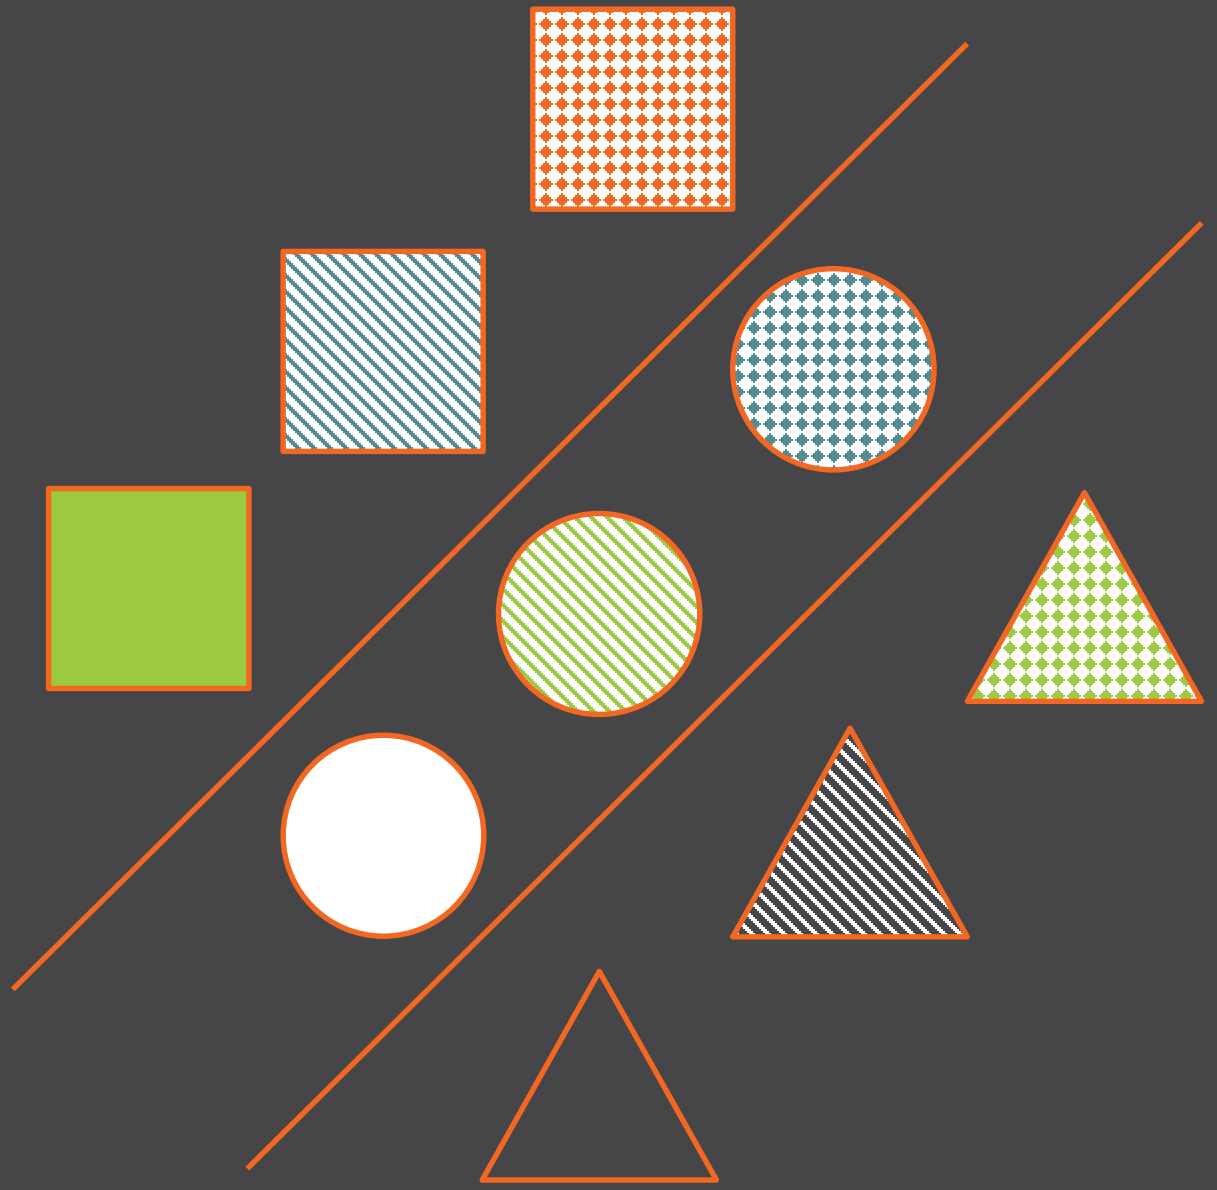

Picking a Shard Key

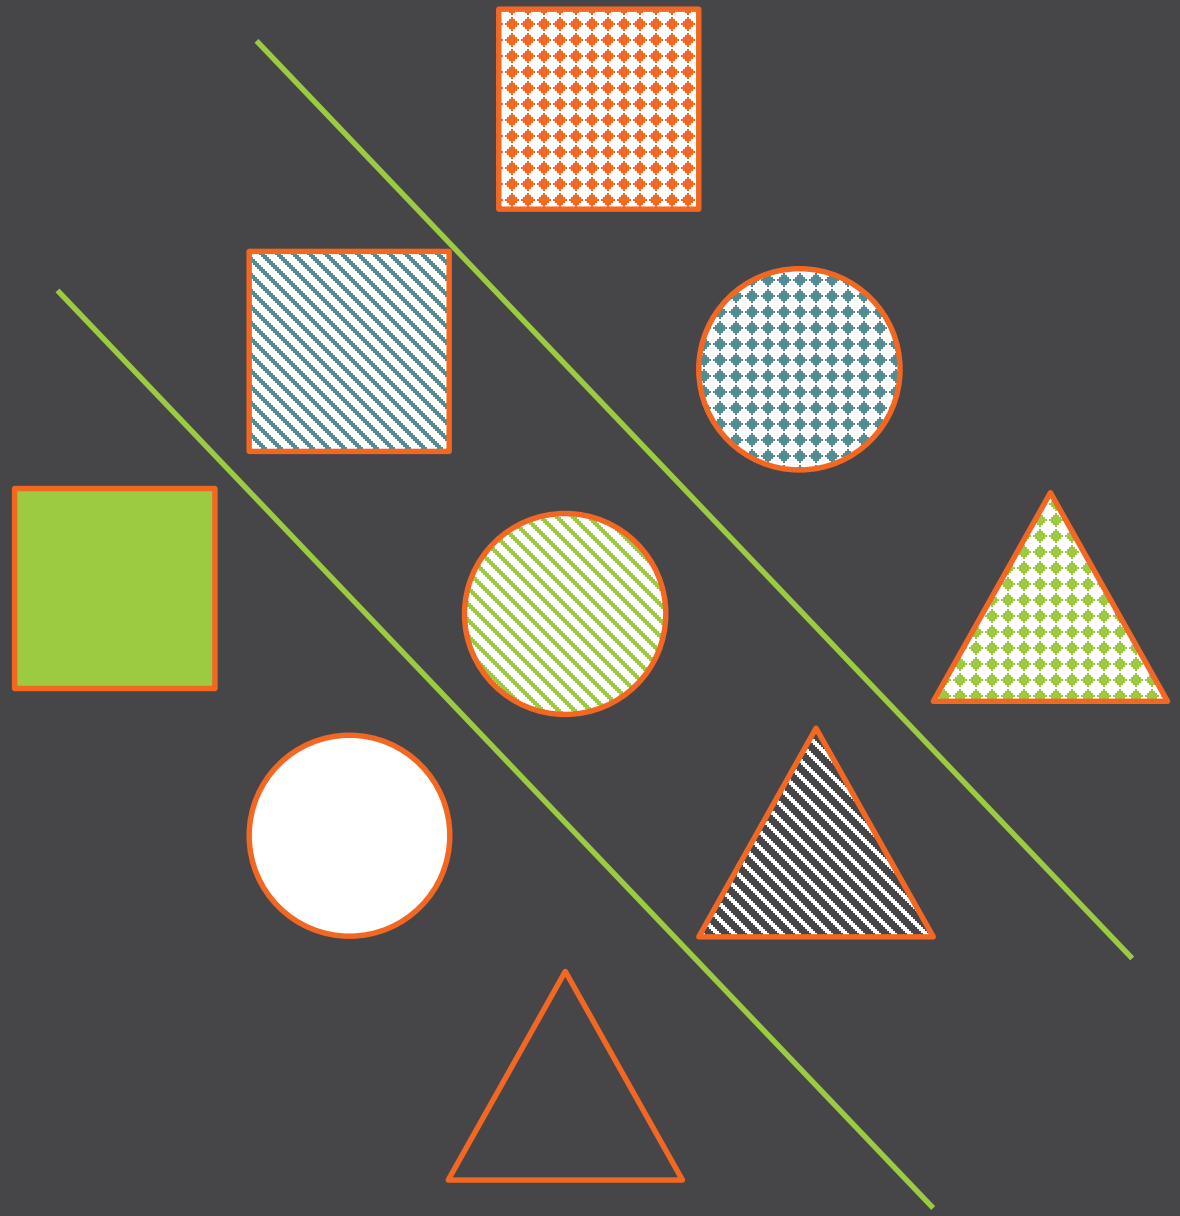

Picking a Shard Key

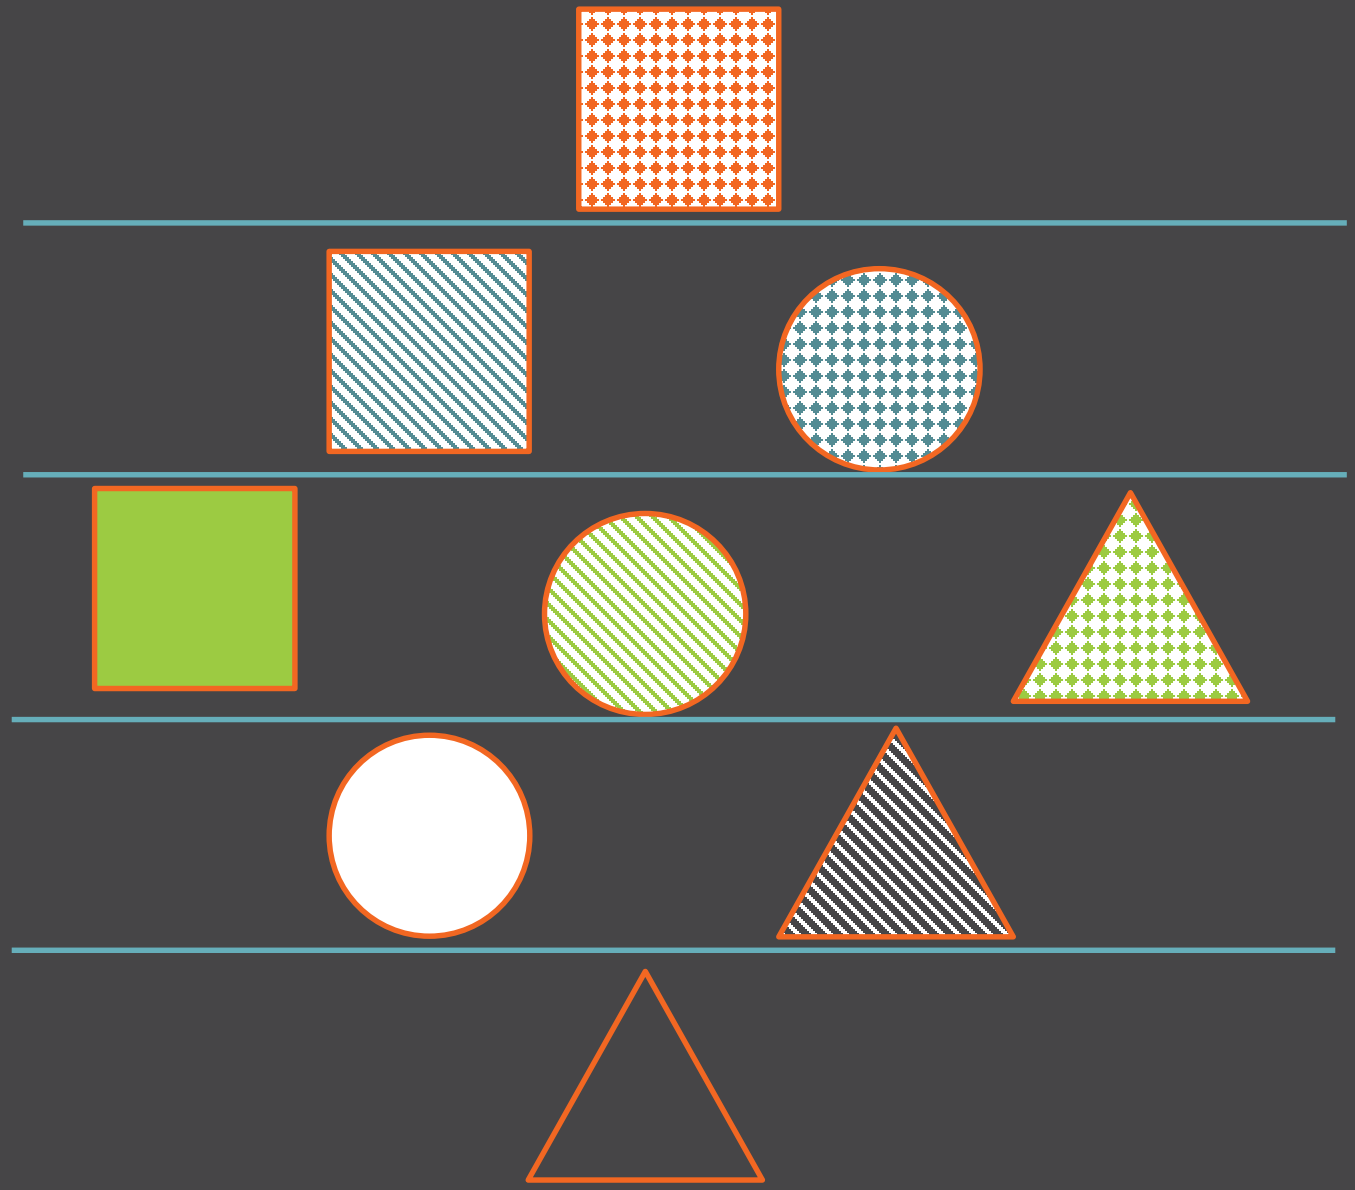

Picking a Shard Key

Sharding Rules

Can't change shard key definition

Can't "un-shard" a collection

Shard key field value can't changed

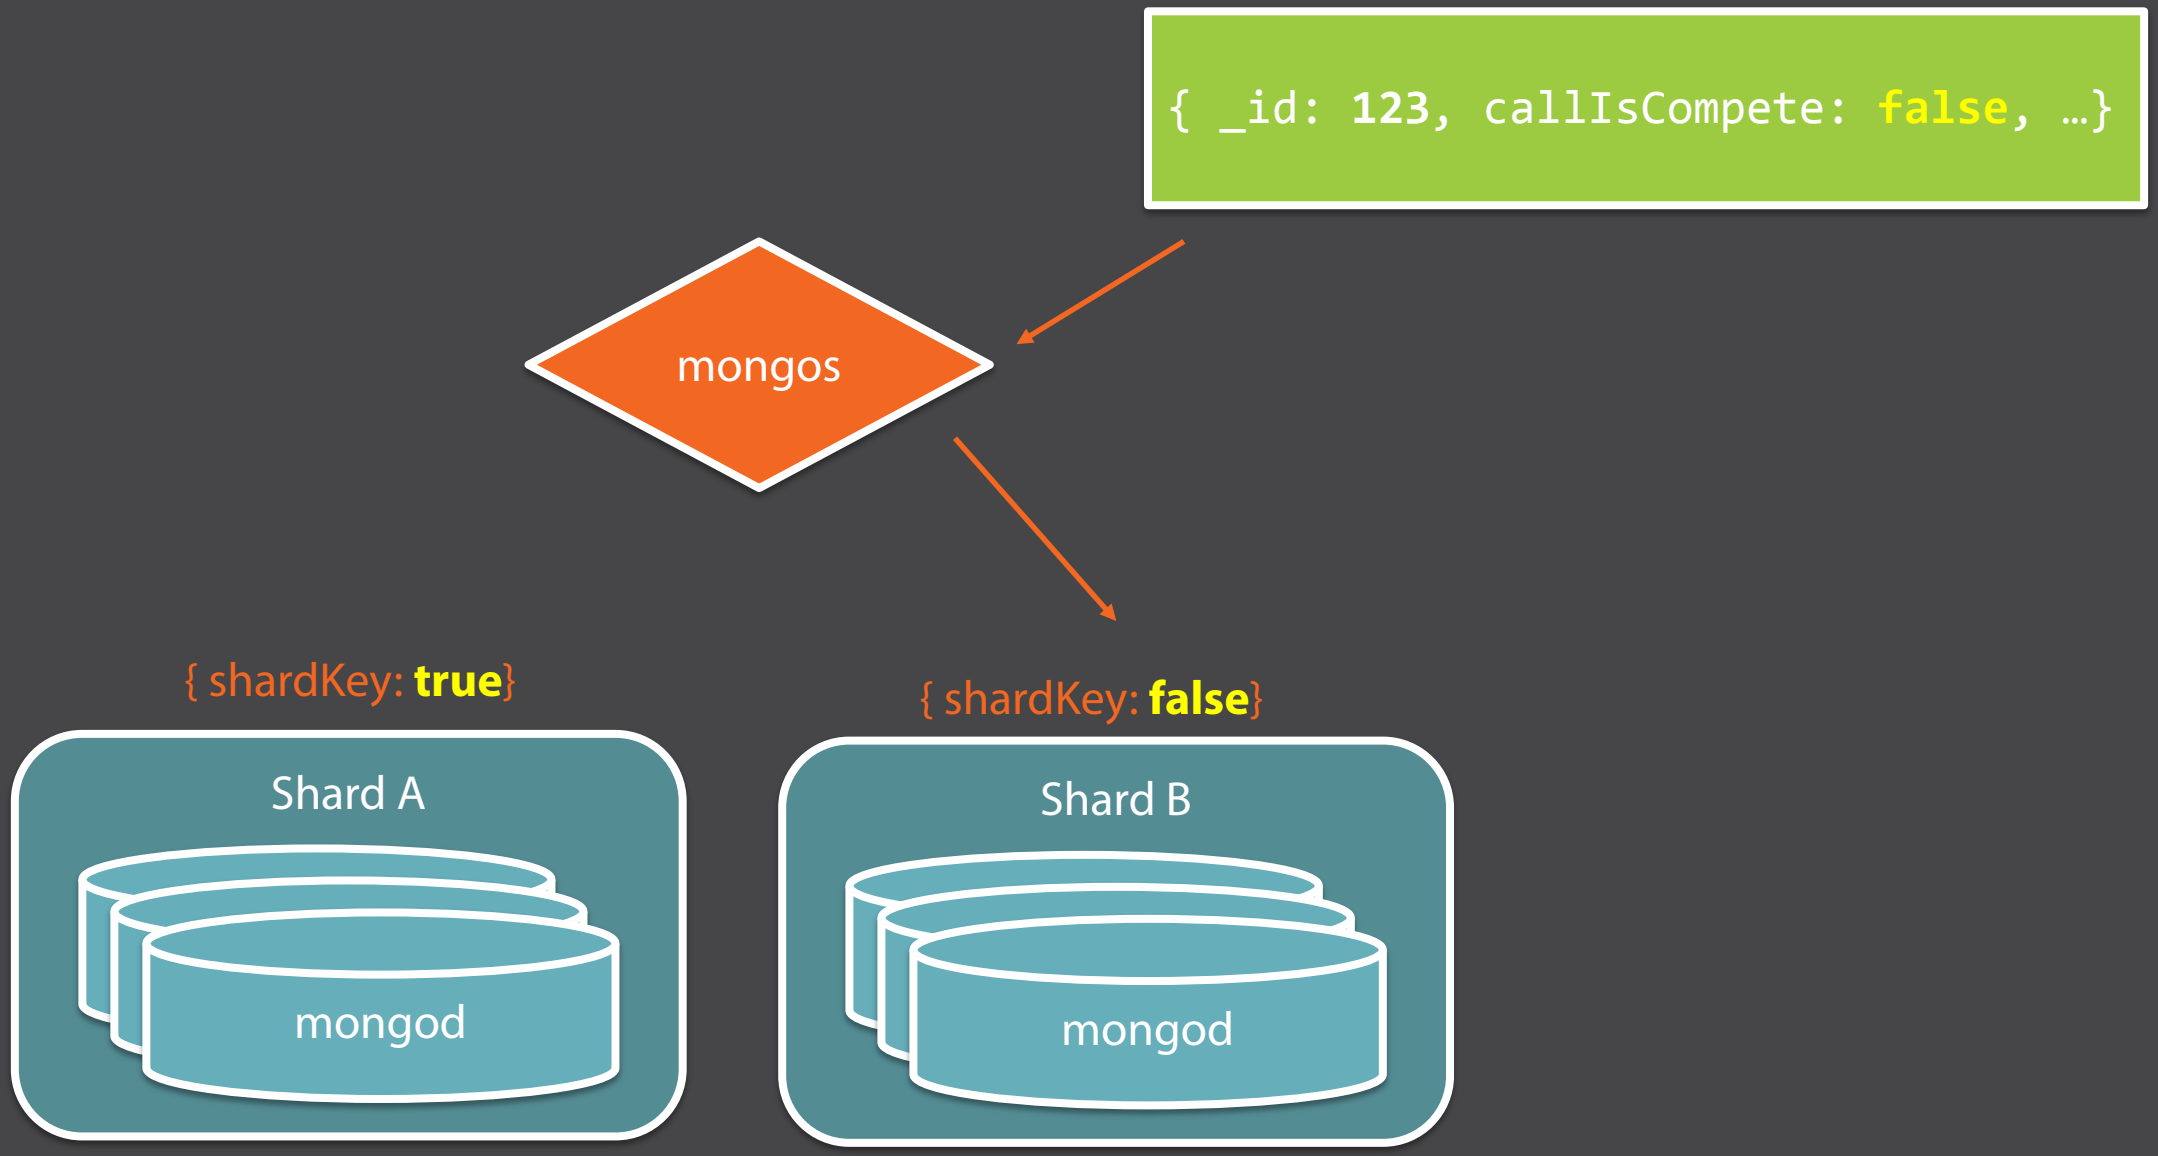

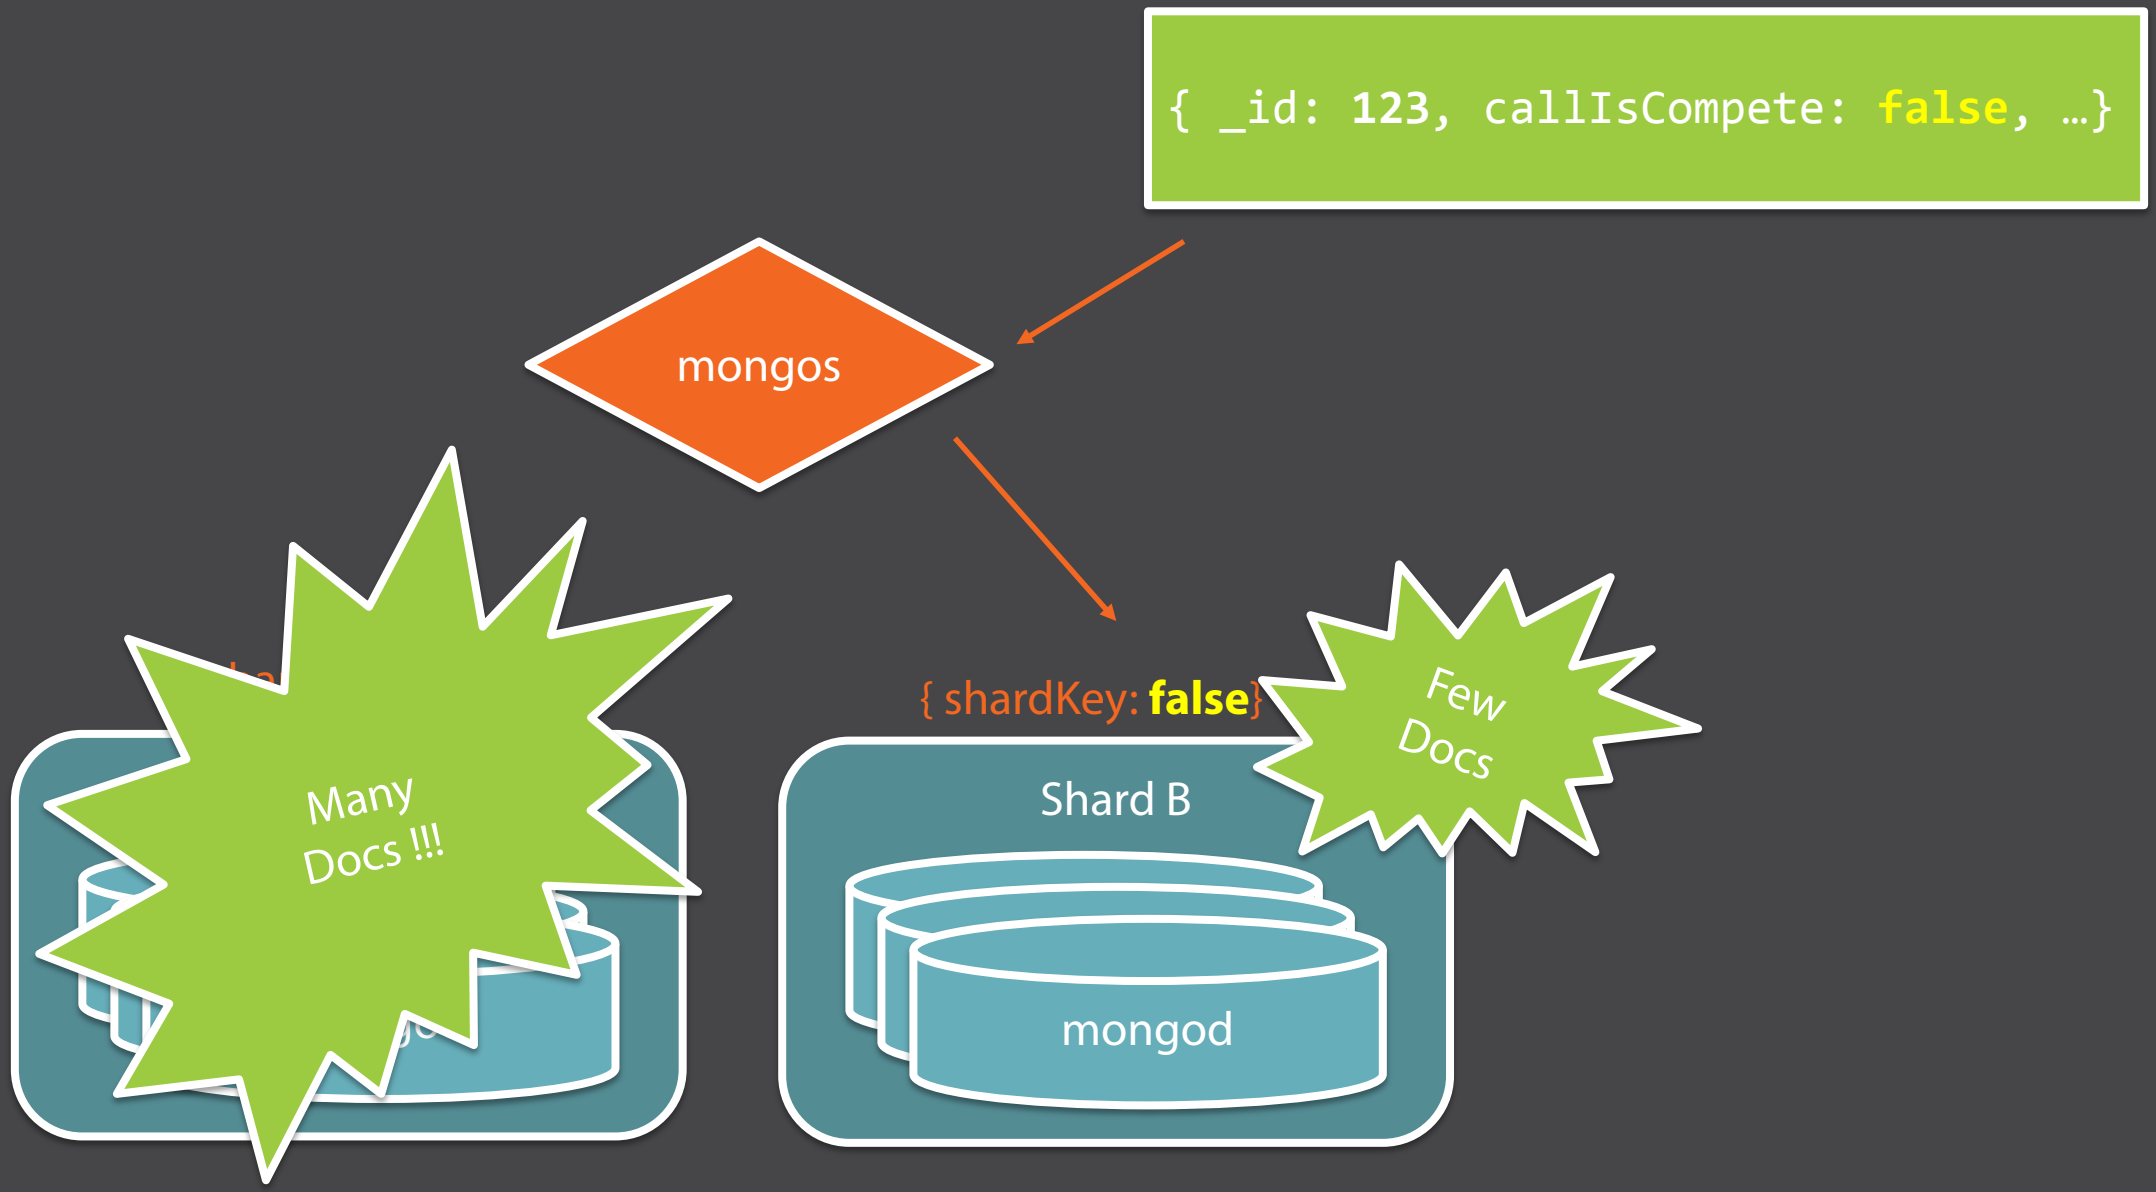

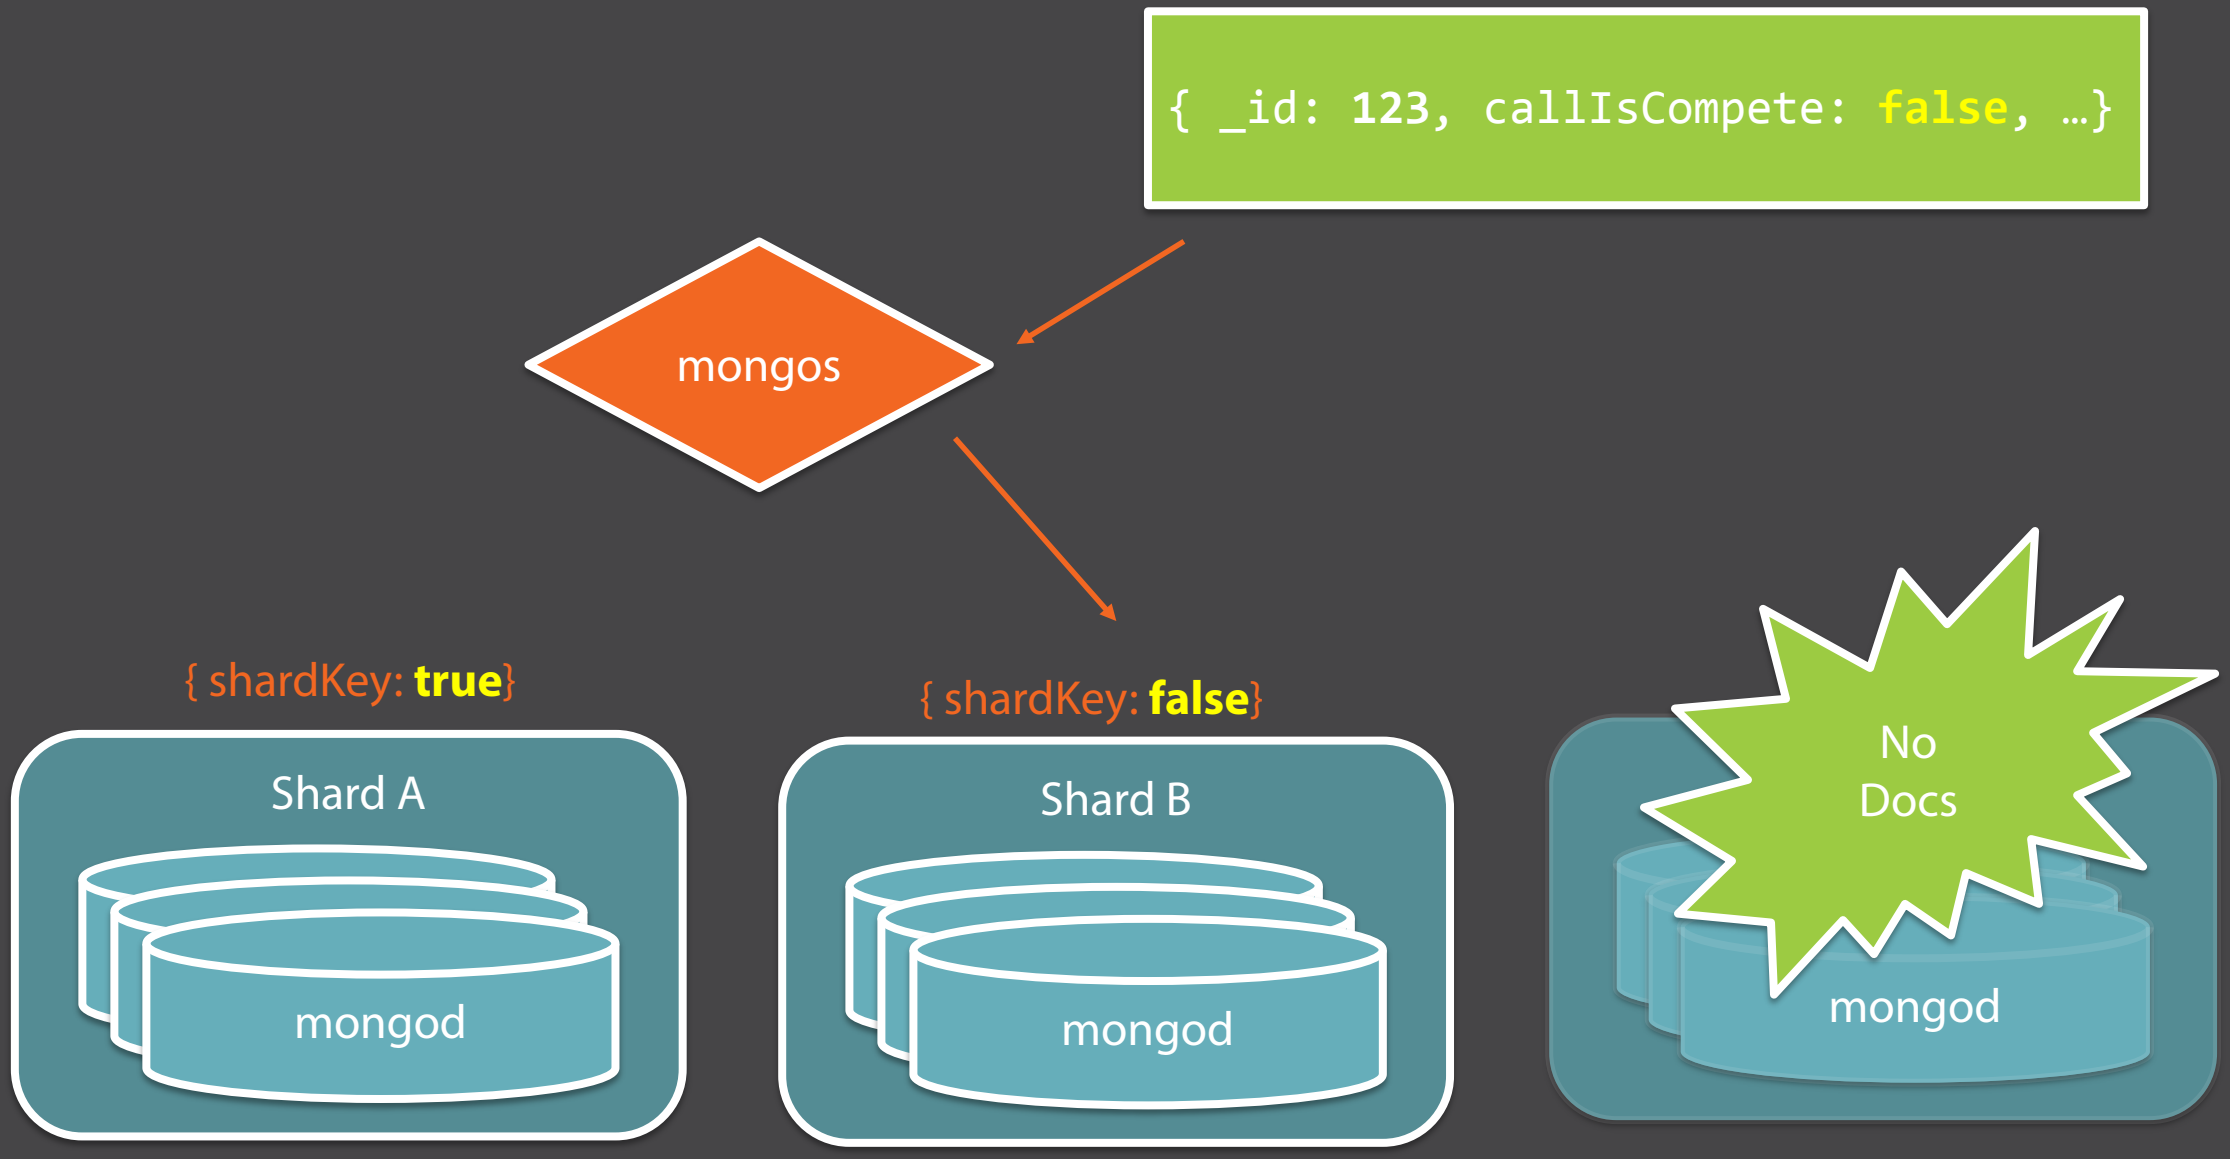

```
Chunk -> {minKey: 1, maxKey: 4}
```

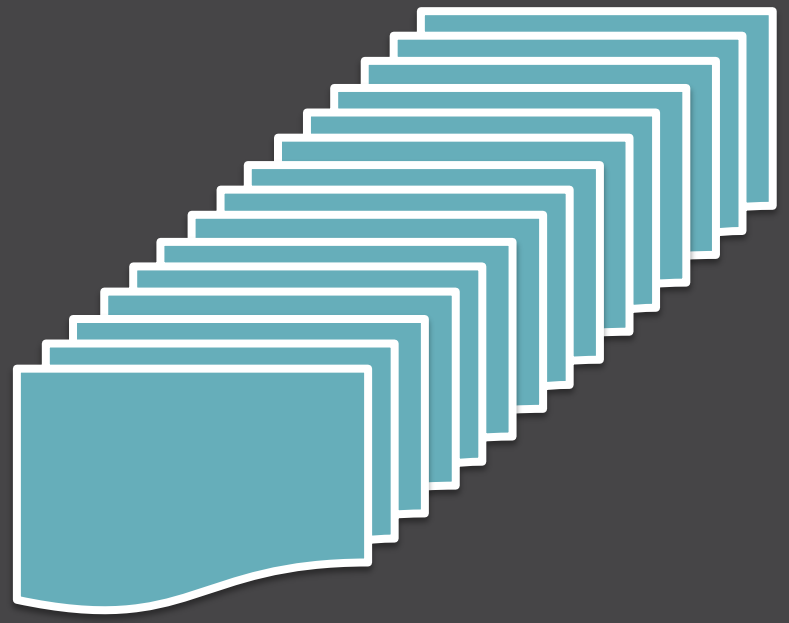

```
Chunk -> {minKey: 1, maxKey: 2}
```

```
Chunk -> {minKey: 2, maxKey: 4}
```

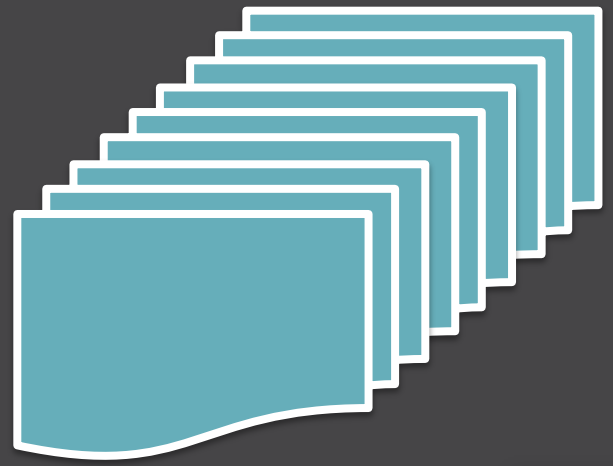

Chunk -> {minKey: 1, maxKey: 2}

Shard A

Chunk -> {minKey: 2, maxKey: 4}

Shard B

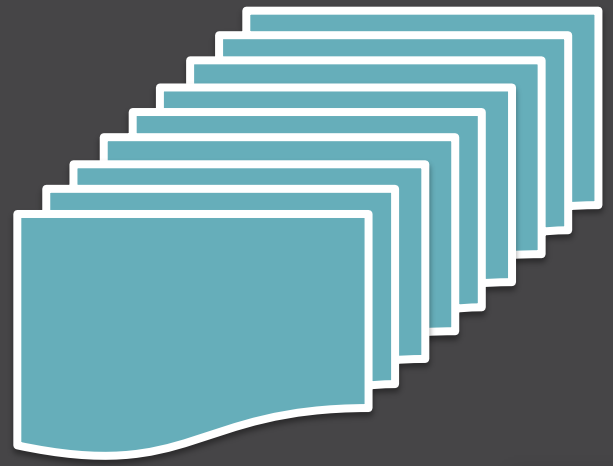

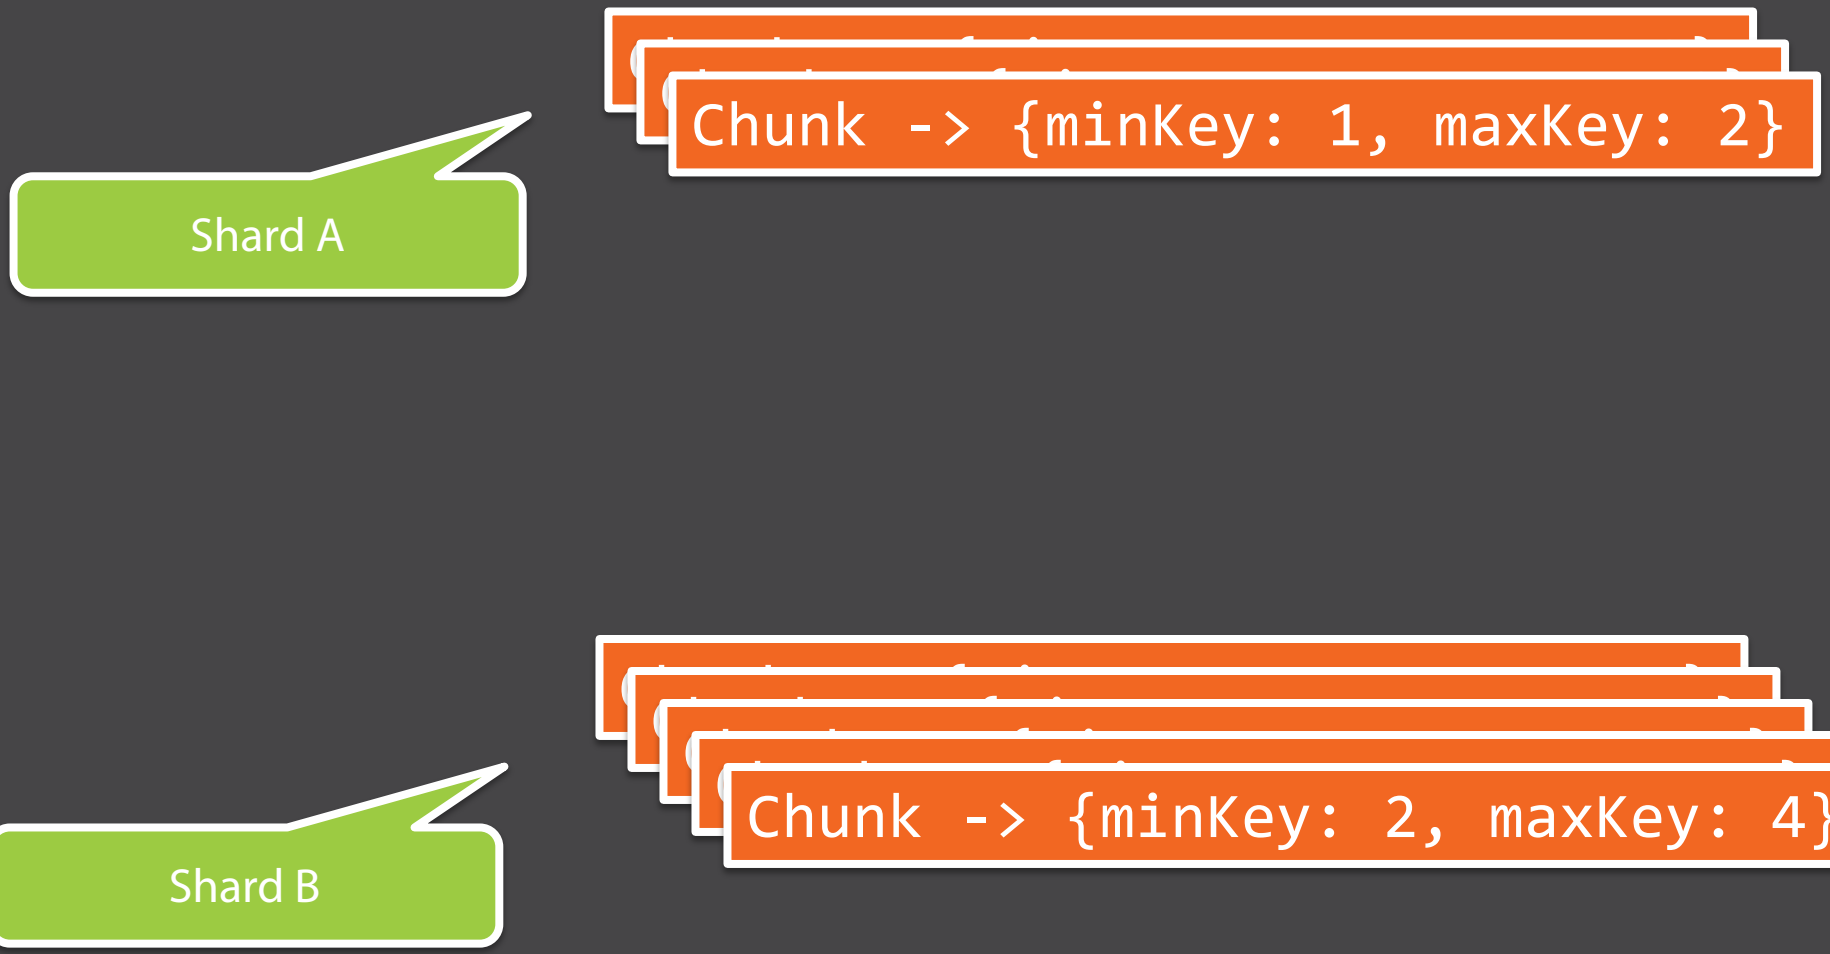

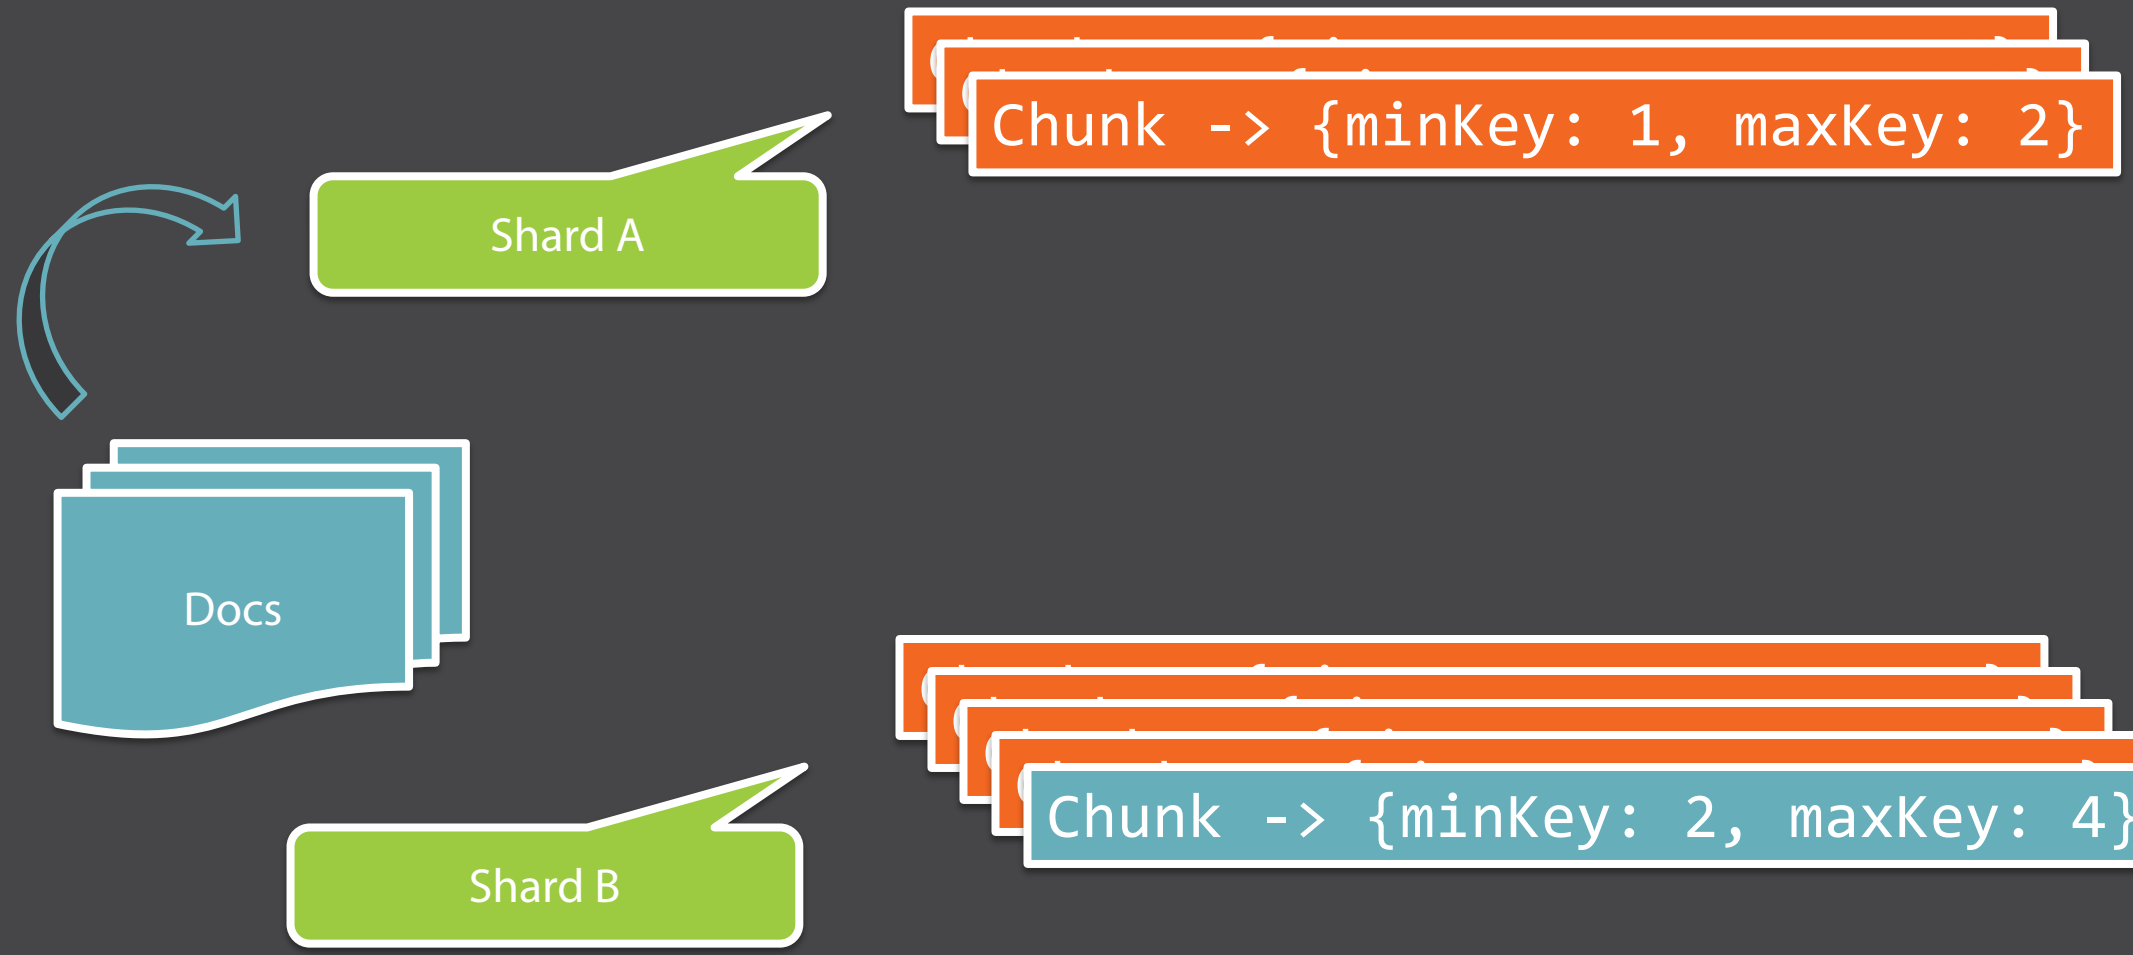

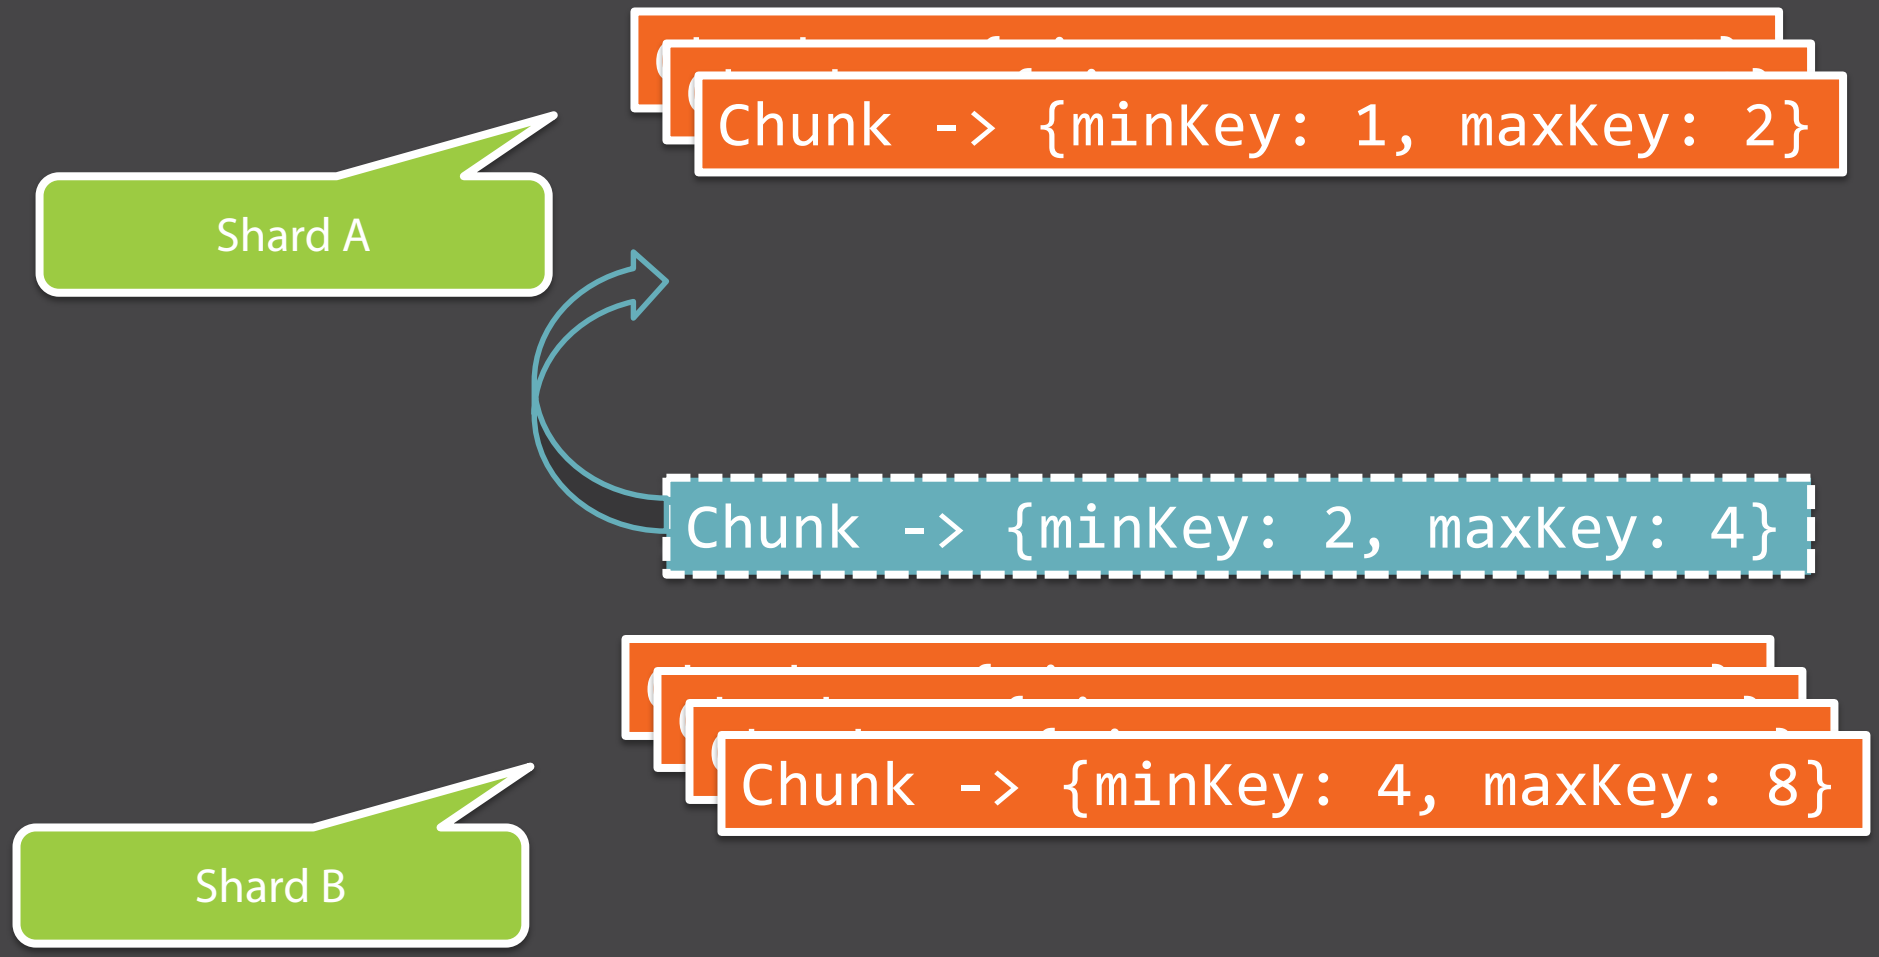

Sample Document

```
{
    _id: ObjectId(),
    from: { country: '01' ,  number: '333-555-1212' },
    to: {  country: '01' , number: '444-555-1212' },
    text:'Hi',
    time: ISODate()};
}
```

Choosing a "Good" Key

Your Access Pattern | Theoretical Granularity | Actual Granularity

Choosing a "Good" Key

Your Access Pattern

Theoretical Granularity

Actual Granularity

Choosing a "Good" Key

Your Access Pattern

Theoretical Granularity

Country + Number
> Billion Keys

Actual Granularity

Choosing a "Good" Key

Your Access Pattern

Theoretical Granularity

Actual Granularity

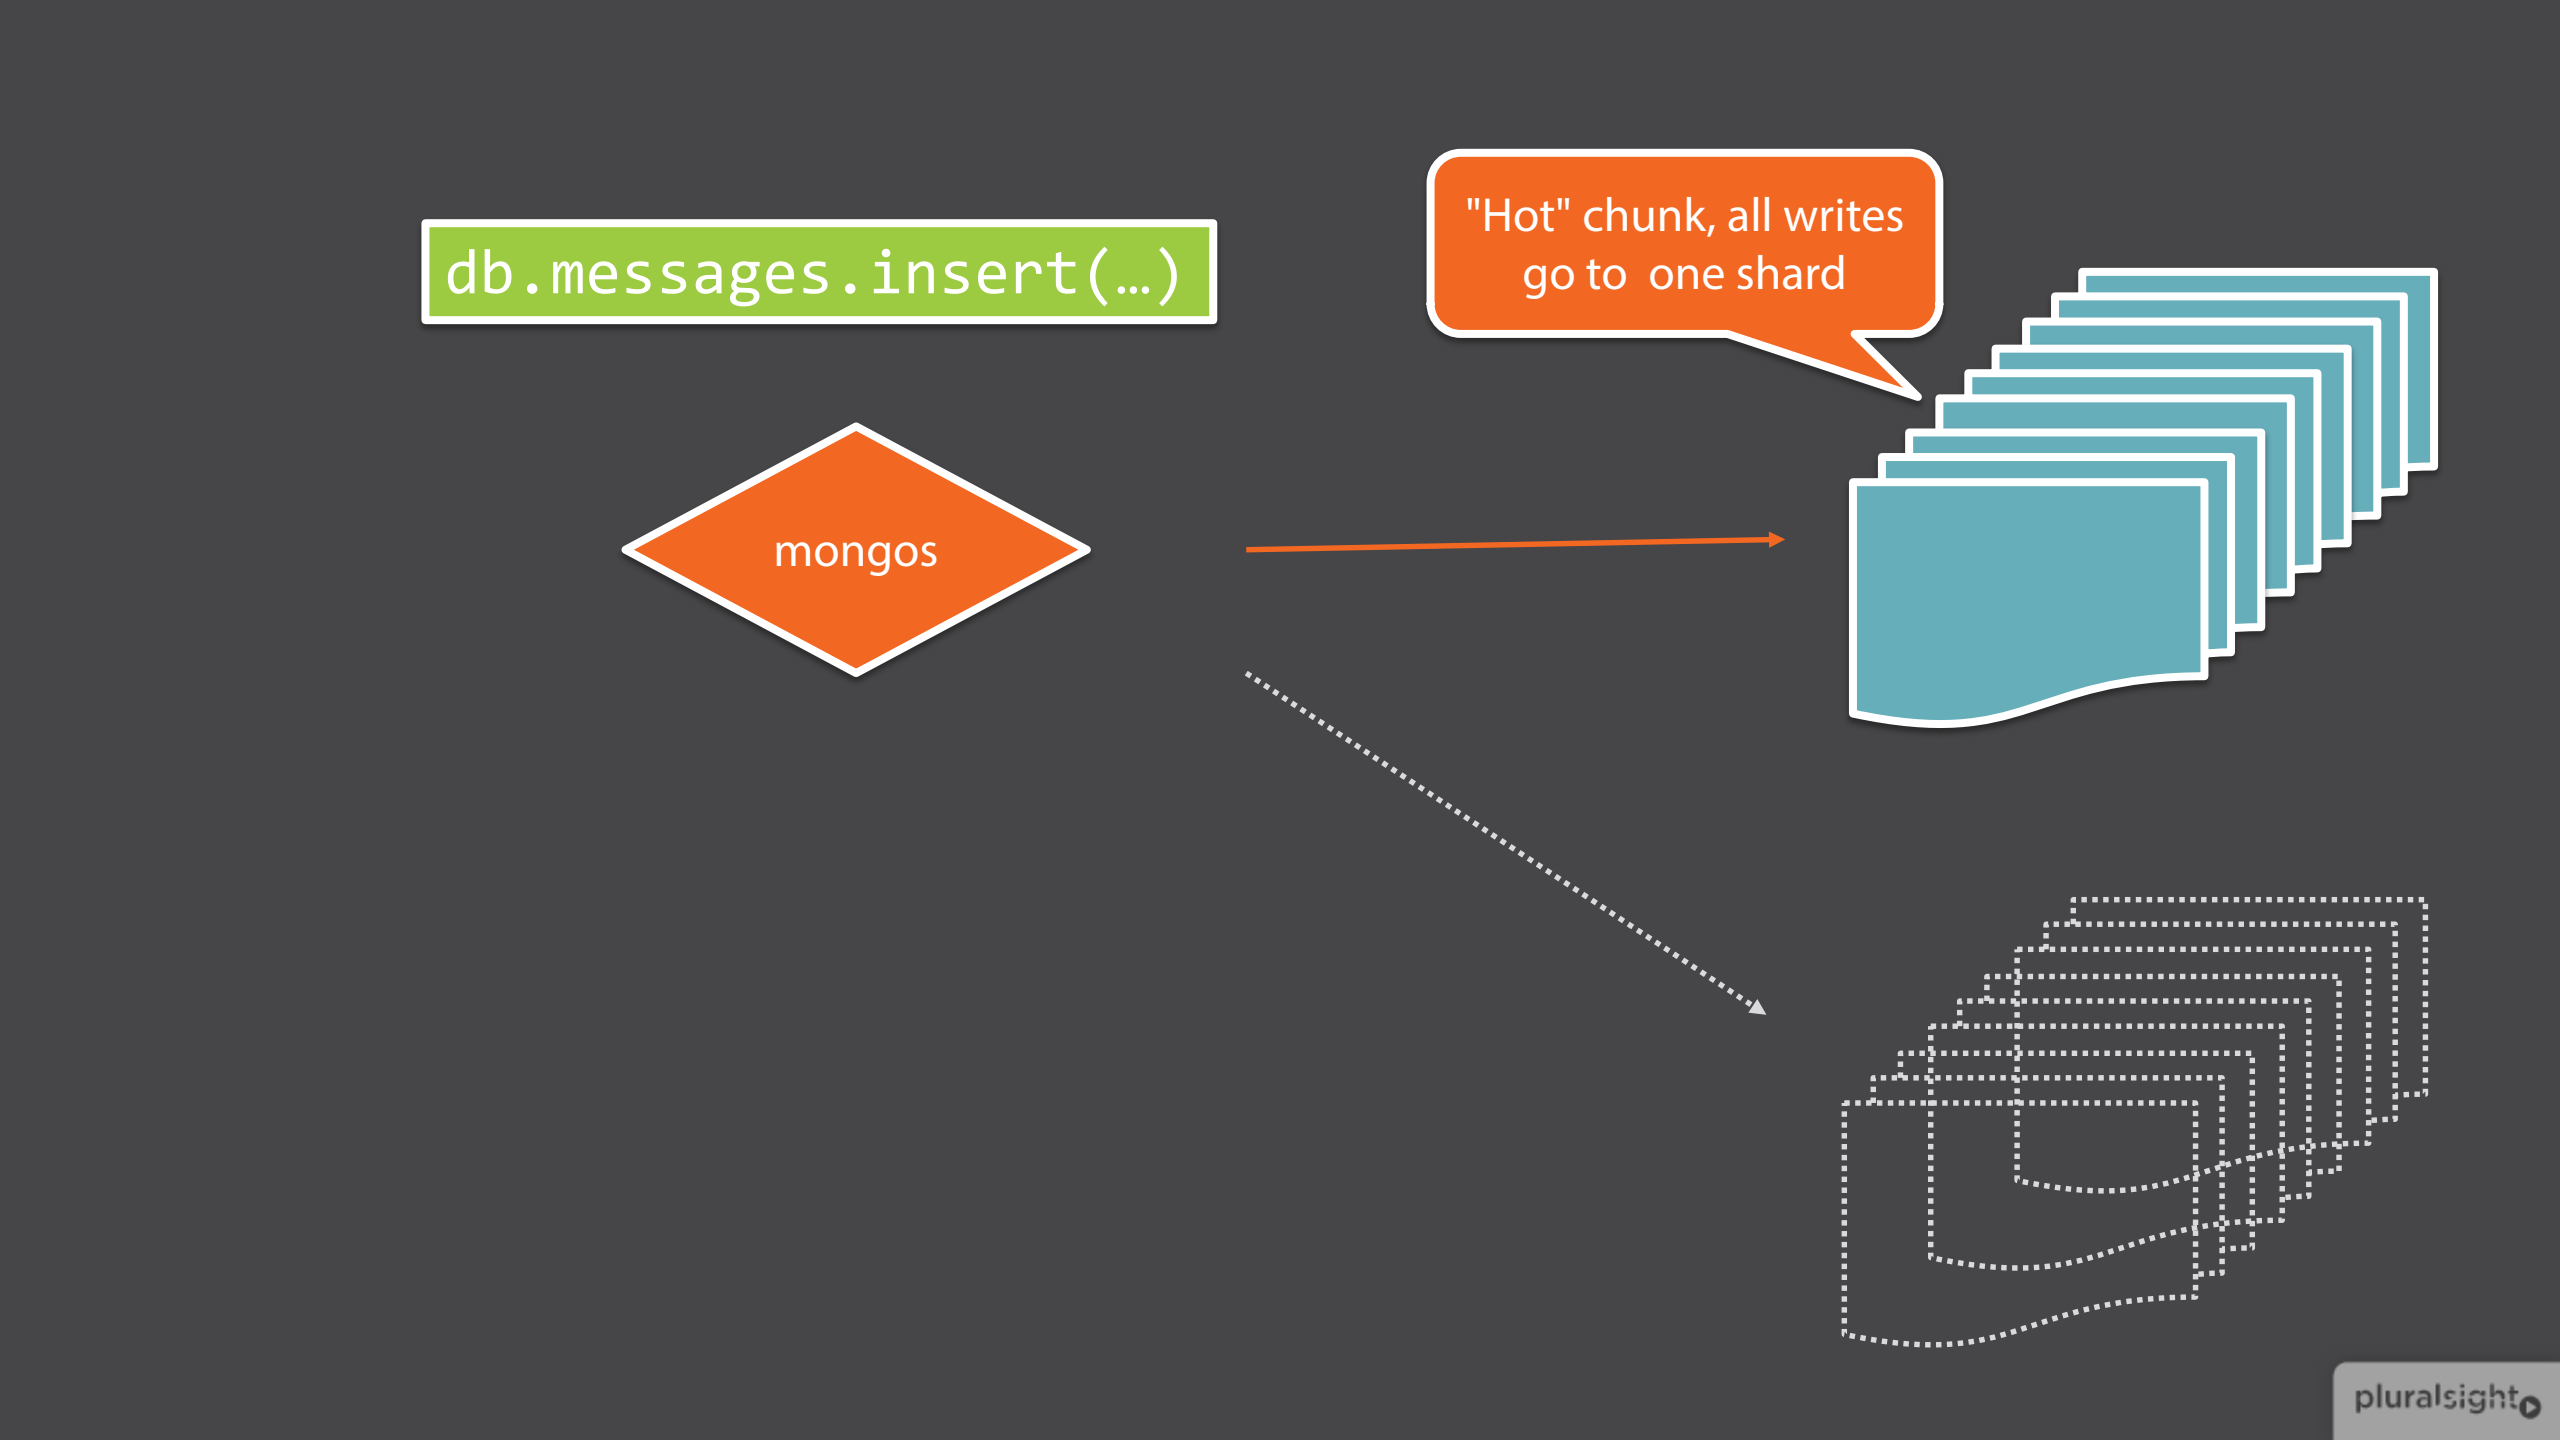
db.messages.insert(…)
"Hot" chunk, all writes go to one shard
mongos

Choosing a "Good" Key

Your Access Pattern

Theoretical Granularity

Actual Granularity

```
sh.shardCollection(
    "dbName.collectionName",
    {"fieldName": "hashed"})
```

Hashed shard key

- Only on a single field
- Helps even distribution across shards

Choosing a "Good" Key

Your Access Pattern

Theoretical Granularity

Actual Granularity


Not Hashed
Shard A
{ April 1 }
{ April 2 }
{ April 3 }
Shard B
{ May 1 }
{ May 2 }
{ May 3 }
Hashed
Shard A
{ "B7Q" /* April 1 */ }
{ "B9T" /* May 2 */ }
{ "B9Z" /* April 3 */ }
Shard A
{ "H2O" /* May 3 */ }
{ "H3X" /* May 1 */ }
{ "I4M" /* April 2 */ }

Choosing a "Good" Key

Establish this …

Your Access Pattern

Theoretical Granularity

Actual Granularity

Find key to satisfy these…

Even Distribution

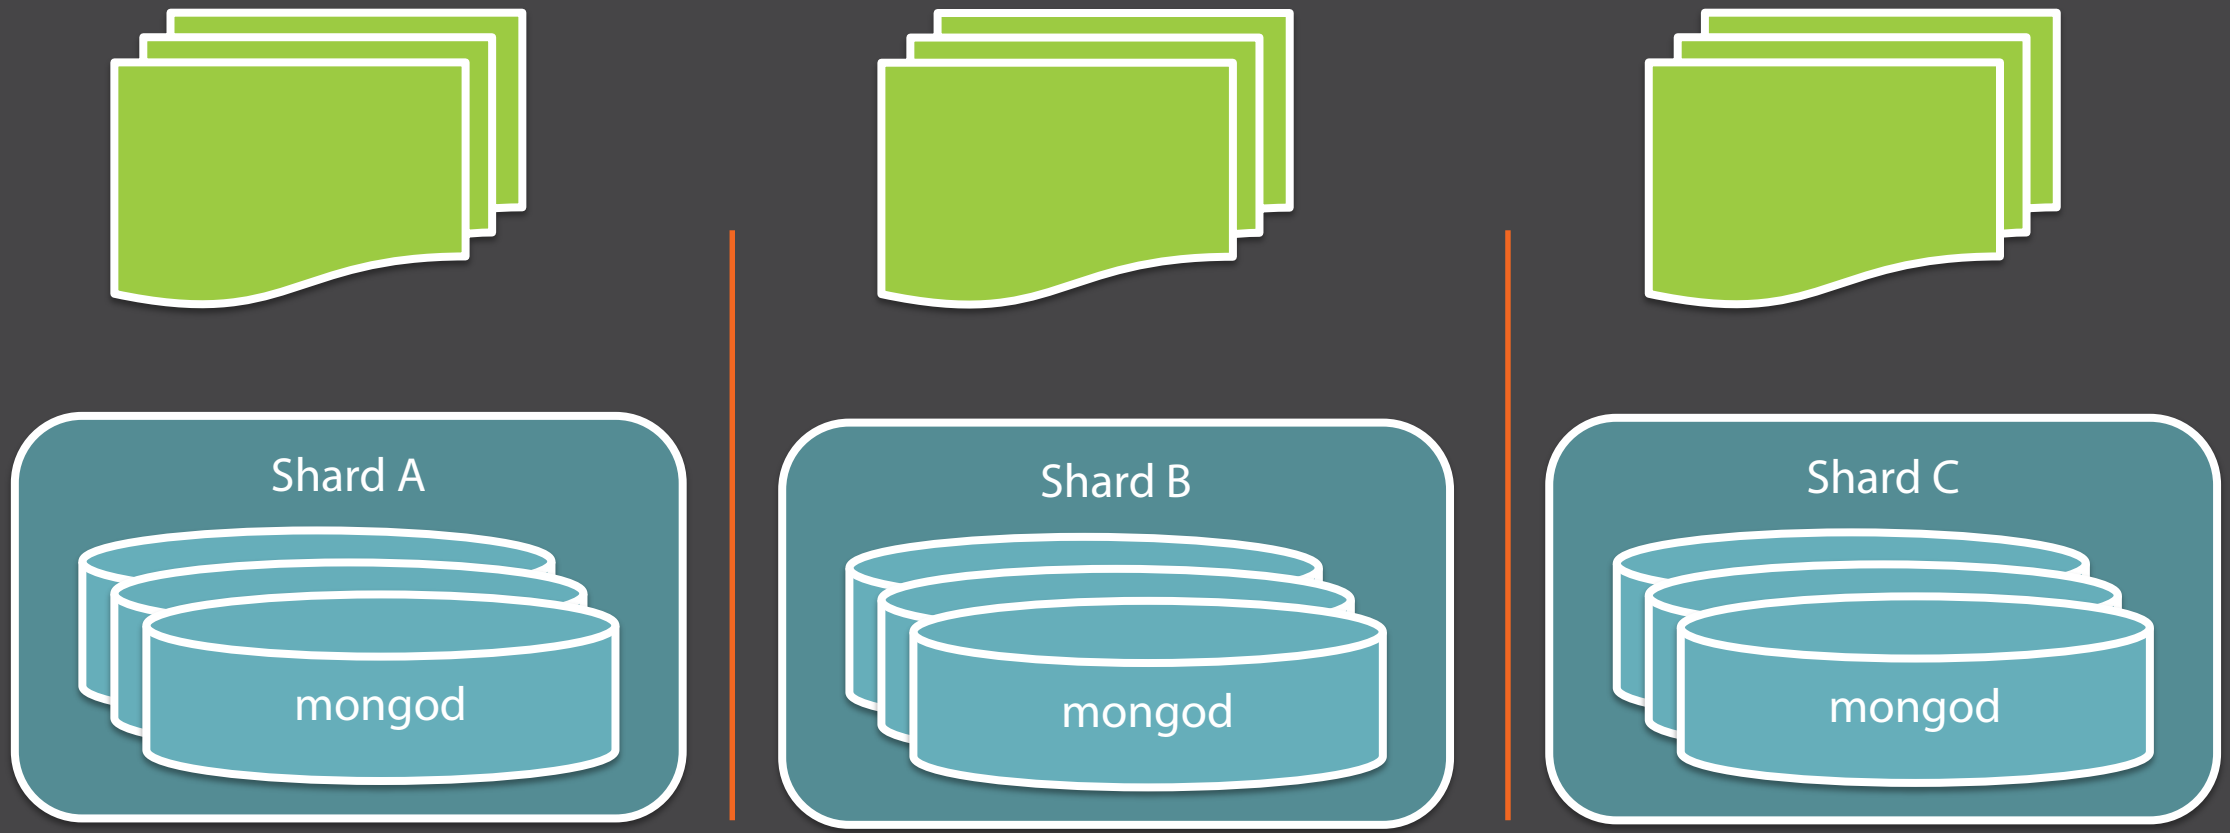

Even Distribution

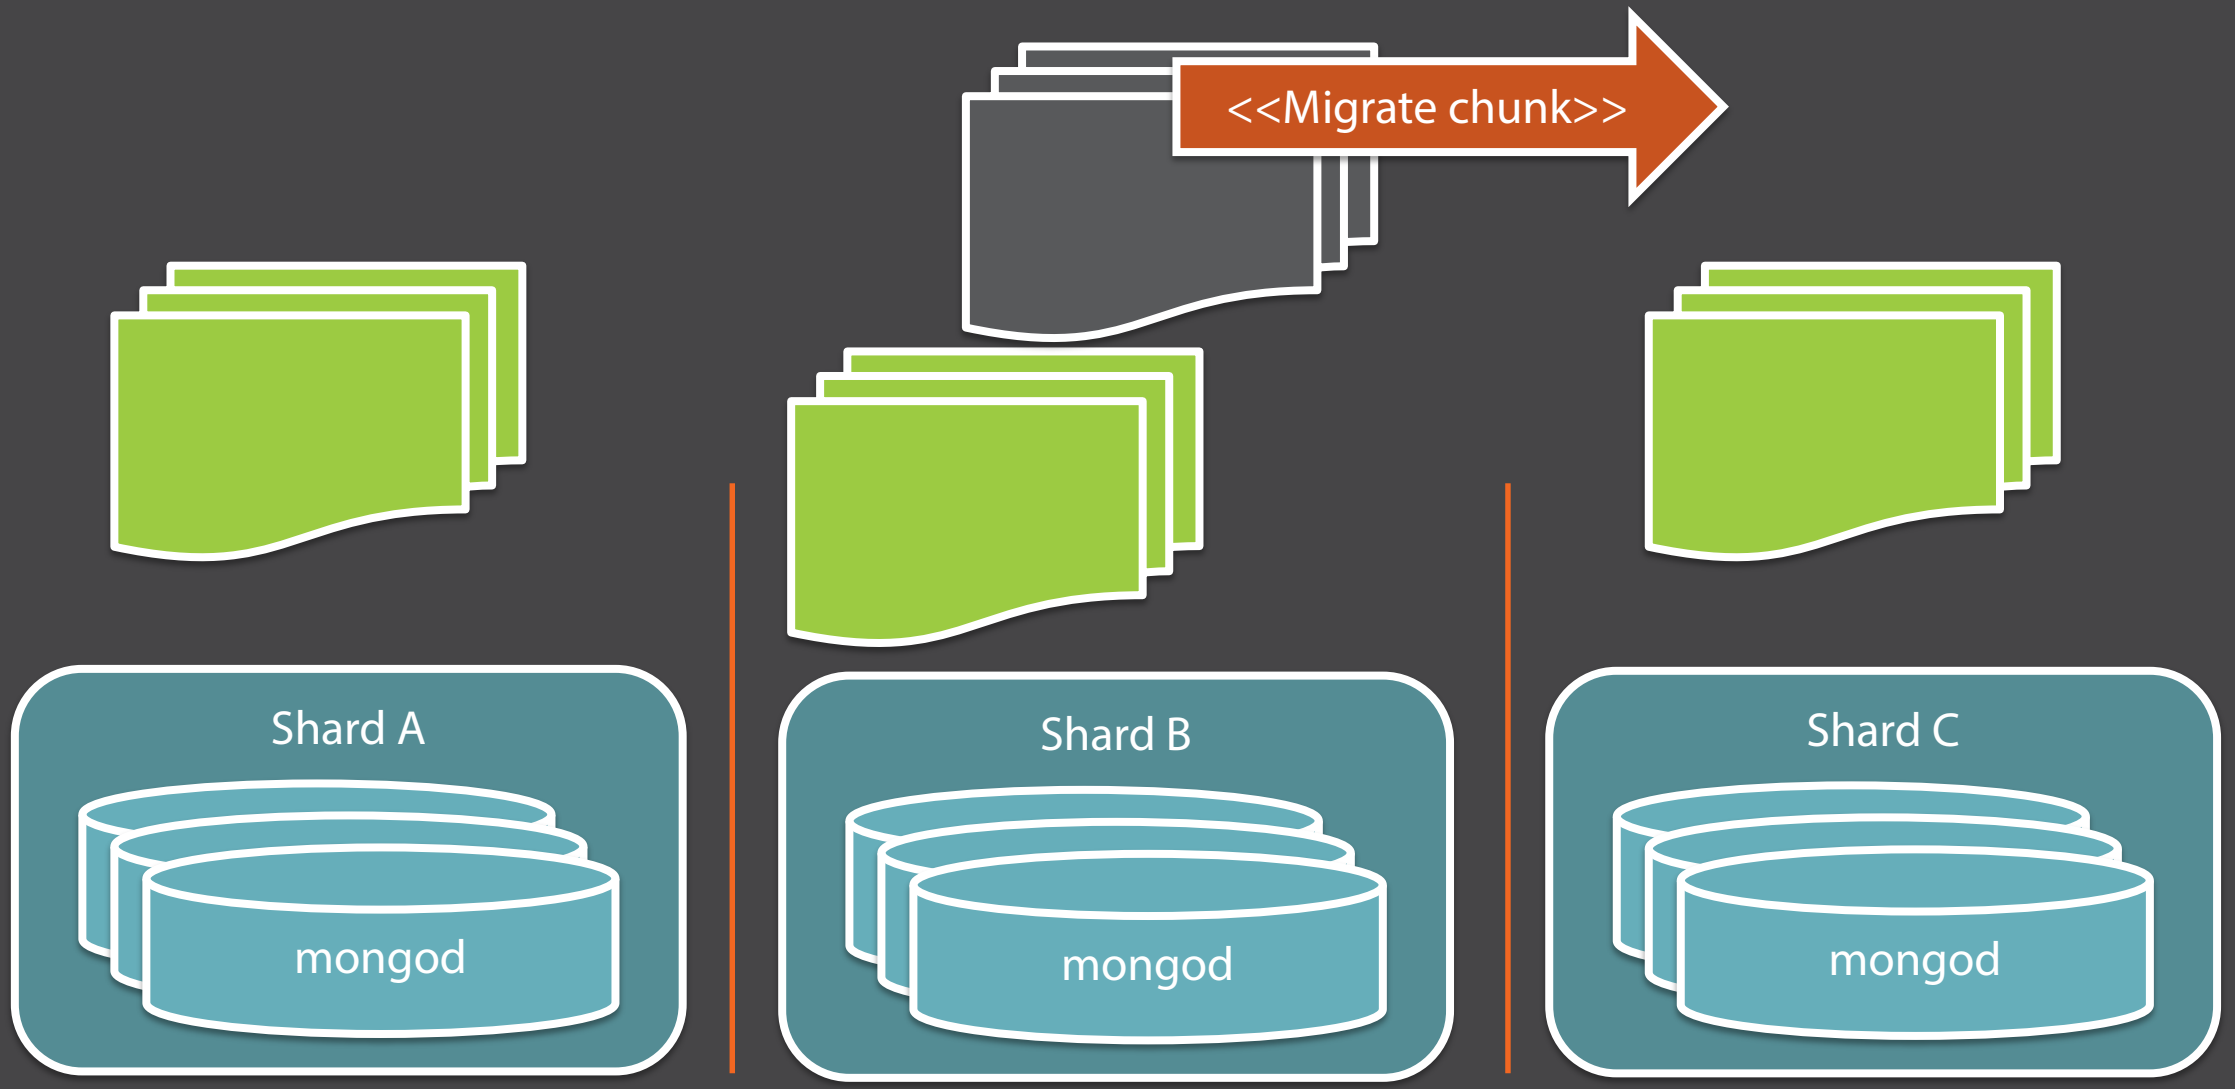

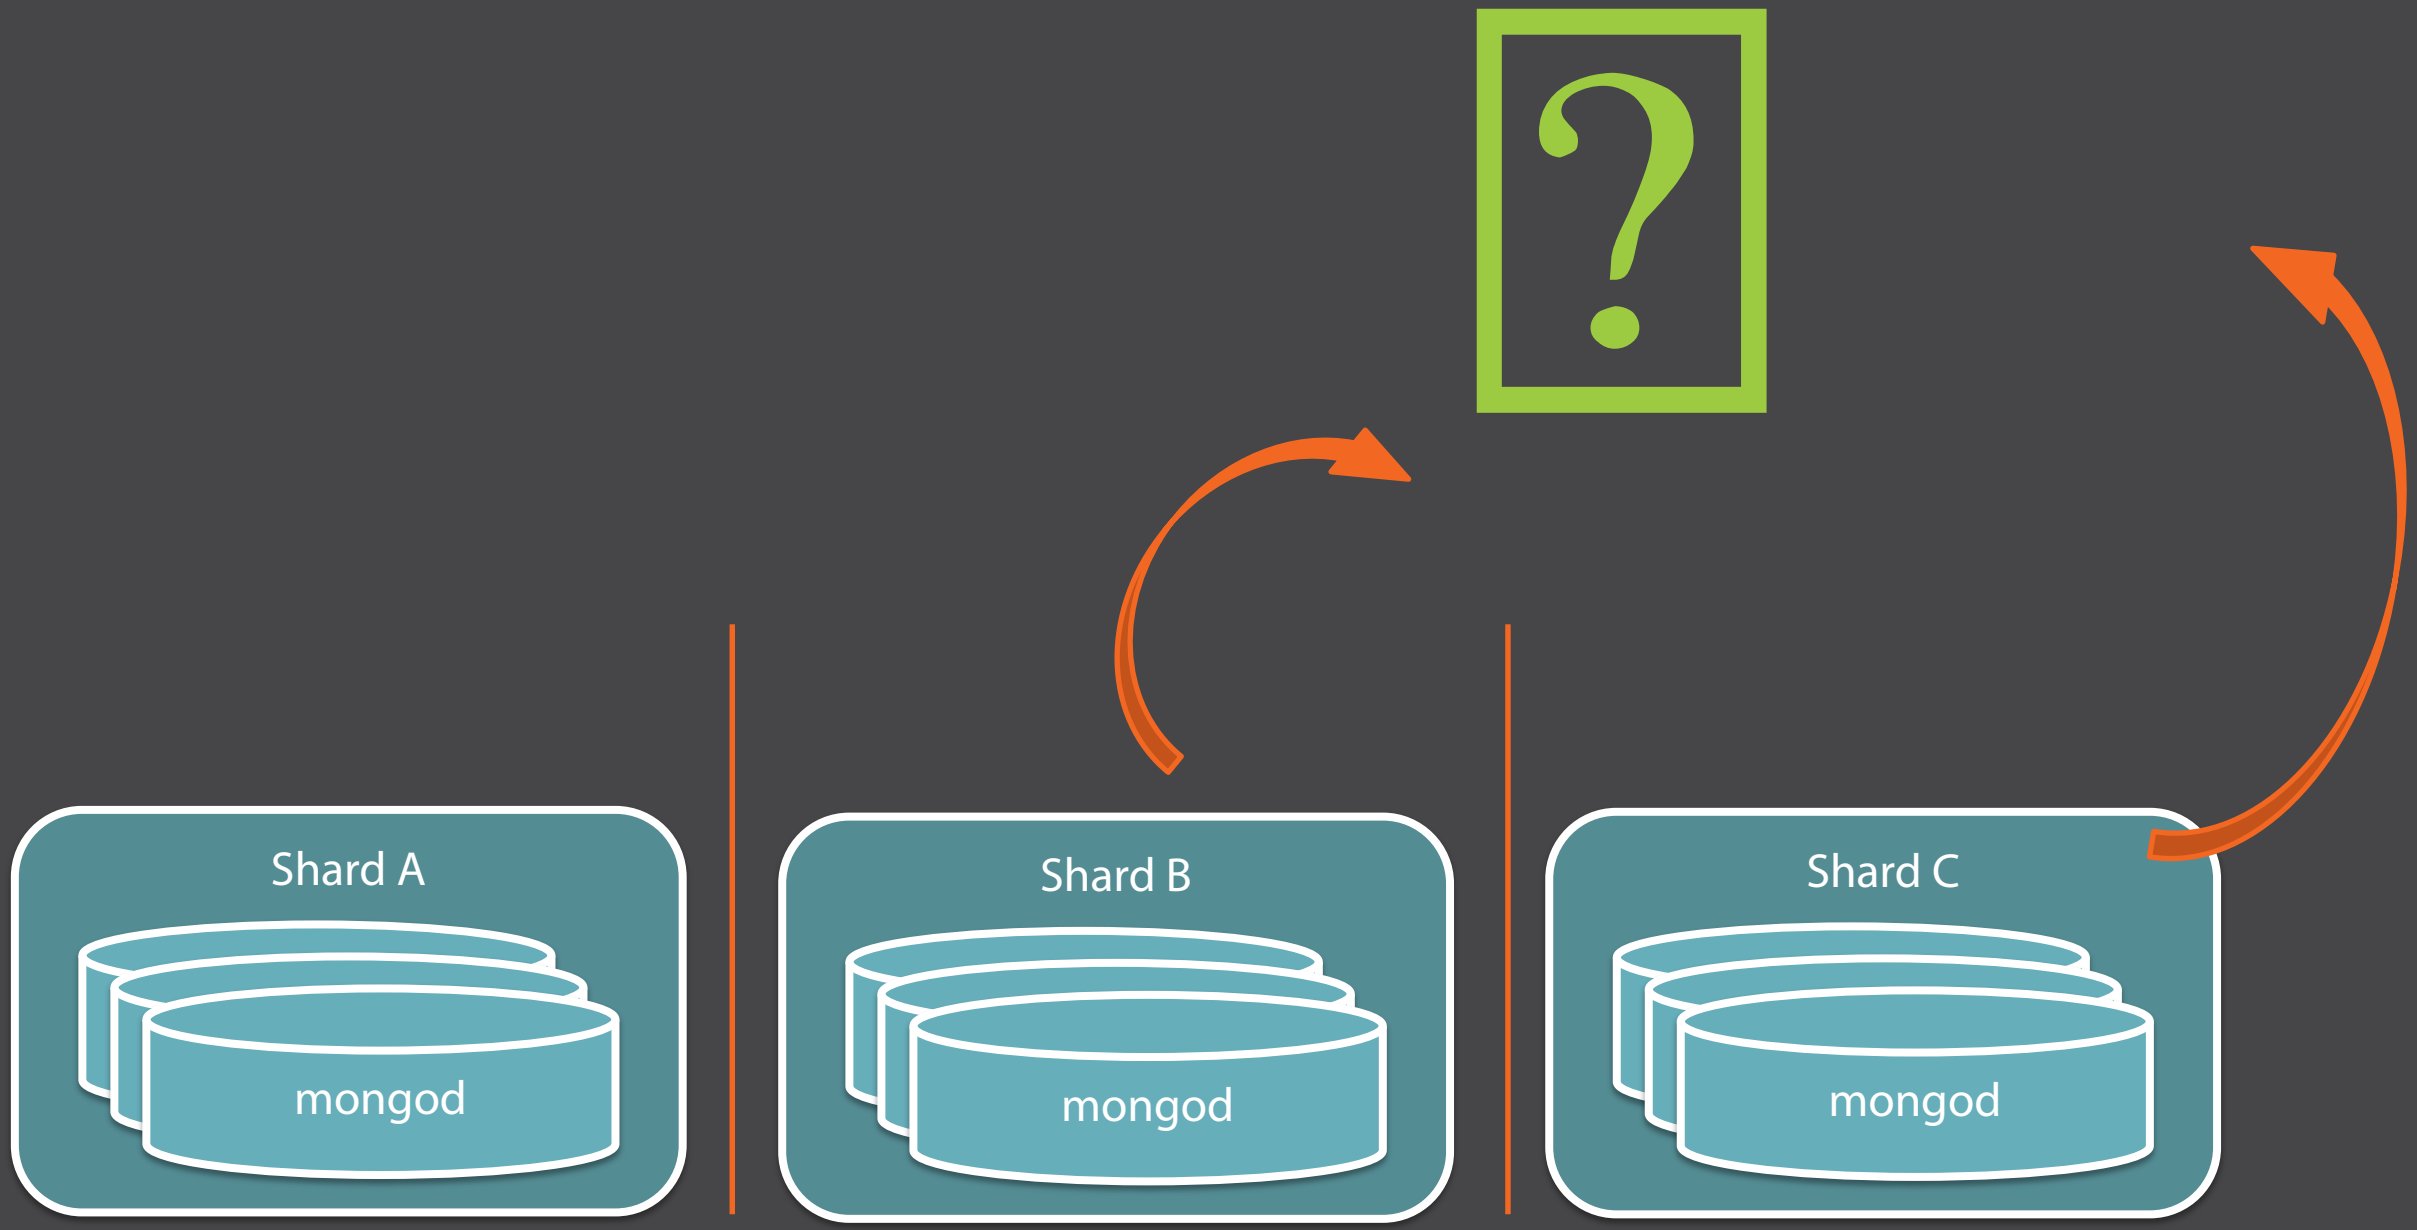

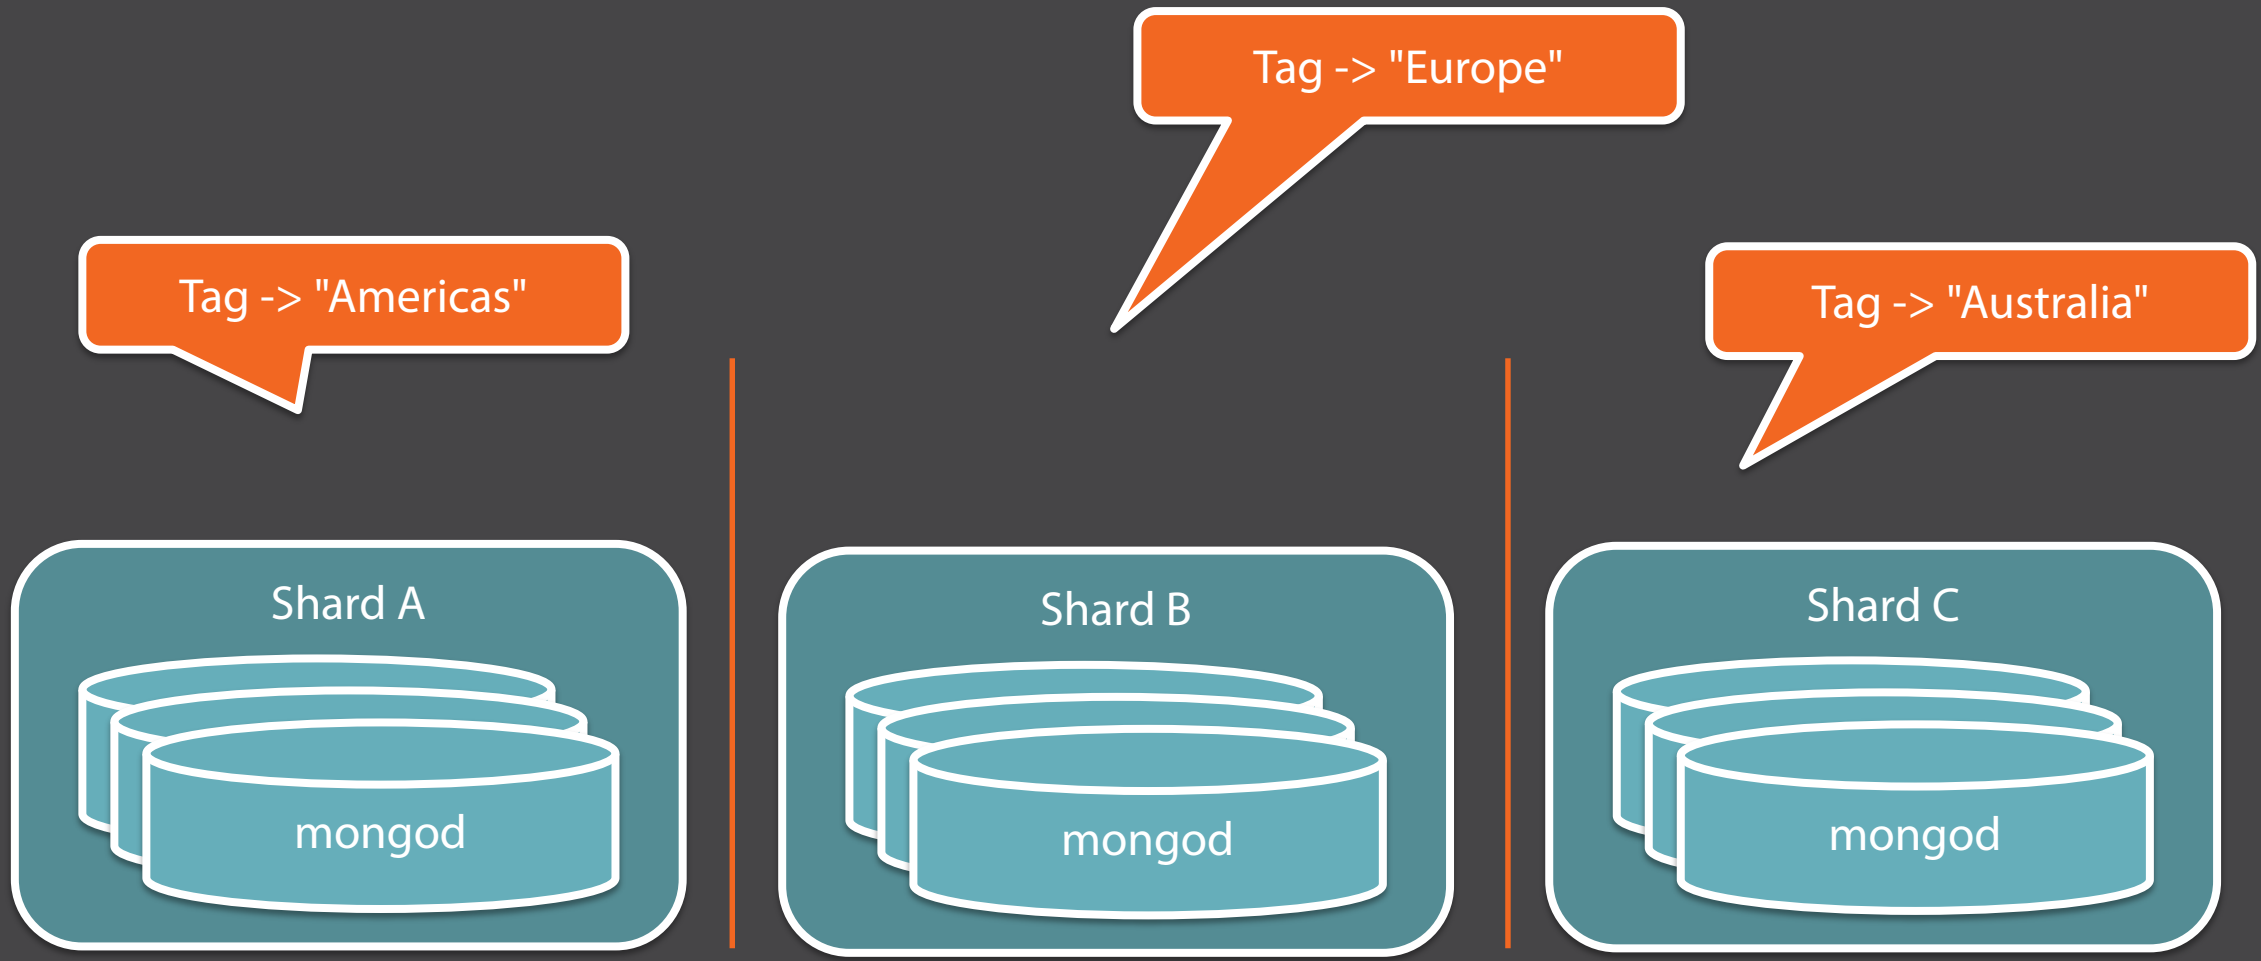

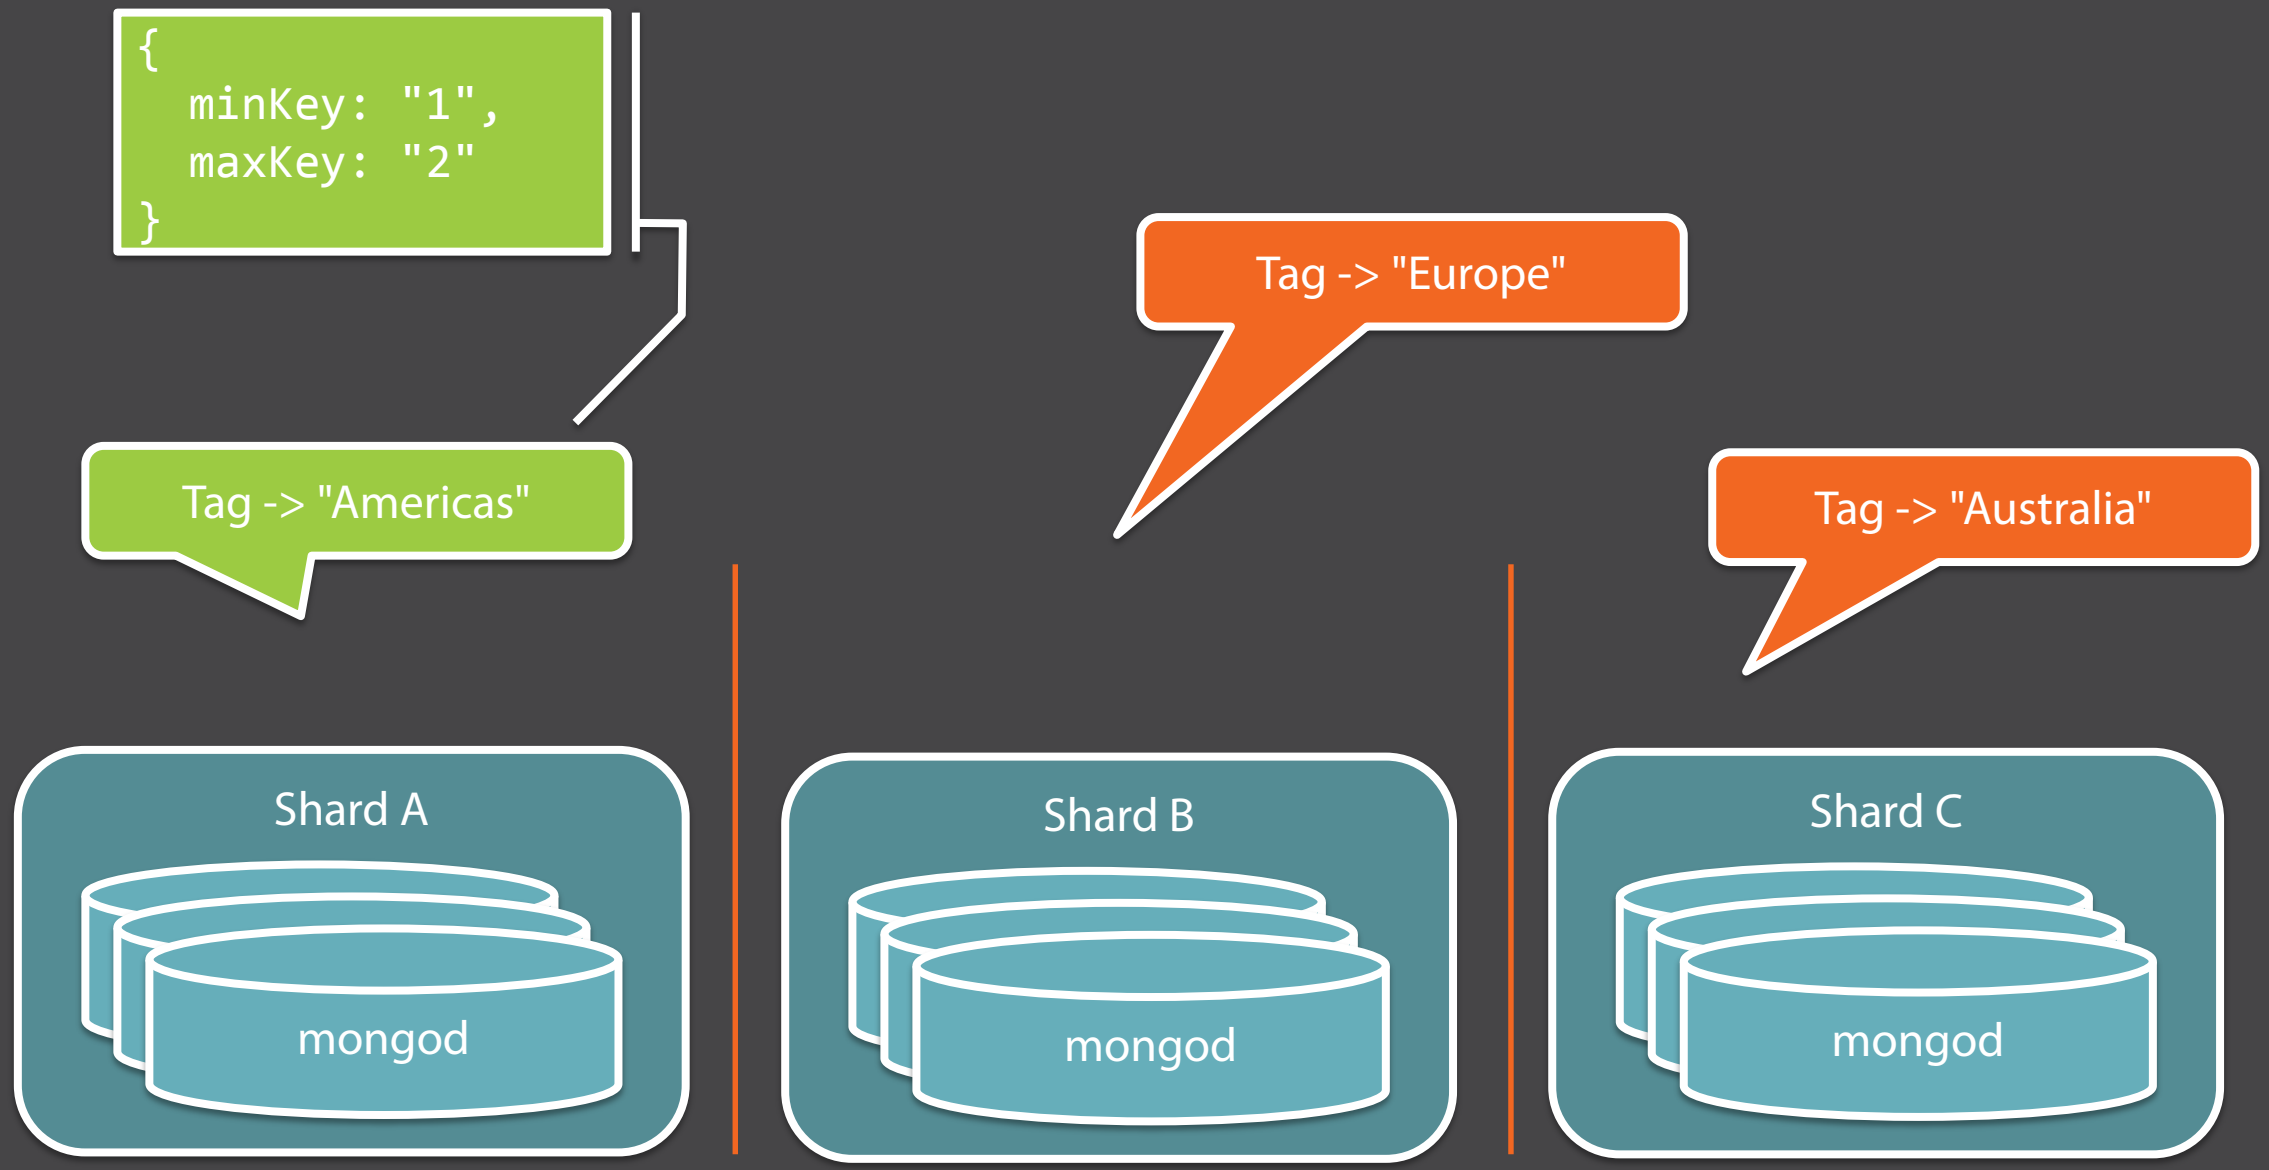

1. Key "**01**" is in range {"01" → "02"}.
2. That range tagged to "Americas".
3. Chunk must live on shard tagged "Americas".

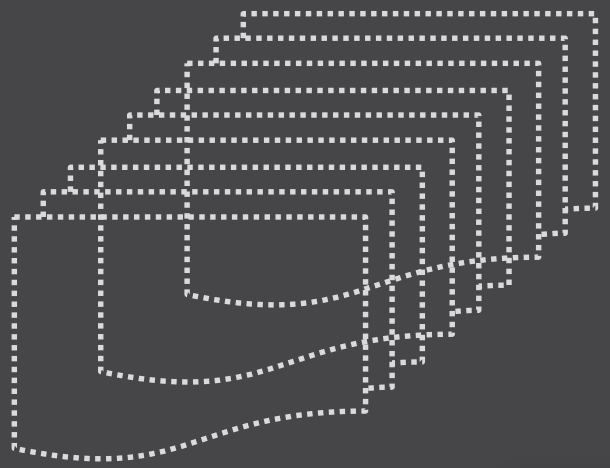

Creating a tag range definition

- Defines a range that is assigned to a tag
- Use only when connected to mongos
- Helps balancer automatically assign chunks to desired shards

Moving a chunk

- Query is equality-match on the shard key
- Moves the whole chunk containing the document found by the query

Shard or Not?

Will it solve a real issue?

Current Capacity?

Application ready for it?

Complexity?

Infrastructure Cost?

When Should I?

Watch for Clues…

Watch for Clues…

Watch for Clues…

Summary